

Business Context

Renewable energy sources play an increasingly important role in the global energy mix, as the effort to reduce the environmental impact of energy production increases.

Out of all the renewable energy alternatives, wind energy is one of the most developed technologies worldwide. The U.S Department of Energy has put together a guide to achieving operational efficiency using predictive maintenance practices.

Predictive maintenance uses sensor information and analysis methods to measure and predict degradation and future component capability. The idea behind predictive maintenance is that failure patterns are predictable and if component failure can be predicted accurately and the component is replaced before it fails, the costs of operation and maintenance will be much lower.

The sensors fitted across different machines involved in the process of energy generation collect data related to various environmental factors (temperature, humidity, wind speed, etc.) and additional features related to various parts of the wind turbine (gearbox, tower, blades, break, etc.).

Objective

"ReneWind" is a company working on improving the machinery/processes involved in the production of wind energy using machine learning and has collected data of generator failure of wind turbines using sensors. They have shared a ciphered version of the data, as the data collected through sensors is confidential (the type of data collected varies with companies). Data has 40 predictors, 20000 observations in the training set and 5000 in the test set.

The objective is to build various classification models, tune them, and find the best one that will help identify failures so that the generators could be repaired before failing/breaking to reduce the overall maintenance cost. The nature of predictions made by the classification model will translate as follows:

- True positives (TP) are failures correctly predicted by the model. These will result in repairing costs.
- False negatives (FN) are real failures where there is no detection by the model. These will result in replacement costs.

- False positives (FP) are detections where there is no failure. These will result in inspection costs.

It is given that the cost of repairing a generator is much less than the cost of replacing it, and the cost of inspection is less than the cost of repair.

"1" in the target variables should be considered as "failure" and "0" represents "No failure".

Problem Statement

In []:

In []:

In []:

In []:

Data Description

The data provided is a transformed version of the original data which was collected using sensors.

- Train.csv - To be used for training and tuning of models.
- Test.csv - To be used only for testing the performance of the final best model.

Both the datasets consist of 40 predictor variables and 1 target variable.

Please read the instructions carefully before starting the project.

This is a commented Jupyter IPython Notebook file in which all the instructions and tasks to be performed are mentioned.

- Blanks '_____' are provided in the notebook that needs to be filled with an appropriate code to get the correct result. With every '_____' blank, there is a comment that briefly describes what needs to be filled in the blank space.

- Identify the task to be performed correctly, and only then proceed to write the required code.
- Fill the code wherever asked by the commented lines like "# write your code here" or "# complete the code". Running incomplete code may throw error.
- Please run the codes in a sequential manner from the beginning to avoid any unnecessary errors.
- Add the results/observations (wherever mentioned) derived from the analysis in the presentation and submit the same.

Installing and Importing the necessary libraries

```
In [ ]: # Installing the libraries with the specified version
        #!pip install --no-deps tensorflow==2.18.0 scikit-learn==1.3.2 matplotlib
```

```
In [1]: #!pip install scikit-learn tensorflow seaborn
```

```
In [3]: import sys
        !{sys.executable} -m pip install scikit-learn
```

Collecting scikit-learn

Downloading scikit_learn-1.8.0-cp313-cp313-macosx_12_0_arm64.whl.metadata (11 kB)

Requirement already satisfied: numpy>=1.24.1 in /opt/miniconda3/lib/python3.13/site-packages (from scikit-learn) (2.3.4)

Collecting scipy>=1.10.0 (from scikit-learn)

Downloading scipy-1.17.1-cp313-cp313-macosx_14_0_arm64.whl.metadata (62 kB)

Collecting joblib>=1.3.0 (from scikit-learn)

Downloading joblib-1.5.3-py3-none-any.whl.metadata (5.5 kB)

Collecting threadpoolctl>=3.2.0 (from scikit-learn)

Downloading threadpoolctl-3.6.0-py3-none-any.whl.metadata (13 kB)

Downloading scikit_learn-1.8.0-cp313-cp313-macosx_12_0_arm64.whl (8.0 MB)

8.0/8.0 MB 6.5 MB/s 0:00:01

1m 6.5 MB/s eta 0:00:01

Downloading joblib-1.5.3-py3-none-any.whl (309 kB)

Downloading scipy-1.17.1-cp313-cp313-macosx_14_0_arm64.whl (20.3 MB)

20.3/20.3 MB 13.8 MB/s 0:00:01

0:014.0 MB/s eta 0:00:01

Downloading threadpoolctl-3.6.0-py3-none-any.whl (18 kB)

Installing collected packages: threadpoolctl, scipy, joblib, scikit-learn

4/4 [scikit-learn]0m 3/4 [scikit-learn]
Successfully installed joblib-1.5.3 scikit-learn-1.8.0 scipy-1.17.1 threadpoolctl-3.6.0

```
In [4]: import sklearn
print(sklearn.__version__)
```

1.8.0

```
In [1]: from sklearn.model_selection import train_test_split
import sklearn

print(sklearn.__version__)
```

1.8.0

```
In [3]: import sys
!{sys.executable} -m pip install seaborn tensorflow
```

Collecting seaborn

Downloading seaborn-0.13.2-py3-none-any.whl.metadata (5.4 kB)

Collecting tensorflow

Using cached tensorflow-2.21.0-cp313-cp313-macosx_12_0_arm64.whl.metadata (4.4 kB)

Requirement already satisfied: numpy!=1.24.0,>=1.20 in /opt/miniconda3/lib/python3.13/site-packages (from seaborn) (2.3.4)

Requirement already satisfied: pandas>=1.2 in /opt/miniconda3/lib/python3.13/site-packages (from seaborn) (2.3.3)

Requirement already satisfied: matplotlib!=3.6.1,>=3.4 in /opt/miniconda3/lib/python3.13/site-packages (from seaborn) (3.10.6)

Collecting absl-py>=1.0.0 (from tensorflow)

Using cached absl_py-2.4.0-py3-none-any.whl.metadata (3.3 kB)

Collecting astunparse>=1.6.0 (from tensorflow)

Using cached astunparse-1.6.3-py2.py3-none-any.whl.metadata (4.4 kB)

Collecting flatbuffers>=25.9.23 (from tensorflow)

Using cached flatbuffers-25.12.19-py2.py3-none-any.whl.metadata (1.0 kB)

Collecting gast!=0.5.0,!0.5.1,!0.5.2,>=0.2.1 (from tensorflow)

Using cached gast-0.7.0-py3-none-any.whl.metadata (1.5 kB)

Collecting google_pasta>=0.1.1 (from tensorflow)

Using cached google_pasta-0.2.0-py3-none-any.whl.metadata (814 bytes)

Collecting libclang>=13.0.0 (from tensorflow)

Using cached libclang-18.1.1-1-py2.py3-none-macosx_11_0_arm64.whl.metadata (5.2 kB)

Collecting opt_einsum>=2.3.2 (from tensorflow)

Using cached opt_einsum-3.4.0-py3-none-any.whl.metadata (6.3 kB)

Requirement already satisfied: packaging in /opt/miniconda3/lib/python3.13/site-packages (from tensorflow) (25.0)

Collecting protobuf<8.0.0,>=6.31.1 (from tensorflow)

Using cached protobuf-7.35.0-cp310-abi3-macosx_10_9_universal2.whl.metadata (595 bytes)

Requirement already satisfied: requests<3,>=2.21.0 in /opt/miniconda3/lib/python3.13/site-packages (from tensorflow) (2.32.5)

Requirement already satisfied: setuptools in /opt/miniconda3/lib/python3.13/site-packages (from tensorflow) (80.9.0)

Requirement already satisfied: six>=1.12.0 in /opt/miniconda3/lib/python3.13/site-packages (from tensorflow) (1.17.0)

```

Collecting termcolor>=1.1.0 (from tensorflow)
  Using cached termcolor-3.3.0-py3-none-any.whl.metadata (6.5 kB)
Requirement already satisfied: typing_extensions>=3.6.6 in /opt/miniconda3/lib/python3.13/site-packages (from tensorflow) (4.15.0)
Collecting wrapt>=1.11.0 (from tensorflow)
  Downloading wrapt-2.2.1-cp313-cp313-macosx_11_0_arm64.whl.metadata (7.4 kB)
Collecting grpcio<2.0,>=1.24.3 (from tensorflow)
  Using cached grpcio-1.80.0-cp313-cp313-macosx_11_0_universal2.whl.metadata (3.8 kB)
Collecting keras>=3.12.0 (from tensorflow)
  Using cached keras-3.14.1-py3-none-any.whl.metadata (6.3 kB)
Collecting h5py<3.15.0,>=3.11.0 (from tensorflow)
  Using cached h5py-3.14.0-cp313-cp313-macosx_11_0_arm64.whl.metadata (2.7 kB)
Collecting ml_dtypes<1.0.0,>=0.5.1 (from tensorflow)
  Using cached ml_dtypes-0.5.4-cp313-cp313-macosx_10_13_universal2.whl.metadata (8.9 kB)
Requirement already satisfied: charset_normalizer<4,>=2 in /opt/miniconda3/lib/python3.13/site-packages (from requests<3,>=2.21.0->tensorflow) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in /opt/miniconda3/lib/python3.13/site-packages (from requests<3,>=2.21.0->tensorflow) (3.7)
Requirement already satisfied: urllib3<3,>=1.21.1 in /opt/miniconda3/lib/python3.13/site-packages (from requests<3,>=2.21.0->tensorflow) (2.5.0)
Requirement already satisfied: certifi>=2017.4.17 in /opt/miniconda3/lib/python3.13/site-packages (from requests<3,>=2.21.0->tensorflow) (2025.10.5)
Requirement already satisfied: wheel<1.0,>=0.23.0 in /opt/miniconda3/lib/python3.13/site-packages (from astunparse>=1.6.0->tensorflow) (0.45.1)
Requirement already satisfied: rich in /opt/miniconda3/lib/python3.13/site-packages (from keras>=3.12.0->tensorflow) (14.2.0)
Collecting namex (from keras>=3.12.0->tensorflow)
  Using cached namex-0.1.0-py3-none-any.whl.metadata (322 bytes)
Collecting optree (from keras>=3.12.0->tensorflow)
  Using cached optree-0.19.1-cp313-cp313-macosx_11_0_arm64.whl.metadata (32 kB)
Requirement already satisfied: contourpy>=1.0.1 in /opt/miniconda3/lib/python3.13/site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.3.1)
Requirement already satisfied: cycler>=0.10 in /opt/miniconda3/lib/python3.13/site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (0.11.0)
Requirement already satisfied: fonttools>=4.22.0 in /opt/miniconda3/lib/python3.13/site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (4.60.1)
Requirement already satisfied: kiwisolver>=1.3.1 in /opt/miniconda3/lib/python3.13/site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.4.8)
Requirement already satisfied: pillow>=8 in /opt/miniconda3/lib/python3.13/site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (12.0.0)

```

Requirement already satisfied: pyparsing<=2.3.1 in /opt/miniconda3/lib/python3.13/site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (3.2.5)

Requirement already satisfied: python-dateutil<=2.7 in /opt/miniconda3/lib/python3.13/site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (2.9.0.post0)

Requirement already satisfied: pytz<=2020.1 in /opt/miniconda3/lib/python3.13/site-packages (from pandas>=1.2->seaborn) (2025.2)

Requirement already satisfied: tzdata<=2022.7 in /opt/miniconda3/lib/python3.13/site-packages (from pandas>=1.2->seaborn) (2025.2)

Requirement already satisfied: markdown-it-py<=2.2.0 in /opt/miniconda3/lib/python3.13/site-packages (from rich->keras>=3.12.0->tensorflow) (4.0.0)

Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /opt/miniconda3/lib/python3.13/site-packages (from rich->keras>=3.12.0->tensorflow) (2.19.1)

Requirement already satisfied: mdurl<=0.1 in /opt/miniconda3/lib/python3.13/site-packages (from markdown-it-py>=2.2.0->rich->keras>=3.12.0->tensorflow) (0.1.2)

Downloading seaborn-0.13.2-py3-none-any.whl (294 kB)

Using cached tensorflow-2.21.0-cp313-cp313-macosx_12_0_arm64.whl (223.8 MB)

Using cached grpcio-1.80.0-cp313-cp313-macosx_11_0_universal2.whl (12.0 MB)

Using cached h5py-3.14.0-cp313-cp313-macosx_11_0_arm64.whl (2.8 MB)

Using cached ml_dtypes-0.5.4-cp313-cp313-macosx_10_13_universal2.whl (676 kB)

Using cached protobuf-7.35.0-cp310-abi3-macosx_10_9_universal2.whl (433 kB)

Using cached absl_py-2.4.0-py3-none-any.whl (135 kB)

Using cached astunparse-1.6.3-py2.py3-none-any.whl (12 kB)

Using cached flatbuffers-25.12.19-py2.py3-none-any.whl (26 kB)

Using cached gast-0.7.0-py3-none-any.whl (22 kB)

Using cached google_pasta-0.2.0-py3-none-any.whl (57 kB)

Using cached keras-3.14.1-py3-none-any.whl (1.6 MB)

Using cached libclang-18.1.1-1-py2.py3-none-macosx_11_0_arm64.whl (25.8 MB)

Using cached opt_einsum-3.4.0-py3-none-any.whl (71 kB)

Using cached termcolor-3.3.0-py3-none-any.whl (7.7 kB)

Downloading wrapt-2.2.1-cp313-cp313-macosx_11_0_arm64.whl (81 kB)

Using cached namex-0.1.0-py3-none-any.whl (5.9 kB)

Using cached optree-0.19.1-cp313-cp313-macosx_11_0_arm64.whl (398 kB)

Installing collected packages: namex, libclang, flatbuffers, wrapt, termcolor, protobuf, optree, opt_einsum, ml_dtypes, h5py, grpcio, google_pasta, gast, astunparse, absl-py, seaborn, keras, tensorflow

18/18 [tensorflow] 37m — 17/18 [tensorflow]

Successfully installed absl-py-2.4.0 astunparse-1.6.3 flatbuffers-25.12.19 gast-0.7.0 google_pasta-0.2.0 grpcio-1.80.0 h5py-3.14.0 keras-3.14.1 libclang-18.1.1 ml_dtypes-0.5.4 namex-0.1.0 opt_einsum-3.4.0 optree-0.19.1 protobuf-7.35.0 seaborn-0.13.2 tensorflow-2.21.0 termcolor-3.3.0 wrapt-2.2.1

```
In [1]: import seaborn as sns
import tensorflow as tf

print("Seaborn OK")
print("TensorFlow version:", tf.__version__)
```

Seaborn OK

TensorFlow version: 2.21.0

In []:

Note:

- After running the above cell, kindly restart the runtime (for Google Colab) or notebook kernel (for Jupyter Notebook), and run all cells sequentially from the next cell.
- On executing the above line of code, you might see a warning regarding package dependencies. This error message can be ignored as the above code ensures that all necessary libraries and their dependencies are maintained to successfully execute the code in **this notebook**.

```
In [2]: # Library for data manipulation and analysis.
import pandas as pd
# Fundamental package for scientific computing.
import numpy as np
#splitting datasets into training and testing sets.
from sklearn.model_selection import train_test_split
#Imports tools for data preprocessing including label encoding, one-hot
from sklearn.preprocessing import LabelEncoder, OneHotEncoder, Standard
#Imports a class for imputing missing values in datasets.
from sklearn.impute import SimpleImputer
#Imports the Matplotlib library for creating visualizations.
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
# Imports the Seaborn library for statistical data visualization.
import seaborn as sns
# Time related functions.
import time
#Imports functions for evaluating the performance of machine learning
from sklearn.metrics import confusion_matrix, f1_score, accuracy_score,
#Imports metrics from
from sklearn import metrics

#Imports the tensorflow, keras and layers.
import tensorflow
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.layers import Dense, Input, Dropout, BatchNormali
from tensorflow.keras import backend
```

```
# to suppress unnecessary warnings
import warnings
warnings.filterwarnings("ignore")
```

Loading the Data

```
In [3]: # uncomment and run the following lines for Google Colab
# from google.colab import drive
# drive.mount('/content/drive')
```

```
In [5]: df = pd.read_csv("Train.csv")    # Complete the code to import the tra
df_test = pd.read_csv("Test.csv")      # Complete the code to import the
```

Data Overview

The initial steps to get an overview of any dataset is to:

- observe the first few rows of the dataset, to check whether the dataset has been loaded properly or not
- get information about the number of rows and columns in the dataset
- find out the data types of the columns to ensure that data is stored in the preferred format and the value of each property is as expected.
- check the statistical summary of the dataset to get an overview of the numerical columns of the data

Checking the shape of the dataset

```
In [6]: # Checking the number of rows and columns in the training data
df.shape # Complete the code to print the shape of the train data
```

```
Out[6]: (20000, 41)
```

```
In [7]: # Checking the number of rows and columns in the test data
df_test.shape # Complete the code to print the shape of the test data
```

```
Out[7]: (5000, 41)
```

```
In [8]: # let's create a copy of the training data
data = df.copy()
```

```
In [9]: # let's create a copy of the testing data
data_test = df_test.copy()
```


Observations

- The training dataset contains **20,000 rows** and **41 columns**.
- The test dataset contains **5,000 rows** and **41 columns**.
- Since both datasets have the same number of columns, the structure of the training and test data is consistent.
- The dataset contains **40 predictor variables** and **1 target variable**, which matches the project description.
- Copies of both datasets were created (`data` and `data_test`) to preserve the original datasets and avoid accidental modification during preprocessing and analysis.

Displaying the first few rows of the dataset

```
In [10]: # let's view the first 5 rows of the data
data.head() # Complete the code to view the first five rows of the tra
```

```
Out[10]:
```

	V1	V2	V3	V4	V5	V6	V7
0	-4.464606	-4.679129	3.101546	0.506130	-0.221083	-2.032511	-2.910870
1	3.365912	3.653381	0.909671	-1.367528	0.332016	2.358938	0.732600
2	-3.831843	-5.824444	0.634031	-2.418815	-1.773827	1.016824	-2.098941
3	1.618098	1.888342	7.046143	-1.147285	0.083080	-1.529780	0.207309
4	-0.111440	3.872488	-3.758361	-2.982897	3.792714	0.544960	0.205433

5 rows x 41 columns

```
In [11]: #viewing first 5 rows of the test data
data_test.head() # Complete the code to view the first five rows of th
```

Out[11]:

	V1	V2	V3	V4	V5	V6	V7
0	-0.613489	-3.819640	2.202302	1.300420	-1.184929	-4.495964	-1.835817
1	0.389608	-0.512341	0.527053	-2.576776	-1.016766	2.235112	-0.441307
2	-0.874861	-0.640632	4.084202	-1.590454	0.525855	-1.957592	-0.695367
3	0.238384	1.458607	4.014528	2.534478	1.196987	-3.117330	-0.924035
4	5.828225	2.768260	-1.234530	2.809264	-1.641648	-1.406698	0.568643

5 rows x 41 columns

Checking the data types of the columns in the dataset

```
In [13]: # let's check the data types of the columns in the dataset
data.dtypes # Complete the code to view the data types of the columns
```

```
Out[13]: V1          float64
          V2          float64
          V3          float64
          V4          float64
          V5          float64
          V6          float64
          V7          float64
          V8          float64
          V9          float64
          V10         float64
          V11         float64
          V12         float64
          V13         float64
          V14         float64
          V15         float64
          V16         float64
          V17         float64
          V18         float64
          V19         float64
          V20         float64
          V21         float64
          V22         float64
          V23         float64
          V24         float64
          V25         float64
          V26         float64
          V27         float64
          V28         float64
          V29         float64
          V30         float64
          V31         float64
          V32         float64
          V33         float64
          V34         float64
          V35         float64
          V36         float64
          V37         float64
          V38         float64
          V39         float64
          V40         float64
          Target      int64
dtype: object
```

- Converting Target column to float

```
In [14]: data['Target'] = data['Target'].astype(float)
```

Now checking for test data

```
In [15]: data_test.dtypes # Complete the code to view the data types of the col
```

```
Out[15]: V1          float64
          V2          float64
          V3          float64
          V4          float64
          V5          float64
          V6          float64
          V7          float64
          V8          float64
          V9          float64
          V10         float64
          V11         float64
          V12         float64
          V13         float64
          V14         float64
          V15         float64
          V16         float64
          V17         float64
          V18         float64
          V19         float64
          V20         float64
          V21         float64
          V22         float64
          V23         float64
          V24         float64
          V25         float64
          V26         float64
          V27         float64
          V28         float64
          V29         float64
          V30         float64
          V31         float64
          V32         float64
          V33         float64
          V34         float64
          V35         float64
          V36         float64
          V37         float64
          V38         float64
          V39         float64
          V40         float64
          Target      int64
dtype: object
```

Converting Target to float

```
In [16]: data_test['Target'] = data_test['Target'].astype(float)
```

Checking for duplicate values

```
In [17]: # let's check for duplicate values in the data
```

```
data.duplicated().sum() # Complete the code to check for duplicate val
```

```
Out[17]: np.int64(0)
```

Checking for missing values

```
In [18]: # let's check for missing values in the data  
data.isnull().sum()
```

```
Out[18]: V1          18  
V2          18  
V3           0  
V4           0  
V5           0  
V6           0  
V7           0  
V8           0  
V9           0  
V10          0  
V11          0  
V12          0  
V13          0  
V14          0  
V15          0  
V16          0  
V17          0  
V18          0  
V19          0  
V20          0  
V21          0  
V22          0  
V23          0  
V24          0  
V25          0  
V26          0  
V27          0  
V28          0  
V29          0  
V30          0  
V31          0  
V32          0  
V33          0  
V34          0  
V35          0  
V36          0  
V37          0  
V38          0  
V39          0  
V40          0  
Target       0  
dtype: int64
```

```
In [21]: # let's check for missing values in the test data
data_test.isnull().sum() #
```

```
Out[21]: V1          5
          V2          6
          V3          0
          V4          0
          V5          0
          V6          0
          V7          0
          V8          0
          V9          0
          V10         0
          V11         0
          V12         0
          V13         0
          V14         0
          V15         0
          V16         0
          V17         0
          V18         0
          V19         0
          V20         0
          V21         0
          V22         0
          V23         0
          V24         0
          V25         0
          V26         0
          V27         0
          V28         0
          V29         0
          V30         0
          V31         0
          V32         0
          V33         0
          V34         0
          V35         0
          V36         0
          V37         0
          V38         0
          V39         0
          V40         0
          Target      0
          dtype: int64
```

Statistical summary of the dataset

```
In [20]: # let's view the statistical summary of the numerical columns in the d
data.describe() # Complete the code to view the statistical summary of
```

Out [20]:

	V1	V2	V3	V4	V
count	19982.000000	19982.000000	20000.000000	20000.000000	20000.000000
mean	-0.271996	0.440430	2.484699	-0.083152	-0.053750
std	3.441625	3.150784	3.388963	3.431595	2.104800
min	-11.876451	-12.319951	-10.708139	-15.082052	-8.603360
25%	-2.737146	-1.640674	0.206860	-2.347660	-1.535600
50%	-0.747917	0.471536	2.255786	-0.135241	-0.101950
75%	1.840112	2.543967	4.566165	2.130615	1.340480
max	15.493002	13.089269	17.090919	13.236381	8.133750

8 rows x 41 columns

In []:

In []:

Data Overview and Initial Findings

Missing Value Analysis

The missing value analysis shows that:

- Only two variables, **V1** and **V2**, contain missing values.
- Both variables have **18 missing observations** each.
- All remaining predictor variables and the target variable contain no missing values.
- Since the dataset contains 20,000 observations, the percentage of missing values is extremely small.

Importance of Missing Value Analysis

Checking for missing values is important because:

- Missing values can negatively impact model performance and may cause errors during training.
- Many machine learning algorithms cannot process null values directly.
- Proper treatment of missing values improves model reliability and prediction accuracy.

- Since the number of missing values is very low, simple imputation techniques such as mean or median imputation can be applied effectively without significantly changing the data distribution.
-

Statistical Summary Observations

The statistical summary provides important insights into the numerical variables in the dataset.

Key observations include:

- The dataset contains **40 predictor variables** and **1 target variable**.
- Several variables have relatively large standard deviations, indicating high variability in sensor readings.
- The wide gap between minimum and maximum values suggests the possible presence of outliers in multiple features.
- The target variable has a mean value of approximately **0.0555**, indicating that only about **5.5%** of observations correspond to generator failures.

Importance of Statistical Summary

The statistical summary is important because:

- It helps understand the distribution and spread of the variables.
- It helps identify potential outliers that may affect model performance.
- It provides insight into whether scaling or normalization may be required before modeling.
- The low proportion of failure cases indicates that the dataset is **highly imbalanced**.

Business and Modeling Implications

Class imbalance is critical in this project because:

- A model trained on imbalanced data may become biased toward predicting the majority class ("No Failure").
- This can result in poor detection of actual failures.
- False Negatives are especially costly because undetected failures can lead to expensive generator replacement and operational downtime.
- Therefore, the final model should prioritize identifying failures effectively by improving Recall and reducing False Negatives, even if it results in slightly

higher inspection costs from False Positives.

- Because of this imbalance, evaluation metrics such as **Recall, Precision, F1-score, and Confusion Matrix** are more meaningful than accuracy alone.

Exploratory Data Analysis

Univariate analysis

In [22]: *# function to plot a boxplot and a histogram along the same scale.*

```
def histogram_boxplot(data, feature, figsize=(12, 7), kde=False, bins=
    """
    Boxplot and histogram combined

    data: dataframe
    feature: dataframe column
    figsize: size of figure (default (12,7))
    kde: whether to the show density curve (default False)
    bins: number of bins for histogram (default None)
    """
    f2, (ax_box2, ax_hist2) = plt.subplots(
        nrows=2, # Number of rows of the subplot grid= 2
        sharex=True, # x-axis will be shared among all subplots
        gridspec_kw={"height_ratios": (0.25, 0.75)},
        figsize=figsize,
    ) # creating the 2 subplots
    sns.boxplot(
        data=data, x=feature, ax=ax_box2, showmeans=True, color="violet"
    ) # boxplot will be created and a star will indicate the mean val
    sns.histplot(
        data=data, x=feature, kde=kde, ax=ax_hist2, bins=bins, palette
    ) if bins else sns.histplot(
        data=data, x=feature, kde=kde, ax=ax_hist2
    ) # For histogram
    ax_hist2.axvline(
        data[feature].mean(), color="green", linestyle="--"
    ) # Add mean to the histogram
    ax_hist2.axvline(
        data[feature].median(), color="black", linestyle="--"
    ) # Add median to the histogram
```

In []:

Exploratory Data Analysis (EDA) –

Univariate Analysis and Business Interpretation

Overview

Univariate analysis was conducted on all numerical variables (V1–V40) and the target variable to better understand:

- the distribution of each feature,
- variability in sensor behavior,
- the presence of outliers,
- skewness,
- and potential preprocessing requirements before machine learning modeling.

Histograms and boxplots were used to evaluate the distribution and spread of each variable individually.

Dataset Characteristics

The dataset contains:

Training Dataset

- 20,000 observations
- 40 numerical predictor variables
- 1 binary target variable

Testing Dataset

- 5,000 observations
- 40 numerical predictor variables
- 1 binary target variable

The target variable is:

- **0 = No Failure**
 - **1 = Failure**
-

Missing Value Analysis

Observations

- Only variables **V1** and **V2** contain missing values.
- Both variables contain only **18 missing observations** each.
- All remaining variables and the target variable contain no missing values.

The percentage of missing values is extremely small relative to the overall dataset size.

Why Missing Value Analysis Is Important

Checking for missing values is important because:

- Many machine learning algorithms cannot process null values directly.
- Missing data can negatively impact prediction accuracy and model stability.
- Improper handling of missing values may introduce bias into the model.

Since missingness is minimal, simple imputation methods such as:

- mean imputation,
- median imputation,
- or SimpleImputer

can effectively handle the issue without significantly altering the data distribution.

Business Impact of Missing Values

Proper treatment of missing sensor readings helps:

- improve prediction reliability,
- reduce operational risk,
- and support more accurate maintenance decisions.

Ignoring missing values could increase the likelihood of inaccurate predictions and unexpected turbine failures.

Univariate Analysis – Variables V1 to V21

Observations

- Most variables from **V1 to V21** follow approximately:
 - bell-shaped,
 - normal,
 - or near-normal distributions.
- Several variables display:
 - slight positive skewness,
 - slight negative skewness,
 - and moderate asymmetry.
- Many variables contain visible outliers on both lower and upper ends of the distributions.
- The mean and median are relatively close for many variables, indicating moderate symmetry and stable sensor behavior.
- Variables such as:
 - **V3**
 - **V12**
 - **V13**
 - **V16**
 - **V21**

show:

- wider spreads,
 - heavier tails,
 - and higher variability.
 - Histograms indicate that most variables are continuous numerical features with varying scales and ranges.
 - Outliers are common across many variables, which is expected in real-world industrial sensor data.
-

Why This Is Important for EDA

Performing univariate analysis is important because it helps:

- understand the behavior of each feature individually,
- identify skewness and unusual patterns,
- detect extreme values and operational anomalies,
- and determine whether preprocessing is required.

This analysis helps determine whether:

- scaling,
- normalization,
- transformation,
- or robust preprocessing techniques

may improve machine learning performance.

Since neural networks are sensitive to feature scale and distribution, understanding variable behavior is critical before model development.

Business Interpretation – Variables V1 to V21

From a business perspective:

- Extreme sensor readings may indicate:
 - abnormal turbine behavior,
 - overheating,
 - vibration abnormalities,
 - mechanical stress,
 - or early signs of equipment degradation.
- Variables with high variability and frequent outliers may contain strong predictive signals for turbine failure.
- Detecting abnormal operating conditions early can help:
 - reduce downtime,
 - lower replacement costs,
 - improve turbine reliability,

- and optimize maintenance planning.
-

Additional Univariate Analysis – Variables V20 to V39

Observations

Variables from **V20 to V39** also display distributions that are mostly:

- approximately normal,
- bell-shaped,
- or near-symmetric.

Several variables exhibit moderate skewness, particularly:

- **V21**
- **V24**
- **V26**
- **V27**
- **V31**
- **V32**
- **V36**

These variables show:

- longer tails,
- asymmetric distributions,
- and larger variability.

Additional observations include:

- Many variables contain substantial outliers on both sides of the distribution.
 - Variables such as **V32** and **V36** display especially wide ranges and heavier tails.
 - Some variables are tightly concentrated around the center, while others are more dispersed.
 - Mean and median lines remain relatively close for several variables despite the presence of outliers.
 - Mild positive skewness is observed in some variables where right tails extend further than left tails.
-

Why This Is Important for EDA

Analyzing these variables is important because:

- it helps identify whether the data satisfies assumptions commonly used in machine learning and statistical modeling,
- highly skewed variables may require preprocessing,
- extreme outliers may strongly influence model performance,
- and wider distributions may indicate abnormal operational conditions.

Variables with heavy tails and large spreads may contain useful information related to:

- turbine instability,
- mechanical stress,
- and generator failure behavior.

This analysis also helps determine whether:

- feature engineering,
- dimensionality reduction,
- or advanced preprocessing

may improve model performance later in the project.

Business Interpretation – Variables V20 to V39

From a business and operational perspective:

- Extreme sensor readings may represent:
 - overheating,
 - vibration irregularities,
 - unstable operating conditions,
 - or early component degradation.
- Variables with larger variability and extreme outliers may be particularly useful for predicting turbine failure events.
- Detecting unusual sensor behavior early can help:
 - reduce costly unplanned downtime,

- improve maintenance scheduling,
- and support predictive maintenance strategies.

Because turbine failures are rare but expensive, identifying subtle abnormal patterns in these variables may significantly improve maintenance decision-making.

Variable V40 Analysis

Observations

Variable **V40** follows an approximately bell-shaped distribution similar to many earlier sensor variables.

Additional observations include:

- The histogram appears relatively symmetric with mild skewness.
 - The mean and median are very close, indicating balanced behavior.
 - Several outliers are visible on both lower and upper tails.
 - Compared to earlier variables, V40 appears slightly more stable with a narrower spread.
-

Importance of Variable Distribution Analysis

Analyzing variable distributions helps identify:

- skewness,
- variability,
- outliers,
- and potential data quality concerns.

These insights help determine:

- whether scaling is required,
 - whether transformations may improve model performance,
 - and which variables may contribute strongly to predictive accuracy.
-

Business Impact of Sensor Distribution Analysis

Stable sensor behavior improves the model's ability to distinguish between:

- healthy turbine operation,
- and abnormal operating conditions.

Variables showing unusual or extreme behavior may act as early warning indicators for turbine failure and predictive maintenance systems.

Target Variable Analysis

Observations

The target variable shows a highly imbalanced distribution.

- Most observations belong to class **0 (No Failure)**.
- Only a small proportion belong to class **1 (Failure)**.

This confirms:

- approximately **94.5% No Failure**
- approximately **5.5% Failure**

The boxplot also highlights the strong imbalance between the two classes.

Why Target Distribution Is Important

1. Class Imbalance Impacts Model Training

Machine learning models may become biased toward the majority class.

For example:

- A model predicting only "No Failure" could still achieve high accuracy.
 - However, it would fail to identify actual turbine failures.
-

2. Accuracy Alone Becomes Misleading

Because failures are rare:

- Accuracy alone is not a reliable evaluation metric.

More meaningful metrics include:

- Recall,
 - Precision,
 - F1-score,
 - ROC-AUC,
 - and confusion matrix analysis.
-

3. Failure Detection Is the Main Business Objective

The business problem prioritizes early failure detection.

Missing a true failure (False Negative) may result in:

- turbine breakdown,
- expensive generator replacement,
- operational downtime,
- and increased maintenance expenses.

Therefore, the final model should prioritize:

- maximizing Recall,
 - minimizing False Negatives,
 - and improving failure detection performance.
-

Overall EDA Findings

Combining the analysis from all variables:

Key Findings

- Most variables follow approximately normal or near-normal distributions.
- Outliers are consistently present across nearly all variables.

- Several variables exhibit moderate skewness and heavy tails.
- Variability differs substantially across sensor variables.
- No strong multimodal distributions are observed.
- The dataset appears numerically stable overall after preprocessing.

Variables such as:

- V16,
- V21,
- V24,
- V31,
- V32,
- and V36

show especially high variability and extreme values, making them potentially important predictors of turbine failure.

Overall Importance of EDA

Exploratory Data Analysis is essential because it helps:

- understand data quality,
- identify missing values,
- detect abnormal patterns,
- analyze feature distributions,
- recognize class imbalance,
- and guide preprocessing and modeling decisions.

EDA ensures:

- appropriate preprocessing strategies,
 - reliable machine learning inputs,
 - and alignment between technical modeling and business objectives.
-

Overall Business Interpretation

The EDA findings suggest that:

- Sensor readings contain meaningful operational patterns.
- Abnormal sensor behavior may act as early warning signals for failure.

- Certain variables may provide strong predictive signals for generator breakdown.
- The dataset is suitable for predictive maintenance modeling.

The strong class imbalance indicates that:

- failure prediction is challenging,
 - but critically important from a business perspective.
-

Overall Business Impact

A successful predictive maintenance model can help ReneWind:

- reduce generator replacement costs,
- minimize turbine downtime,
- improve operational efficiency,
- optimize maintenance scheduling,
- reduce unnecessary inspections,
- and extend generator lifespan.

The largest business risk comes from False Negatives, where actual failures are missed.

Therefore, the final predictive model should focus heavily on:

- failure detection capability,
- high Recall,
- and minimizing missed failures.

The long-term business goal is to transition from:

- reactive maintenance

to:

- proactive predictive maintenance.
-

Recommended Next Steps

Based on the EDA findings, the next stages should include:

Data Preprocessing

1. Missing value treatment
2. Feature scaling and normalization

Handling Class Imbalance

3. Techniques such as:
 - SMOTE,
 - class weighting,
 - threshold tuning

Model Development

4. Build and compare:
 - Logistic Regression
 - Decision Trees
 - Random Forest
 - Gradient Boosting
 - XGBoost
 - Neural Networks

Model Evaluation

5. Evaluate models using:
 - Recall
 - Precision
 - F1-score
 - ROC-AUC
 - Confusion Matrix

Model Optimization

6. Perform hyperparameter tuning to improve failure detection performance

In []:

In []:

In []:


In []:


In []:


In []:


Variables V1 to V29


```
In [23]: for feature in df.columns:  
         histogram_boxplot(df, feature, figsize=(12, 7), kde=False, bins=No
```



























































In []:

Checking the distrubution of Target variable

```
In [24]: # For train data
df["Target"].value_counts(1)
```

```
Out[24]: Target  
0      0.9445  
1      0.0555  
Name: proportion, dtype: float64
```

```
In [25]: # For test data  
df_test["Target"].value_counts(1) # Complete the code to display the p
```

```
Out[25]: Target  
0      0.9436  
1      0.0564  
Name: proportion, dtype: float64
```

Distribution of the Target Variable

Observations

The distribution of the target variable in both the training and testing datasets is highly imbalanced.

Training Dataset

- Class 0 (No Failure): 94.45%
- Class 1 (Failure): 5.55%

Testing Dataset

- Class 0 (No Failure): 94.36%
- Class 1 (Failure): 5.64%

The proportions in the training and testing datasets are very similar, indicating that both datasets follow a consistent distribution.

This means:

- turbine failures are relatively rare events,
- while the majority of turbines operate under normal conditions.

Why This Is Important

Analyzing the distribution of the target variable is extremely important because:

1. The Dataset Is Highly Imbalanced

The model will be trained on a dataset where the majority class ("No Failure") dominates the observations.

Without proper handling:

- the model may become biased toward predicting the majority class,
- resulting in poor failure detection performance.

For example:

- a model predicting only "No Failure" could still achieve over 94% accuracy,
 - but it would completely fail to identify actual turbine failures.
-

2. Accuracy Alone Becomes Misleading

Because failure events are rare:

- accuracy is not an appropriate standalone evaluation metric.

More meaningful evaluation metrics include:

- Recall,
- Precision,
- F1-score,
- ROC-AUC,
- and confusion matrix analysis.

These metrics provide a better understanding of how effectively the model identifies failures.

3. Class Imbalance Affects Model Optimization

The imbalance suggests that additional techniques may be needed during model development, such as:

- class weighting,
- SMOTE (Synthetic Minority Oversampling),
- threshold tuning,
- or specialized evaluation strategies.

These methods help improve the model's ability to detect rare failure events.

Business Impact

The target variable distribution has major business implications because turbine failures are rare but extremely costly.

Cost of False Negatives

A False Negative occurs when:

- the model predicts "No Failure,"
- but the turbine actually fails.

This can lead to:

- generator breakdown,
- expensive replacement costs,
- operational downtime,
- loss of energy production,
- and increased maintenance expenses.

False Negatives are the most critical business risk in this project.

Cost of False Positives

A False Positive occurs when:

- the model predicts failure,
- but the turbine is actually healthy.

This leads to:

- inspection costs,
- unnecessary maintenance checks,
- and minor operational inefficiencies.

However, False Positives are significantly less expensive than missed failures.

Business Objective

From a business perspective, the primary objective is:

- to detect as many true failures as possible,
- minimize False Negatives,
- and reduce costly turbine breakdowns.

Therefore, the final predictive maintenance model should prioritize:

- high Recall,
- strong failure detection capability,
- and early identification of abnormal turbine behavior.

The ultimate goal is to help ReneWind transition from:

- reactive maintenance

to:

- proactive predictive maintenance.

In []:

In []:


In []:

Bivariate Analysis

Correlation Check

```
In [26]: cols_list = df.select_dtypes(include=np.number).columns.tolist()
cols_list.remove("Target")

plt.figure(figsize=(20, 20))
sns.heatmap(
    df[cols_list].corr(), annot=True, vmin=-1, vmax=1, fmt=".2f", cmap
)
plt.show()
```



In []:

Bivariate Analysis – Correlation Heatmap

Observations from the Correlation Matrix

The heatmap displays the pairwise correlation coefficients between all numerical variables (V1–V40).

Correlation values range from:

- **+1** → strong positive relationship

- **-1** → strong negative relationship
- **0** → little or no linear relationship

Key Observations

- Most variables show:
 - weak to moderate correlations,
 - indicating that many sensor variables capture different operational characteristics of the turbines.
- Several variable pairs exhibit:
 - strong positive correlations,
 - where variables increase together.

Examples include:

- V1 with V25 and V27
- V11 with V29
- V24 with V32
- V30 with V33
- Some variable pairs show strong negative correlations, meaning:
 - as one variable increases, the other decreases.

Examples include:

- V2 with V14
- V18 with V35
- V25 with V31
- No extremely perfect correlations are observed between most variables, which is beneficial because:
 - the dataset does not appear heavily dominated by duplicate information.
- Certain clusters of variables display similar correlation patterns, suggesting that:
 - some sensors may be monitoring related turbine behaviors or interconnected mechanical systems.

Why Correlation Analysis Is Important in EDA

Correlation analysis is one of the most important parts of Exploratory Data Analysis because it helps identify relationships between variables.

1. Detecting Multicollinearity

Highly correlated variables may contain redundant information.

This is important because:

- multicollinearity can negatively affect certain machine learning models,
- especially linear models such as Logistic Regression.

Strong multicollinearity may:

- reduce model interpretability,
 - increase model instability,
 - and inflate coefficient importance.
-

2. Understanding Relationships Between Sensors

Correlation analysis helps identify:

- which turbine sensors move together,
- which variables behave independently,
- and which measurements may represent similar operational conditions.

This provides insight into:

- turbine mechanics,
 - operational dependencies,
 - and system interactions.
-

3. Supporting Feature Selection

Highly correlated features may:

- provide overlapping information,

- or contribute limited additional predictive value.

This analysis helps determine whether:

- feature reduction,
- dimensionality reduction,
- or feature engineering

may improve model performance.

4. Improving Machine Learning Performance

Understanding variable relationships helps:

- reduce noise,
- improve model stability,
- and optimize preprocessing decisions.

For example:

- strongly correlated features may benefit tree-based models,
 - while redundant variables may hurt simpler linear models.
-

Business Interpretation

From a business and operational perspective:

- Correlated variables may represent connected turbine subsystems such as:
 - temperature,
 - pressure,
 - vibration,
 - rotational speed,
 - or generator load.
- Strong positive correlations may indicate:
 - components that react together during normal operation,
 - or shared operational behavior across turbine systems.
- Strong negative correlations may indicate:
 - balancing mechanisms,

- inverse operational relationships,
 - or compensating system responses.
 - Changes in correlated sensor groups may help identify:
 - abnormal turbine conditions,
 - mechanical degradation,
 - or early failure patterns.
-

Business Impact

Understanding correlations between turbine sensors can significantly improve predictive maintenance strategies.

Operational Benefits

A better understanding of sensor relationships can help:

- detect abnormal operating conditions earlier,
 - improve failure prediction accuracy,
 - reduce unnecessary maintenance inspections,
 - and optimize monitoring systems.
-

Maintenance Optimization

Correlation analysis may help engineers:

- identify which sensors provide the most useful operational information,
 - reduce redundant monitoring,
 - and focus attention on high-risk turbine components.
-

Cost Reduction

Improved predictive modeling can help reduce:

- generator replacement costs,
- unplanned downtime,
- operational interruptions,
- and maintenance expenses.

Overall Conclusion

The correlation heatmap suggests that:

- the dataset contains meaningful relationships between turbine sensor variables,
- moderate-to-strong correlations exist between several features,
- and the data appears suitable for machine learning modeling.

The absence of excessive perfect correlations is beneficial because:

- the model is less likely to suffer from severe redundancy issues,
- while still retaining useful inter-variable relationships that may improve turbine failure prediction.

In []:

In []:

In []:

In []:

Data Preprocessing

Data Preparation for Modeling

```
In [31]: # Dividing train data into X and y
X = data.drop(columns = ["Target"] , axis=1) # Complete the code to re
y = data["Target"] # Complete the code to select the column named 'Tar
```

Since we already have a separate test set, we don't need to divide data into train, valiation and test

```
In [32]: # Splitting data into training and validation set:

X_train, X_val, y_train, y_val = train_test_split(
    X, y, test_size=0.2, random_state=1, stratify=y # Complete the cod
)
```

```
In [33]: # Checking the number of rows and columns in the X_train data
```

```
X_train.shape
```

```
Out[33]: (16000, 40)
```

```
In [34]: # Checking the number of rows and columns in the X_val data  
X_val.shape
```

```
Out[34]: (4000, 40)
```

```
In [35]: # Dividing test data into X_test and y_test  
X_test = data_test.drop(columns = ['Target'] , axis= 1) # Complete the  
y_test = data_test["Target"] # Complete the code to select the target
```

```
In [36]: # Checking the number of rows and columns in the X_test data  
X_test.shape
```

```
Out[36]: (5000, 40)
```

Missing Value Imputation

Data Preparation for Modeling

Observations

The dataset was successfully divided into training, validation, and testing datasets for machine learning model development.

Training Dataset

- Shape: **(16000, 40)**
- Contains:
 - 16,000 observations
 - 40 predictor variables

Validation Dataset

- Shape: **(4000, 40)**
- Contains:
 - 4,000 observations
 - 40 predictor variables

Testing Dataset

- Shape: **(5000, 40)**
- Contains:
 - 5,000 observations
 - 40 predictor variables

The target variable (Target) was successfully separated from the predictor variables before model training.

Why This Is Important

Preparing the data correctly before modeling is one of the most critical steps in machine learning.

1. Separation of Features and Target Variable

The predictor variables (X) and target variable (y) were separated to:

- allow the model to learn relationships between sensor readings and turbine failures,
- avoid including the target variable as an input feature,
- and ensure proper supervised learning.

This is necessary because the target variable represents the output the model is trying to predict.

2. Splitting the Dataset

The training dataset was split into:

- training data (80%)
- validation data (20%)

This is important because:

Training Data

The model learns patterns from the training dataset.

Validation Data

The validation dataset is used to:

- evaluate model performance during training,
- tune hyperparameters,
- compare models,
- and prevent overfitting.

The separate testing dataset is reserved for final unbiased model evaluation.

3. Stratified Sampling

The parameter:

`stratify=y`

ensures that the **class** distribution remains consistent across:

- training,
- validation,
- **and** testing datasets.

This **is** especially important because the target variable **is** highly imbalanced.

Without stratification:

- one dataset might contain too few failure cases,
- leading to unreliable model evaluation **and** poor generalization.

Maintaining similar **class** distributions across datasets helps ensure that:

- the model learns realistic patterns,
- evaluation metrics remain reliable,
- **and** model performance reflects real-world operating conditions.

4. Preventing Data Leakage

The testing dataset was kept completely separate **from** the training **and** validation process.

This **is** important because:

- the model should never learn **from** unseen test data during training,

- **and** the testing dataset should only be used **for** final evaluation.

Preventing data leakage ensures:

- unbiased model evaluation,
- realistic performance measurement,
- **and** better generalization to future turbine data.

Business Impact

Proper data preparation directly impacts the reliability **and** effectiveness of the predictive maintenance system.

Improved Failure Detection

Well-structured training **and** validation datasets help the model:

- learn meaningful turbine behavior patterns,
- identify abnormal operating conditions,
- **and** improve failure prediction accuracy.

Reliable Model Evaluation

Using separate validation **and** testing datasets ensures:

- fair model evaluation,
- reliable performance measurement,
- **and** stronger confidence **in** deployment readiness.

This **is** critical because predictive maintenance systems must perform consistently on unseen real-world turbine data.

Reduced Operational Risk

Accurate failure prediction can help ReneWind:

- reduce unexpected turbine breakdowns,
- minimize operational downtime,
- lower replacement costs,
- **and** improve maintenance scheduling efficiency.

Support for Predictive Maintenance Strategy

Properly prepared data increases the likelihood of building a

robust predictive maintenance system capable of:

- detecting failures early,
- reducing costly emergency repairs,
- extending turbine lifespan,
- **and** improving operational reliability.

The ultimate business goal **is** to transition from:

- reactive maintenance

to:

- proactive predictive maintenance.

- There were few missing values in V1 and V2, we will impute them using the median.
- And to avoid data leakage we will impute missing values after splitting train data into train and validation sets.

```
In [37]: imputer = SimpleImputer(strategy="median")
```

```
In [38]: # Fit and transform the train data
X_train = pd.DataFrame(imputer.fit_transform(X_train), columns=X_train.columns)

# Transform the validation data
X_val = pd.DataFrame(imputer.transform(X_val), columns=X_train.columns)

# Transform the test data
X_test = pd.DataFrame(imputer.transform(X_test), columns=X_train.columns)
```

```
In [39]: # Checking that no column has missing values in train or test sets
print(X_train.isna().sum())
print("-" * 30)
print(X_val.isna().sum())
print("-" * 30)
print(X_test.isna().sum())
```

```
V1      0
V2      0
V3      0
V4      0
V5      0
V6      0
V7      0
V8      0
V9      0
V10     0
V11     0
V12     0
V13     0
```

V14	0
V15	0
V16	0
V17	0
V18	0
V19	0
V20	0
V21	0
V22	0
V23	0
V24	0
V25	0
V26	0
V27	0
V28	0
V29	0
V30	0
V31	0
V32	0
V33	0
V34	0
V35	0
V36	0
V37	0
V38	0
V39	0
V40	0

dtype: int64

V1	0
V2	0
V3	0
V4	0
V5	0
V6	0
V7	0
V8	0
V9	0
V10	0
V11	0
V12	0
V13	0
V14	0
V15	0
V16	0
V17	0
V18	0
V19	0
V20	0
V21	0
V22	0
V23	0

```
V24    0
V25    0
V26    0
V27    0
V28    0
V29    0
V30    0
V31    0
V32    0
V33    0
V34    0
V35    0
V36    0
V37    0
V38    0
V39    0
V40    0
dtype: int64
```

```
V1      0
V2      0
V3      0
V4      0
V5      0
V6      0
V7      0
V8      0
V9      0
V10     0
V11     0
V12     0
V13     0
V14     0
V15     0
V16     0
V17     0
V18     0
V19     0
V20     0
V21     0
V22     0
V23     0
V24     0
V25     0
V26     0
V27     0
V28     0
V29     0
V30     0
V31     0
V32     0
V33     0
```

```
V34    0
V35    0
V36    0
V37    0
V38    0
V39    0
V40    0
dtype: int64
```

```
In [40]: y_train = y_train.to_numpy()
         y_val = y_val.to_numpy()
         y_test = y_test.to_numpy()
```

Missing Value Treatment

What This Output Says

The missing values in the dataset were handled using **median imputation**.

Earlier analysis showed that only **V1** and **V2** had missing values. After applying the imputer, the output shows that every column in the training, validation, and test datasets now has **0 missing values**.

This confirms that missing value treatment was completed successfully.

Why This Is Important

Missing value treatment is important because:

- most machine learning models cannot process missing values directly,
- missing values can cause errors during model training,
- incomplete data can reduce model performance,
- and proper imputation helps maintain consistency across train, validation, and test datasets.

Median imputation was used because it is more robust to outliers than mean imputation. This is useful here because the EDA showed that several sensor variables contain outliers.

Preventing Data Leakage

The imputer was fitted only on the training data using:

```
imputer.fit_transform(X_train)
```

Model Building

Model Evaluation Criterion

- Write down the metric of choice with rationale here

We are now done with pre-processing and evaluation criterion, so let's start building the model.

Utility Functions

```
In [42]: def plot(history, name):
        """
        Function to plot loss/accuracy

        history: an object which stores the metrics and losses.
        name: can be one of Loss or Accuracy
        """
        fig, ax = plt.subplots() #Creating a subplot with figure and axes.
        plt.plot(history.history[name]) #Plotting the train accuracy or tr
        plt.plot(history.history['val_'+name]) #Plotting the validation ac

        plt.title('Model ' + name.capitalize()) #Defining the title of the
        plt.ylabel(name.capitalize()) #Capitalizing the first letter.
        plt.xlabel('Epoch') #Defining the label for the x-axis.
        fig.legend(['Train', 'Validation'], loc="outside right upper") #De
```

```
In [43]: # defining a function to compute different metrics to check performanc
def model_performance_classification(
    model, predictors, target, threshold=0.5
):
    """
    Function to compute different metrics to check classification mode

    model: classifier
    predictors: independent variables
    target: dependent variable
    threshold: threshold for classifying the observation as class 1
    """
```

```

# checking which probabilities are greater than threshold
pred = model.predict(predictors) > threshold
# pred_temp = model.predict(predictors) > threshold
# # rounding off the above values to get classes
# pred = np.round(pred_temp)

acc = accuracy_score(target, pred) # to compute Accuracy
recall = recall_score(target, pred, average='macro') # to compute
precision = precision_score(target, pred, average='macro') # to c
f1 = f1_score(target, pred, average='macro') # to compute F1-scor

# creating a dataframe of metrics
df_perf = pd.DataFrame(
    {"Accuracy": acc, "Recall": recall, "Precision": precision, "F
)

return df_perf

```

Initial Model Building (Model 0)

- Let's start with a neural network consisting of
 - just one hidden layer of 7 neurons respectively
 - activation function of ReLU.
 - SGD as the optimizer

```

In [44]: # defining the batch size and # epochs upfront as we'll be using the s
epochs = 56 # Complete the code to enter the number of epochs to be
batch_size = 256 # Complete the code to enter the batch size to be

```

```

In [45]: # clears the current Keras session, resetting all layers and models pr
tf.keras.backend.clear_session()

```

```

In [46]: #Initializing the neural network
model_0 = Sequential()
model_0.add(Dense( 64 ,activation="relu",input_dim=X_train.shape[1]))
model_0.add(Dense( 1 ,activation="sigmoid")) # Complete the code to de

```

```

In [47]: model_0.summary()

```

Model: "sequential"

Layer (type)	Output Shape	
dense (Dense)	(None, 64)	
dense_1 (Dense)	(None, 1)	

Total params: 2,689 (10.50 KB)

Trainable params: 2,689 (10.50 KB)

Non-trainable params: 0 (0.00 B)

In []:

In []:

In []:

```
In [48]: optimizer = tf.keras.optimizers.SGD() # defining SGD as the optimizer
# model_0.compile(loss='binary_crossentropy', optimizer=optimizer, met
# model_0.compile(loss='binary_crossentropy', optimizer=optimizer, met
# model_0.compile(loss='binary_crossentropy', optimizer=optimizer, met
# model_0.compile(loss='binary_crossentropy', optimizer=optimizer, met
```

```
In [50]: optimizer = tf.keras.optimizers.SGD()
```

```
model_0.compile(
    loss="binary_crossentropy",
    optimizer=optimizer,
    metrics=["Recall"]
)
```

```
In [51]: start = time.time()

history = model_0.fit(
    X_train,
    y_train,
    validation_data=(X_val, y_val),
    batch_size=batch_size,
    epochs=epochs
)

end = time.time()

print("Time taken in seconds ", end - start)
```

Epoch 1/56

63/63 ————— 0s 2ms/step - Recall: 0.0890 - loss: 0.2253
- val_Recall: 0.1892 - val_loss: 0.1616

Epoch 2/56

63/63 ————— 0s 509us/step - Recall: 0.2815 - loss: 0.1524
- val_Recall: 0.3829 - val_loss: 0.1338

Epoch 3/56

63/63 ————— 0s 565us/step - Recall: 0.4178 - loss: 0.1312
- val_Recall: 0.5045 - val_loss: 0.1215

Epoch 4/56

63/63 ————— 0s 577us/step - Recall: 0.4797 - loss: 0.1193
- val_Recall: 0.5270 - val_loss: 0.1126

Epoch 5/56
63/63  **0s** 584us/step - Recall: 0.5214 - loss: 0.111
1 - val_Recall: 0.5811 - val_loss: 0.1066
Epoch 6/56
63/63  **0s** 572us/step - Recall: 0.5586 - loss: 0.105
1 - val_Recall: 0.6216 - val_loss: 0.1022
Epoch 7/56
63/63  **0s** 560us/step - Recall: 0.5777 - loss: 0.100
3 - val_Recall: 0.6171 - val_loss: 0.0983
Epoch 8/56
63/63  **0s** 575us/step - Recall: 0.6002 - loss: 0.096
4 - val_Recall: 0.6351 - val_loss: 0.0953
Epoch 9/56
63/63  **0s** 598us/step - Recall: 0.6160 - loss: 0.093
2 - val_Recall: 0.6532 - val_loss: 0.0929
Epoch 10/56
63/63  **0s** 592us/step - Recall: 0.6351 - loss: 0.090
5 - val_Recall: 0.6757 - val_loss: 0.0908
Epoch 11/56
63/63  **0s** 527us/step - Recall: 0.6453 - loss: 0.088
1 - val_Recall: 0.6982 - val_loss: 0.0891
Epoch 12/56
63/63  **0s** 523us/step - Recall: 0.6622 - loss: 0.086
1 - val_Recall: 0.6982 - val_loss: 0.0875
Epoch 13/56
63/63  **0s** 530us/step - Recall: 0.6723 - loss: 0.084
3 - val_Recall: 0.7027 - val_loss: 0.0861
Epoch 14/56
63/63  **0s** 520us/step - Recall: 0.6869 - loss: 0.082
8 - val_Recall: 0.7027 - val_loss: 0.0848
Epoch 15/56
63/63  **0s** 536us/step - Recall: 0.6982 - loss: 0.081
4 - val_Recall: 0.7027 - val_loss: 0.0836
Epoch 16/56
63/63  **0s** 541us/step - Recall: 0.7005 - loss: 0.080
1 - val_Recall: 0.7162 - val_loss: 0.0827
Epoch 17/56
63/63  **0s** 547us/step - Recall: 0.7106 - loss: 0.078
9 - val_Recall: 0.7252 - val_loss: 0.0818
Epoch 18/56
63/63  **0s** 518us/step - Recall: 0.7173 - loss: 0.077
8 - val_Recall: 0.7387 - val_loss: 0.0810
Epoch 19/56
63/63  **0s** 535us/step - Recall: 0.7241 - loss: 0.076
9 - val_Recall: 0.7387 - val_loss: 0.0803
Epoch 20/56
63/63  **0s** 515us/step - Recall: 0.7275 - loss: 0.075
9 - val_Recall: 0.7432 - val_loss: 0.0796
Epoch 21/56
63/63  **0s** 548us/step - Recall: 0.7354 - loss: 0.075
1 - val_Recall: 0.7387 - val_loss: 0.0789
Epoch 22/56

63/63 ————— **0s** 510us/step - Recall: 0.7342 - loss: 0.074
3 - val_Recall: 0.7387 - val_loss: 0.0783
Epoch 23/56

63/63 ————— **0s** 559us/step - Recall: 0.7399 - loss: 0.073
6 - val_Recall: 0.7568 - val_loss: 0.0778
Epoch 24/56

63/63 ————— **0s** 502us/step - Recall: 0.7477 - loss: 0.072
9 - val_Recall: 0.7568 - val_loss: 0.0772
Epoch 25/56

63/63 ————— **0s** 502us/step - Recall: 0.7511 - loss: 0.072
3 - val_Recall: 0.7613 - val_loss: 0.0768
Epoch 26/56

63/63 ————— **0s** 513us/step - Recall: 0.7568 - loss: 0.071
7 - val_Recall: 0.7658 - val_loss: 0.0763
Epoch 27/56

63/63 ————— **0s** 533us/step - Recall: 0.7601 - loss: 0.071
1 - val_Recall: 0.7703 - val_loss: 0.0759
Epoch 28/56

63/63 ————— **0s** 513us/step - Recall: 0.7646 - loss: 0.070
6 - val_Recall: 0.7748 - val_loss: 0.0755
Epoch 29/56

63/63 ————— **0s** 520us/step - Recall: 0.7680 - loss: 0.070
0 - val_Recall: 0.7703 - val_loss: 0.0751
Epoch 30/56

63/63 ————— **0s** 547us/step - Recall: 0.7703 - loss: 0.069
6 - val_Recall: 0.7793 - val_loss: 0.0748
Epoch 31/56

63/63 ————— **0s** 506us/step - Recall: 0.7725 - loss: 0.069
1 - val_Recall: 0.7748 - val_loss: 0.0743
Epoch 32/56

63/63 ————— **0s** 524us/step - Recall: 0.7714 - loss: 0.068
6 - val_Recall: 0.7748 - val_loss: 0.0740
Epoch 33/56

63/63 ————— **0s** 584us/step - Recall: 0.7714 - loss: 0.068
1 - val_Recall: 0.7838 - val_loss: 0.0740
Epoch 34/56

63/63 ————— **0s** 488us/step - Recall: 0.7804 - loss: 0.067
7 - val_Recall: 0.7748 - val_loss: 0.0733
Epoch 35/56

63/63 ————— **0s** 491us/step - Recall: 0.7770 - loss: 0.067
3 - val_Recall: 0.7838 - val_loss: 0.0731
Epoch 36/56

63/63 ————— **0s** 525us/step - Recall: 0.7815 - loss: 0.066
9 - val_Recall: 0.7793 - val_loss: 0.0728
Epoch 37/56

63/63 ————— **0s** 491us/step - Recall: 0.7793 - loss: 0.066
5 - val_Recall: 0.7883 - val_loss: 0.0726
Epoch 38/56

63/63 ————— **0s** 522us/step - Recall: 0.7838 - loss: 0.066
2 - val_Recall: 0.7883 - val_loss: 0.0722
Epoch 39/56

63/63 ————— **0s** 559us/step - Recall: 0.7815 - loss: 0.065

```
8 - val_Recall: 0.7883 - val_loss: 0.0721
Epoch 40/56
63/63 ————— 0s 539us/step - Recall: 0.7849 - loss: 0.065
5 - val_Recall: 0.7928 - val_loss: 0.0718
Epoch 41/56
63/63 ————— 0s 552us/step - Recall: 0.7883 - loss: 0.065
1 - val_Recall: 0.7928 - val_loss: 0.0715
Epoch 42/56
63/63 ————— 0s 565us/step - Recall: 0.7883 - loss: 0.064
8 - val_Recall: 0.7928 - val_loss: 0.0713
Epoch 43/56
63/63 ————— 0s 561us/step - Recall: 0.7905 - loss: 0.064
5 - val_Recall: 0.7928 - val_loss: 0.0711
Epoch 44/56
63/63 ————— 0s 616us/step - Recall: 0.7928 - loss: 0.064
2 - val_Recall: 0.7928 - val_loss: 0.0708
Epoch 45/56
63/63 ————— 0s 529us/step - Recall: 0.7917 - loss: 0.063
9 - val_Recall: 0.7928 - val_loss: 0.0706
Epoch 46/56
63/63 ————— 0s 523us/step - Recall: 0.7939 - loss: 0.063
6 - val_Recall: 0.7928 - val_loss: 0.0705
Epoch 47/56
63/63 ————— 0s 587us/step - Recall: 0.7973 - loss: 0.063
3 - val_Recall: 0.7928 - val_loss: 0.0703
Epoch 48/56
63/63 ————— 0s 545us/step - Recall: 0.7962 - loss: 0.063
0 - val_Recall: 0.7928 - val_loss: 0.0701
Epoch 49/56
63/63 ————— 0s 658us/step - Recall: 0.7973 - loss: 0.062
7 - val_Recall: 0.7928 - val_loss: 0.0700
Epoch 50/56
63/63 ————— 0s 612us/step - Recall: 0.7984 - loss: 0.062
5 - val_Recall: 0.7928 - val_loss: 0.0697
Epoch 51/56
63/63 ————— 0s 593us/step - Recall: 0.8007 - loss: 0.062
2 - val_Recall: 0.7928 - val_loss: 0.0695
Epoch 52/56
63/63 ————— 0s 620us/step - Recall: 0.7995 - loss: 0.061
9 - val_Recall: 0.7928 - val_loss: 0.0694
Epoch 53/56
63/63 ————— 0s 593us/step - Recall: 0.7973 - loss: 0.061
7 - val_Recall: 0.7928 - val_loss: 0.0692
Epoch 54/56
63/63 ————— 0s 544us/step - Recall: 0.8018 - loss: 0.061
4 - val_Recall: 0.7928 - val_loss: 0.0690
Epoch 55/56
63/63 ————— 0s 511us/step - Recall: 0.8018 - loss: 0.061
2 - val_Recall: 0.7928 - val_loss: 0.0689
Epoch 56/56
63/63 ————— 0s 562us/step - Recall: 0.8029 - loss: 0.061
0 - val_Recall: 0.7973 - val_loss: 0.0687
```

Time taken in seconds 2.64277720451355

In []:

```
In [52]: start = time.time()
history = model_0.fit(X_train, y_train, validation_data=(X_val,y_val)
end=time.time()
```

Epoch 1/56

63/63  **0s** 1ms/step - Recall: 0.8029 - loss: 0.0607
- val_Recall: 0.7973 - val_loss: 0.0686

Epoch 2/56

63/63  **0s** 558us/step - Recall: 0.8063 - loss: 0.0605
- val_Recall: 0.7973 - val_loss: 0.0684

Epoch 3/56

63/63  **0s** 521us/step - Recall: 0.8063 - loss: 0.0603
- val_Recall: 0.7973 - val_loss: 0.0682

Epoch 4/56

63/63  **0s** 552us/step - Recall: 0.8041 - loss: 0.0601
- val_Recall: 0.7973 - val_loss: 0.0682

Epoch 5/56

63/63  **0s** 534us/step - Recall: 0.8074 - loss: 0.0598
- val_Recall: 0.7973 - val_loss: 0.0680

Epoch 6/56

63/63  **0s** 845us/step - Recall: 0.8074 - loss: 0.0597
- val_Recall: 0.7973 - val_loss: 0.0679

Epoch 7/56

63/63  **0s** 508us/step - Recall: 0.8063 - loss: 0.0595
- val_Recall: 0.7973 - val_loss: 0.0678

Epoch 8/56

63/63  **0s** 557us/step - Recall: 0.8086 - loss: 0.0593
- val_Recall: 0.7973 - val_loss: 0.0676

Epoch 9/56

63/63  **0s** 539us/step - Recall: 0.8108 - loss: 0.0591
- val_Recall: 0.7973 - val_loss: 0.0675

Epoch 10/56

63/63  **0s** 511us/step - Recall: 0.8086 - loss: 0.0589
- val_Recall: 0.7973 - val_loss: 0.0674

Epoch 11/56

63/63  **0s** 553us/step - Recall: 0.8119 - loss: 0.0587
- val_Recall: 0.7973 - val_loss: 0.0672

Epoch 12/56

63/63  **0s** 509us/step - Recall: 0.8131 - loss: 0.0585
- val_Recall: 0.7973 - val_loss: 0.0671

Epoch 13/56

63/63  **0s** 552us/step - Recall: 0.8119 - loss: 0.0583
- val_Recall: 0.7973 - val_loss: 0.0670

Epoch 14/56

63/63  **0s** 546us/step - Recall: 0.8119 - loss: 0.0582
- val_Recall: 0.7973 - val_loss: 0.0669

Epoch 15/56

63/63  **0s** 539us/step - Recall: 0.8142 - loss: 0.0580
- val_Recall: 0.7973 - val_loss: 0.0668

Epoch 16/56
63/63  **0s** 523us/step - Recall: 0.8142 - loss: 0.0578 - val_Recall: 0.7973 - val_loss: 0.0667

Epoch 17/56
63/63  **0s** 582us/step - Recall: 0.8142 - loss: 0.0577 - val_Recall: 0.7973 - val_loss: 0.0666

Epoch 18/56
63/63  **0s** 534us/step - Recall: 0.8164 - loss: 0.0575 - val_Recall: 0.8018 - val_loss: 0.0665

Epoch 19/56
63/63  **0s** 535us/step - Recall: 0.8187 - loss: 0.0573 - val_Recall: 0.8018 - val_loss: 0.0664

Epoch 20/56
63/63  **0s** 542us/step - Recall: 0.8164 - loss: 0.0572 - val_Recall: 0.8018 - val_loss: 0.0663

Epoch 21/56
63/63  **0s** 521us/step - Recall: 0.8198 - loss: 0.0570 - val_Recall: 0.8018 - val_loss: 0.0662

Epoch 22/56
63/63  **0s** 515us/step - Recall: 0.8164 - loss: 0.0569 - val_Recall: 0.8063 - val_loss: 0.0661

Epoch 23/56
63/63  **0s** 569us/step - Recall: 0.8198 - loss: 0.0567 - val_Recall: 0.8108 - val_loss: 0.0661

Epoch 24/56
63/63  **0s** 540us/step - Recall: 0.8221 - loss: 0.0566 - val_Recall: 0.8108 - val_loss: 0.0660

Epoch 25/56
63/63  **0s** 538us/step - Recall: 0.8221 - loss: 0.0564 - val_Recall: 0.8108 - val_loss: 0.0658

Epoch 26/56
63/63  **0s** 526us/step - Recall: 0.8176 - loss: 0.0563 - val_Recall: 0.8108 - val_loss: 0.0658

Epoch 27/56
63/63  **0s** 512us/step - Recall: 0.8232 - loss: 0.0561 - val_Recall: 0.8108 - val_loss: 0.0657

Epoch 28/56
63/63  **0s** 570us/step - Recall: 0.8221 - loss: 0.0560 - val_Recall: 0.8108 - val_loss: 0.0656

Epoch 29/56
63/63  **0s** 550us/step - Recall: 0.8187 - loss: 0.0559 - val_Recall: 0.8108 - val_loss: 0.0656

Epoch 30/56
63/63  **0s** 559us/step - Recall: 0.8221 - loss: 0.0557 - val_Recall: 0.8108 - val_loss: 0.0654

Epoch 31/56
63/63  **0s** 579us/step - Recall: 0.8221 - loss: 0.0556 - val_Recall: 0.8153 - val_loss: 0.0654

Epoch 32/56
63/63  **0s** 511us/step - Recall: 0.8221 - loss: 0.0555 - val_Recall: 0.8153 - val_loss: 0.0653

Epoch 33/56

63/63 ————— **0s** 553us/step - Recall: 0.8232 - loss: 0.055
3 - val_Recall: 0.8153 - val_loss: 0.0652
Epoch 34/56

63/63 ————— **0s** 545us/step - Recall: 0.8221 - loss: 0.055
2 - val_Recall: 0.8153 - val_loss: 0.0652
Epoch 35/56

63/63 ————— **0s** 542us/step - Recall: 0.8232 - loss: 0.055
1 - val_Recall: 0.8153 - val_loss: 0.0651
Epoch 36/56

63/63 ————— **0s** 548us/step - Recall: 0.8221 - loss: 0.055
0 - val_Recall: 0.8153 - val_loss: 0.0650
Epoch 37/56

63/63 ————— **0s** 522us/step - Recall: 0.8243 - loss: 0.054
8 - val_Recall: 0.8153 - val_loss: 0.0650
Epoch 38/56

63/63 ————— **0s** 567us/step - Recall: 0.8243 - loss: 0.054
7 - val_Recall: 0.8153 - val_loss: 0.0649
Epoch 39/56

63/63 ————— **0s** 588us/step - Recall: 0.8232 - loss: 0.054
6 - val_Recall: 0.8153 - val_loss: 0.0649
Epoch 40/56

63/63 ————— **0s** 519us/step - Recall: 0.8243 - loss: 0.054
5 - val_Recall: 0.8153 - val_loss: 0.0648
Epoch 41/56

63/63 ————— **0s** 527us/step - Recall: 0.8243 - loss: 0.054
4 - val_Recall: 0.8153 - val_loss: 0.0648
Epoch 42/56

63/63 ————— **0s** 555us/step - Recall: 0.8232 - loss: 0.054
3 - val_Recall: 0.8153 - val_loss: 0.0647
Epoch 43/56

63/63 ————— **0s** 532us/step - Recall: 0.8243 - loss: 0.054
1 - val_Recall: 0.8198 - val_loss: 0.0647
Epoch 44/56

63/63 ————— **0s** 540us/step - Recall: 0.8232 - loss: 0.054
0 - val_Recall: 0.8198 - val_loss: 0.0647
Epoch 45/56

63/63 ————— **0s** 555us/step - Recall: 0.8255 - loss: 0.053
9 - val_Recall: 0.8198 - val_loss: 0.0646
Epoch 46/56

63/63 ————— **0s** 524us/step - Recall: 0.8243 - loss: 0.053
8 - val_Recall: 0.8198 - val_loss: 0.0646
Epoch 47/56

63/63 ————— **0s** 530us/step - Recall: 0.8243 - loss: 0.053
7 - val_Recall: 0.8198 - val_loss: 0.0646
Epoch 48/56

63/63 ————— **0s** 555us/step - Recall: 0.8266 - loss: 0.053
6 - val_Recall: 0.8198 - val_loss: 0.0644
Epoch 49/56

63/63 ————— **0s** 536us/step - Recall: 0.8255 - loss: 0.053
5 - val_Recall: 0.8198 - val_loss: 0.0644
Epoch 50/56

63/63 ————— **0s** 563us/step - Recall: 0.8243 - loss: 0.053

```


4 - val_Recall: 0.8243 - val_loss: 0.0644
Epoch 51/56
63/63 ————— 0s 576us/step - Recall: 0.8277 - loss: 0.053
3 - val_Recall: 0.8243 - val_loss: 0.0643
Epoch 52/56
63/63 ————— 0s 518us/step - Recall: 0.8288 - loss: 0.053
2 - val_Recall: 0.8198 - val_loss: 0.0643
Epoch 53/56
63/63 ————— 0s 569us/step - Recall: 0.8243 - loss: 0.053
1 - val_Recall: 0.8198 - val_loss: 0.0642
Epoch 54/56
63/63 ————— 0s 544us/step - Recall: 0.8300 - loss: 0.053
0 - val_Recall: 0.8243 - val_loss: 0.0642
Epoch 55/56
63/63 ————— 0s 531us/step - Recall: 0.8255 - loss: 0.052
9 - val_Recall: 0.8288 - val_loss: 0.0642
Epoch 56/56
63/63 ————— 0s 537us/step - Recall: 0.8300 - loss: 0.052
8 - val_Recall: 0.8243 - val_loss: 0.0641

```

```
In [53]: print("Time taken in seconds ",end=start)
```

Time taken in seconds 2.328244924545288

```
In [54]: plot(history, 'loss')
```



```
In [55]: model_0_train_perf = model_performance_classification(model_0, X_train)
```

```
model_0_train_perf
```

500/500 ————— 0s 137us/step

```
Out [55]:
```

	Accuracy	Recall	Precision	F1 Score
0	0.989313	0.913786	0.982432	0.945145

```
In [56]: model_0_val_perf = model_performance_classification(model_0,X_val,y_val)
model_0_val_perf
```

125/125 ————— 0s 233us/step

```
Out [56]:
```

	Accuracy	Recall	Precision	F1 Score
0	0.98875	0.911368	0.97901	0.942291

Let's check the classification reports.

```
In [57]: y_train_pred_0 = model_0.predict(X_train)
y_val_pred_0 = model_0.predict(X_val)
```

500/500 ————— 0s 158us/step

125/125 ————— 0s 151us/step

```
In [58]: print("Classification Report - Train data Model_0",end="\n\n")
cr_train_model_0 = classification_report(y_train,y_train_pred_0>0.5)
print(cr_train_model_0)
```

Classification Report - Train data Model_0

	precision	recall	f1-score	support
0.0	0.99	1.00	0.99	15112
1.0	0.97	0.83	0.90	888
accuracy			0.99	16000
macro avg	0.98	0.91	0.95	16000
weighted avg	0.99	0.99	0.99	16000

```
In [59]: print("Classification Report - Validation data Model_0",end="\n\n")
cr_val_model_0 = classification_report(y_val,y_val_pred_0>0.5)
print(cr_val_model_0)
```

Classification Report – Validation data Model_0

	precision	recall	f1-score	support
0.0	0.99	1.00	0.99	3778
1.0	0.97	0.82	0.89	222
accuracy			0.99	4000
macro avg	0.98	0.91	0.94	4000
weighted avg	0.99	0.99	0.99	4000

In []:

Initial Neural Network Model (Model 0) – Results and Interpretation

Model 0 Performance Results

The initial neural network model (Model 0) was evaluated using both the training and validation datasets.

Training Dataset Results

Metric	No Failure (0)	Failure (1)
Precision	0.99	0.97
Recall	1.00	0.83
F1-Score	0.99	0.90

Overall Training Performance

- Accuracy: **99%**
- Macro Average Recall: **91%**
- Macro Average F1-Score: **95%**

Validation Dataset Results

Metric	No Failure (0)	Failure (1)
Precision	0.99	0.97
Recall	1.00	0.82
F1-Score	0.99	0.89

Overall Validation Performance

- Accuracy: **99%**
 - Macro Average Recall: **91%**
 - Macro Average F1-Score: **94%**
-

Interpretation of Results

Strong Overall Model Performance

The model achieved extremely strong overall performance on both the training and validation datasets.

Key observations include:

- Overall accuracy reached approximately **99%**.
- Precision for failure prediction was approximately **97%**.
- Recall for failure prediction was approximately:
 - **83%** on training data,
 - **82%** on validation data.
- F1-score for the failure class remained high at approximately:
 - **0.90** for training,
 - **0.89** for validation.

These results indicate that the model successfully learned meaningful patterns from the turbine sensor data.

Interpretation of Failure Detection Results

Recall for Failure Class

The Recall score for class 1 (Failure) is particularly important.

A Recall of approximately:

- **82–83%**

means that:

- the model correctly identifies most turbine failures,
- but still misses approximately:
 - **17–18%** of actual failures.

This indicates good failure detection capability, although some failures remain undetected.

Precision for Failure Class

The Precision score for failures is approximately:

- **97%**

This means:

- when the model predicts a turbine failure,
- it is usually correct.

As a result:

- the number of false alarms is relatively low,
 - reducing unnecessary inspections and maintenance activity.
-

F1-Score Interpretation

The F1-score combines:

- Precision,
- and Recall.

A high F1-score (~0.89–0.90) indicates that the model maintains a strong balance between:

- accurately identifying failures,
- and minimizing false alarms.

Validation Results and Generalization

The training and validation results are very similar.

This is important because it suggests:

- the model is generalizing well,
- the learning process is stable,
- and there is limited evidence of severe overfitting.

A large gap between training and validation performance would indicate overfitting, but that is not strongly observed here.

Why These Results Are Important

These results are important because the project focuses on predictive maintenance for wind turbines.

The model's ability to correctly identify failures directly impacts:

- operational reliability,
 - maintenance efficiency,
 - and overall business cost reduction.
-

Business Impact

Reduced Generator Failure Risk

The model successfully identifies most turbine failures before complete breakdown occurs.

This enables:

- earlier intervention,
 - preventive maintenance,
 - and reduced catastrophic equipment failures.
-

Lower Maintenance and Replacement Costs

Because failures can be identified early:

- generators can often be repaired instead of replaced.

This can significantly reduce:

- replacement expenses,
 - emergency repair costs,
 - and operational losses.
-

Reduced Operational Downtime

Accurate failure prediction helps:

- minimize unexpected turbine shutdowns,
- improve equipment availability,
- and maintain stable energy production.

This improves operational efficiency and overall revenue stability.

Reduced Unnecessary Maintenance

The high Precision score (~97%) means:

- most predicted failures are actual failures.

This reduces:

- unnecessary inspections,
 - unnecessary maintenance activities,
 - and wasted operational resources.
-

Remaining Business Risk

Although the model performs very well overall, some actual failures are still missed.

These missed failures are called:

- False Negatives.

False Negatives remain the largest business risk because they may lead to:

- turbine breakdown,
- expensive generator replacement,
- operational downtime,
- and production losses.

Improving Recall further would help reduce this risk even more.

Overall Conclusion

Model 0 provides a very strong baseline neural network model for predictive maintenance.

The results demonstrate:

- excellent overall accuracy,
- strong failure detection capability,
- high precision,
- stable validation performance,
- and strong business value potential.

The model shows significant promise for helping ReneWind:

- reduce maintenance costs,
- improve turbine reliability,
- minimize downtime,
- and transition toward proactive predictive maintenance.

In []:

In []:

In []:

```
In [59]: print("Classification Report - Validation data Model_0",end="\n\n")
cr_val_model_0 = classification_report(y_val,y_val_pred_0>0.5)
print(cr_val_model_0)
```

Classification Report – Validation data Model_0

	precision	recall	f1-score	support
0.0	0.99	1.00	0.99	3778
1.0	0.97	0.82	0.89	222
accuracy			0.99	4000
macro avg	0.98	0.91	0.94	4000
weighted avg	0.99	0.99	0.99	4000

Model Performance Improvement

Model 1

- Let's try adding another layer to see if we can improve our model's performance.

```
In [60]: # clears the current Keras session, resetting all layers and models pr
tf.keras.backend.clear_session()
```

```
In [61]: #Initializing the neural network
model_1 = Sequential()
model_1.add(Dense(64 ,activation="relu",input_dim=X_train.shape[1])) #
model_1.add(Dense( 32,activation="relu")) # Complete the code to defin
model_1.add(Dense(1,activation="sigmoid")) # Complete the code to defi
```

```
In [62]: model_1.summary()
```

Model: "sequential"

Layer (type)	Output Shape	
dense (Dense)	(None, 64)	
dense_1 (Dense)	(None, 32)	
dense_2 (Dense)	(None, 1)	

Total params: 4,737 (18.50 KB)

Trainable params: 4,737 (18.50 KB)

Non-trainable params: 0 (0.00 B)

```
In [63]: optimizer = tf.keras.optimizers.SGD() # defining SGD as the optimize
```

```
# model_1.compile(loss='binary_crossentropy', optimizer=optimizer, met
# model_1.compile(loss='binary_crossentropy', optimizer=optimizer, met
# model_1.compile(loss='binary_crossentropy', optimizer=optimizer, met
# model_1.compile(loss='binary_crossentropy', optimizer=optimizer, met
```

In [65]: `optimizer = tf.keras.optimizers.SGD()`

```
model_1.compile(
    loss="binary_crossentropy",
    optimizer=optimizer,
    metrics=["Recall"]
)
```

In [66]: `model_1.summary()`

```
optimizer = tf.keras.optimizers.SGD()

model_1.compile(
    loss='binary_crossentropy',
    optimizer=optimizer,
    metrics=['Recall']
)

start = time.time()

history = model_1.fit(
    X_train,
    y_train,
    validation_data=(X_val, y_val),
    batch_size=batch_size,
    epochs=epochs
)

end = time.time()

print("Time taken in seconds ", end - start)
```

Model: "sequential"

Layer (type)	Output Shape	
dense (Dense)	(None, 64)	
dense_1 (Dense)	(None, 32)	
dense_2 (Dense)	(None, 1)	

Total params: 4,737 (18.50 KB)

Trainable params: 4,737 (18.50 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/56
63/63 ————— **0s** 2ms/step – Recall: 0.0935 – loss: 0.2164
– val_Recall: 0.1622 – val_loss: 0.1737

Epoch 2/56
63/63 ————— **0s** 611us/step – Recall: 0.1858 – loss: 0.1643
– val_Recall: 0.2613 – val_loss: 0.1505

Epoch 3/56
63/63 ————— **0s** 549us/step – Recall: 0.2939 – loss: 0.1434
– val_Recall: 0.3604 – val_loss: 0.1360

Epoch 4/56
63/63 ————— **0s** 564us/step – Recall: 0.3806 – loss: 0.1295
– val_Recall: 0.4234 – val_loss: 0.1252

Epoch 5/56
63/63 ————— **0s** 597us/step – Recall: 0.4493 – loss: 0.1190
– val_Recall: 0.4775 – val_loss: 0.1169

Epoch 6/56
63/63 ————— **0s** 574us/step – Recall: 0.4831 – loss: 0.1109
– val_Recall: 0.5360 – val_loss: 0.1103

Epoch 7/56
63/63 ————— **0s** 536us/step – Recall: 0.5214 – loss: 0.1045
– val_Recall: 0.5676 – val_loss: 0.1050

Epoch 8/56
63/63 ————— **0s** 594us/step – Recall: 0.5574 – loss: 0.0993
– val_Recall: 0.5946 – val_loss: 0.1008

Epoch 9/56
63/63 ————— **0s** 544us/step – Recall: 0.5800 – loss: 0.0950
– val_Recall: 0.6216 – val_loss: 0.0974

Epoch 10/56
63/63 ————— **0s** 556us/step – Recall: 0.6160 – loss: 0.0915
– val_Recall: 0.6667 – val_loss: 0.0944

Epoch 11/56
63/63 ————— **0s** 594us/step – Recall: 0.6363 – loss: 0.0884
– val_Recall: 0.6712 – val_loss: 0.0920

Epoch 12/56
63/63 ————— **0s** 543us/step – Recall: 0.6396 – loss: 0.0858
– val_Recall: 0.7072 – val_loss: 0.0900

Epoch 13/56
63/63 ————— **0s** 537us/step – Recall: 0.6734 – loss: 0.0834
– val_Recall: 0.7072 – val_loss: 0.0880

Epoch 14/56
63/63 ————— **0s** 539us/step – Recall: 0.6802 – loss: 0.0813
– val_Recall: 0.7342 – val_loss: 0.0865

Epoch 15/56
63/63 ————— **0s** 528us/step – Recall: 0.6937 – loss: 0.0794
– val_Recall: 0.7432 – val_loss: 0.0850

Epoch 16/56
63/63 ————— **0s** 597us/step – Recall: 0.7128 – loss: 0.0778
– val_Recall: 0.7477 – val_loss: 0.0836

Epoch 17/56
63/63 ————— **0s** 522us/step – Recall: 0.7218 – loss: 0.0762
– val_Recall: 0.7477 – val_loss: 0.0825

Epoch 18/56

63/63 ————— **0s** 560us/step - Recall: 0.7264 - loss: 0.074
8 - val_Recall: 0.7477 - val_loss: 0.0814
Epoch 19/56

63/63 ————— **0s** 551us/step - Recall: 0.7320 - loss: 0.073
5 - val_Recall: 0.7477 - val_loss: 0.0804
Epoch 20/56

63/63 ————— **0s** 556us/step - Recall: 0.7432 - loss: 0.072
4 - val_Recall: 0.7523 - val_loss: 0.0795
Epoch 21/56

63/63 ————— **0s** 544us/step - Recall: 0.7534 - loss: 0.071
3 - val_Recall: 0.7523 - val_loss: 0.0786
Epoch 22/56

63/63 ————— **0s** 558us/step - Recall: 0.7545 - loss: 0.070
3 - val_Recall: 0.7523 - val_loss: 0.0779
Epoch 23/56

63/63 ————— **0s** 526us/step - Recall: 0.7624 - loss: 0.069
3 - val_Recall: 0.7568 - val_loss: 0.0772
Epoch 24/56

63/63 ————— **0s** 519us/step - Recall: 0.7624 - loss: 0.068
4 - val_Recall: 0.7568 - val_loss: 0.0765
Epoch 25/56

63/63 ————— **0s** 525us/step - Recall: 0.7725 - loss: 0.067
6 - val_Recall: 0.7568 - val_loss: 0.0759
Epoch 26/56

63/63 ————— **0s** 537us/step - Recall: 0.7782 - loss: 0.066
8 - val_Recall: 0.7613 - val_loss: 0.0753
Epoch 27/56

63/63 ————— **0s** 524us/step - Recall: 0.7793 - loss: 0.066
0 - val_Recall: 0.7568 - val_loss: 0.0748
Epoch 28/56

63/63 ————— **0s** 534us/step - Recall: 0.7860 - loss: 0.065
3 - val_Recall: 0.7568 - val_loss: 0.0743
Epoch 29/56

63/63 ————— **0s** 570us/step - Recall: 0.7905 - loss: 0.064
6 - val_Recall: 0.7613 - val_loss: 0.0738
Epoch 30/56

63/63 ————— **0s** 554us/step - Recall: 0.7950 - loss: 0.064
0 - val_Recall: 0.7613 - val_loss: 0.0733
Epoch 31/56

63/63 ————— **0s** 532us/step - Recall: 0.7984 - loss: 0.063
4 - val_Recall: 0.7658 - val_loss: 0.0728
Epoch 32/56

63/63 ————— **0s** 548us/step - Recall: 0.8041 - loss: 0.062
9 - val_Recall: 0.7658 - val_loss: 0.0725
Epoch 33/56

63/63 ————— **0s** 560us/step - Recall: 0.8052 - loss: 0.062
3 - val_Recall: 0.7658 - val_loss: 0.0721
Epoch 34/56

63/63 ————— **0s** 583us/step - Recall: 0.8074 - loss: 0.061
8 - val_Recall: 0.7658 - val_loss: 0.0717
Epoch 35/56

63/63 ————— **0s** 528us/step - Recall: 0.8086 - loss: 0.061

2 - val_Recall: 0.7703 - val_loss: 0.0714
Epoch 36/56
63/63  0s 554us/step - Recall: 0.8142 - loss: 0.060
8 - val_Recall: 0.7748 - val_loss: 0.0711
Epoch 37/56
63/63  0s 539us/step - Recall: 0.8119 - loss: 0.060
3 - val_Recall: 0.7793 - val_loss: 0.0707
Epoch 38/56
63/63  0s 596us/step - Recall: 0.8164 - loss: 0.059
9 - val_Recall: 0.7883 - val_loss: 0.0704
Epoch 39/56
63/63  0s 591us/step - Recall: 0.8176 - loss: 0.059
4 - val_Recall: 0.7883 - val_loss: 0.0702
Epoch 40/56
63/63  0s 553us/step - Recall: 0.8164 - loss: 0.059
0 - val_Recall: 0.7883 - val_loss: 0.0699
Epoch 41/56
63/63  0s 604us/step - Recall: 0.8187 - loss: 0.058
6 - val_Recall: 0.7883 - val_loss: 0.0697
Epoch 42/56
63/63  0s 558us/step - Recall: 0.8209 - loss: 0.058
2 - val_Recall: 0.7928 - val_loss: 0.0694
Epoch 43/56
63/63  0s 539us/step - Recall: 0.8187 - loss: 0.057
8 - val_Recall: 0.7928 - val_loss: 0.0691
Epoch 44/56
63/63  0s 628us/step - Recall: 0.8209 - loss: 0.057
4 - val_Recall: 0.7928 - val_loss: 0.0689
Epoch 45/56
63/63  0s 545us/step - Recall: 0.8209 - loss: 0.057
1 - val_Recall: 0.7973 - val_loss: 0.0687
Epoch 46/56
63/63  0s 543us/step - Recall: 0.8243 - loss: 0.056
7 - val_Recall: 0.7973 - val_loss: 0.0685
Epoch 47/56
63/63  0s 551us/step - Recall: 0.8277 - loss: 0.056
4 - val_Recall: 0.8018 - val_loss: 0.0682
Epoch 48/56
63/63  0s 568us/step - Recall: 0.8277 - loss: 0.056
0 - val_Recall: 0.8018 - val_loss: 0.0680
Epoch 49/56
63/63  0s 584us/step - Recall: 0.8266 - loss: 0.055
7 - val_Recall: 0.8063 - val_loss: 0.0678
Epoch 50/56
63/63  0s 635us/step - Recall: 0.8300 - loss: 0.055
4 - val_Recall: 0.8063 - val_loss: 0.0676
Epoch 51/56
63/63  0s 519us/step - Recall: 0.8300 - loss: 0.055
1 - val_Recall: 0.8063 - val_loss: 0.0674
Epoch 52/56
63/63  0s 581us/step - Recall: 0.8311 - loss: 0.054
7 - val_Recall: 0.8063 - val_loss: 0.0671

Epoch 53/56

63/63  **0s** 536us/step - Recall: 0.8322 - loss: 0.0544 - val_Recall: 0.8063 - val_loss: 0.0669

Epoch 54/56

63/63  **0s** 560us/step - Recall: 0.8322 - loss: 0.0541 - val_Recall: 0.8018 - val_loss: 0.0668

Epoch 55/56

63/63  **0s** 550us/step - Recall: 0.8333 - loss: 0.0538 - val_Recall: 0.8063 - val_loss: 0.0665

Epoch 56/56

63/63  **0s** 548us/step - Recall: 0.8333 - loss: 0.0536 - val_Recall: 0.8108 - val_loss: 0.0663

Time taken in seconds 2.6005098819732666

In []:

In []:

In []:

In []:

In []:

In [67]: `start = time.time()
history = model_1.fit(X_train, y_train, validation_data=(X_val,y_val)
end=time.time()`

Epoch 1/56

63/63  **0s** 934us/step - Recall: 0.8356 - loss: 0.0533 - val_Recall: 0.8108 - val_loss: 0.0662

Epoch 2/56

63/63  **0s** 566us/step - Recall: 0.8356 - loss: 0.0530 - val_Recall: 0.8108 - val_loss: 0.0661

Epoch 3/56

63/63  **0s** 554us/step - Recall: 0.8367 - loss: 0.0528 - val_Recall: 0.8153 - val_loss: 0.0657

Epoch 4/56

63/63  **0s** 564us/step - Recall: 0.8367 - loss: 0.0525 - val_Recall: 0.8153 - val_loss: 0.0656

Epoch 5/56

63/63  **0s** 535us/step - Recall: 0.8390 - loss: 0.0523 - val_Recall: 0.8108 - val_loss: 0.0654

Epoch 6/56




63/63  **0s** 532us/step - Recall: 0.8367 - loss: 0.0520 - val_Recall: 0.8108 - val_loss: 0.0653

Epoch 7/56

63/63  **0s** 566us/step - Recall: 0.8390 - loss: 0.0518 - val_Recall: 0.8108 - val_loss: 0.0653

Epoch 8/56

63/63  **0s** 538us/step - Recall: 0.8390 - loss: 0.051

6 - val_Recall: 0.8108 - val_loss: 0.0651
Epoch 9/56
63/63  0s 551us/step - Recall: 0.8401 - loss: 0.051
4 - val_Recall: 0.8108 - val_loss: 0.0649
Epoch 10/56
63/63  0s 532us/step - Recall: 0.8401 - loss: 0.051
1 - val_Recall: 0.8108 - val_loss: 0.0648
Epoch 11/56
63/63  0s 514us/step - Recall: 0.8412 - loss: 0.050
9 - val_Recall: 0.8108 - val_loss: 0.0647
Epoch 12/56
63/63  0s 538us/step - Recall: 0.8412 - loss: 0.050
7 - val_Recall: 0.8108 - val_loss: 0.0646
Epoch 13/56
63/63  0s 546us/step - Recall: 0.8423 - loss: 0.050
5 - val_Recall: 0.8108 - val_loss: 0.0644
Epoch 14/56
63/63  0s 556us/step - Recall: 0.8435 - loss: 0.050
3 - val_Recall: 0.8108 - val_loss: 0.0644
Epoch 15/56
63/63  0s 662us/step - Recall: 0.8401 - loss: 0.050
1 - val_Recall: 0.8108 - val_loss: 0.0642
Epoch 16/56
63/63  0s 626us/step - Recall: 0.8423 - loss: 0.049
9 - val_Recall: 0.8108 - val_loss: 0.0641
Epoch 17/56
63/63  0s 620us/step - Recall: 0.8412 - loss: 0.049
7 - val_Recall: 0.8108 - val_loss: 0.0640
Epoch 18/56
63/63  0s 711us/step - Recall: 0.8435 - loss: 0.049
5 - val_Recall: 0.8108 - val_loss: 0.0637
Epoch 19/56
63/63  0s 573us/step - Recall: 0.8457 - loss: 0.049
3 - val_Recall: 0.8108 - val_loss: 0.0637
Epoch 20/56
63/63  0s 589us/step - Recall: 0.8468 - loss: 0.049
1 - val_Recall: 0.8108 - val_loss: 0.0637
Epoch 21/56
63/63  0s 558us/step - Recall: 0.8435 - loss: 0.049
0 - val_Recall: 0.8108 - val_loss: 0.0636
Epoch 22/56
63/63  0s 569us/step - Recall: 0.8412 - loss: 0.048
8 - val_Recall: 0.8108 - val_loss: 0.0635
Epoch 23/56
63/63  0s 547us/step - Recall: 0.8457 - loss: 0.048
6 - val_Recall: 0.8108 - val_loss: 0.0633
Epoch 24/56
63/63  0s 562us/step - Recall: 0.8446 - loss: 0.048
5 - val_Recall: 0.8153 - val_loss: 0.0631
Epoch 25/56
63/63  0s 551us/step - Recall: 0.8480 - loss: 0.048
3 - val_Recall: 0.8108 - val_loss: 0.0632

Epoch 26/56
63/63  **0s** 539us/step - Recall: 0.8480 - loss: 0.048
1 - val_Recall: 0.8108 - val_loss: 0.0630
Epoch 27/56
63/63  **0s** 523us/step - Recall: 0.8457 - loss: 0.048
0 - val_Recall: 0.8198 - val_loss: 0.0628
Epoch 28/56
63/63  **0s** 503us/step - Recall: 0.8525 - loss: 0.047
8 - val_Recall: 0.8153 - val_loss: 0.0628
Epoch 29/56
63/63  **0s** 530us/step - Recall: 0.8525 - loss: 0.047
7 - val_Recall: 0.8108 - val_loss: 0.0628
Epoch 30/56
63/63  **0s** 505us/step - Recall: 0.8457 - loss: 0.047
5 - val_Recall: 0.8198 - val_loss: 0.0626
Epoch 31/56
63/63  **0s** 495us/step - Recall: 0.8536 - loss: 0.047
4 - val_Recall: 0.8153 - val_loss: 0.0626
Epoch 32/56
63/63  **0s** 474us/step - Recall: 0.8525 - loss: 0.047
2 - val_Recall: 0.8153 - val_loss: 0.0625
Epoch 33/56
63/63  **0s** 483us/step - Recall: 0.8547 - loss: 0.047
1 - val_Recall: 0.8153 - val_loss: 0.0625
Epoch 34/56
63/63  **0s** 475us/step - Recall: 0.8536 - loss: 0.046
9 - val_Recall: 0.8198 - val_loss: 0.0624
Epoch 35/56
63/63  **0s** 488us/step - Recall: 0.8559 - loss: 0.046
8 - val_Recall: 0.8198 - val_loss: 0.0623
Epoch 36/56
63/63  **0s** 476us/step - Recall: 0.8536 - loss: 0.046
6 - val_Recall: 0.8198 - val_loss: 0.0622
Epoch 37/56
63/63  **0s** 467us/step - Recall: 0.8581 - loss: 0.046
5 - val_Recall: 0.8198 - val_loss: 0.0622
Epoch 38/56
63/63  **0s** 482us/step - Recall: 0.8559 - loss: 0.046
4 - val_Recall: 0.8198 - val_loss: 0.0621
Epoch 39/56
63/63  **0s** 486us/step - Recall: 0.8559 - loss: 0.046
2 - val_Recall: 0.8198 - val_loss: 0.0620
Epoch 40/56
63/63  **0s** 470us/step - Recall: 0.8581 - loss: 0.046
1 - val_Recall: 0.8198 - val_loss: 0.0619
Epoch 41/56
63/63  **0s** 485us/step - Recall: 0.8570 - loss: 0.046
0 - val_Recall: 0.8198 - val_loss: 0.0618
Epoch 42/56
63/63  **0s** 477us/step - Recall: 0.8570 - loss: 0.045
9 - val_Recall: 0.8333 - val_loss: 0.0618
Epoch 43/56

```


63/63 ————— 0s 471us/step - Recall: 0.8536 - loss: 0.045
7 - val_Recall: 0.8333 - val_loss: 0.0617
Epoch 44/56
63/63 ————— 0s 478us/step - Recall: 0.8581 - loss: 0.045
6 - val_Recall: 0.8288 - val_loss: 0.0616
Epoch 45/56
63/63 ————— 0s 471us/step - Recall: 0.8581 - loss: 0.045
5 - val_Recall: 0.8333 - val_loss: 0.0615
Epoch 46/56
63/63 ————— 0s 471us/step - Recall: 0.8592 - loss: 0.045
3 - val_Recall: 0.8333 - val_loss: 0.0616
Epoch 47/56
63/63 ————— 0s 475us/step - Recall: 0.8581 - loss: 0.045
3 - val_Recall: 0.8333 - val_loss: 0.0615
Epoch 48/56
63/63 ————— 0s 481us/step - Recall: 0.8592 - loss: 0.045
1 - val_Recall: 0.8333 - val_loss: 0.0615
Epoch 49/56
63/63 ————— 0s 521us/step - Recall: 0.8604 - loss: 0.045
0 - val_Recall: 0.8378 - val_loss: 0.0614
Epoch 50/56
63/63 ————— 0s 468us/step - Recall: 0.8581 - loss: 0.044
9 - val_Recall: 0.8423 - val_loss: 0.0614
Epoch 51/56
63/63 ————— 0s 521us/step - Recall: 0.8604 - loss: 0.044
8 - val_Recall: 0.8423 - val_loss: 0.0613
Epoch 52/56
63/63 ————— 0s 512us/step - Recall: 0.8626 - loss: 0.044
7 - val_Recall: 0.8423 - val_loss: 0.0613
Epoch 53/56
63/63 ————— 0s 497us/step - Recall: 0.8615 - loss: 0.044
5 - val_Recall: 0.8423 - val_loss: 0.0611
Epoch 54/56
63/63 ————— 0s 512us/step - Recall: 0.8604 - loss: 0.044
5 - val_Recall: 0.8423 - val_loss: 0.0612
Epoch 55/56
63/63 ————— 0s 511us/step - Recall: 0.8604 - loss: 0.044
3 - val_Recall: 0.8423 - val_loss: 0.0612
Epoch 56/56
63/63 ————— 0s 493us/step - Recall: 0.8637 - loss: 0.044
2 - val_Recall: 0.8468 - val_loss: 0.0611

```

```
In [68]: print("Time taken in seconds ",end=start)
```

Time taken in seconds 2.2073919773101807

```
In [69]: plot(history, 'loss')
```



```
In [70]: model_1_train_perf = model_performance_classification(model_1,X_train,
model_1_train_perf)
```

500/500 ————— 0s 155us/step

```
Out[70]:
```

	Accuracy	Recall	Precision	F1 Score
0	0.9915	0.931903	0.985853	0.957088

```
In [71]: model_1_val_perf = model_performance_classification(model_1,X_val,y_val,
model_1_val_perf)
```

125/125 ————— 0s 198us/step

```
Out[71]:
```

	Accuracy	Recall	Precision	F1 Score
0	0.9895	0.922365	0.975123	0.946991

```
In [72]: y_train_pred_1 = model_1.predict(X_train)
y_val_pred_1 = model_1.predict(X_val)
```

500/500 ————— 0s 158us/step
125/125 ————— 0s 158us/step

```
In [73]: print("Classification Report - Train data Model_1", end="\n\n")
cr_train_model_1 = classification_report(y_train,y_train_pred_1 > 0.5)
print(cr_train_model_1)
```

Classification Report – Train data Model_1

	precision	recall	f1-score	support
0.0	0.99	1.00	1.00	15112
1.0	0.98	0.86	0.92	888
accuracy			0.99	16000
macro avg	0.99	0.93	0.96	16000
weighted avg	0.99	0.99	0.99	16000

```
In [74]: print("Classification Report – Validation data Model_1", end="\n\n")
cr_val_model_1 = classification_report(y_val,y_val_pred_1 > 0.5)
print(cr_val_model_1)
```

Classification Report – Validation data Model_1

	precision	recall	f1-score	support
0.0	0.99	1.00	0.99	3778
1.0	0.96	0.85	0.90	222
accuracy			0.99	4000
macro avg	0.98	0.92	0.95	4000
weighted avg	0.99	0.99	0.99	4000

Model 1 – Neural Network Performance Improvement

Model Architecture

Model 1 was built as a more complex neural network than Model 0 by adding an additional hidden layer.

Architecture

- Input features: 40 sensor variables
- Hidden Layer 1: 64 neurons with ReLU activation
- Hidden Layer 2: 32 neurons with ReLU activation
- Output Layer: 1 neuron with Sigmoid activation
- Total trainable parameters: 4,737

The additional hidden layer allows the model to learn more complex nonlinear relationships between turbine sensor readings and generator failure patterns.

Model Loss Results

The loss plot shows that both training loss and validation loss decrease steadily across epochs.

Key Observations

- Training loss decreases consistently over time.
- Validation loss also decreases, though it remains slightly higher than training loss.
- There is no major divergence between training and validation loss curves.

Interpretation

This suggests that:

- the model is learning effectively,
- training is stable,
- the model is not severely overfitting,
- and the neural network is capturing meaningful patterns from the turbine sensor data.

Validation Performance Results

The validation classification report shows the model's performance on unseen validation data.

Class	Precision	Recall	F1-Score	Support
No Failure (0)	0.99	1.00	0.99	3,778
Failure (1)	0.96	0.85	0.90	222

Overall Validation Results

- Accuracy: 99%
- Macro Average Precision: 0.98
- Macro Average Recall: 0.92
- Macro Average F1-Score: 0.95
- Weighted Average F1-Score: 0.99

What the Results Say

Model 1 performs very strongly on the validation data.

For the majority class, No Failure, the model performs almost perfectly:

- Precision is 0.99
- Recall is 1.00
- F1-score is 0.99

This means the model is highly effective at identifying normal turbine operations.

For the failure class, which is the most important class for this project:

- Precision is 0.96
- Recall is 0.85
- F1-score is 0.90

A failure precision of 0.96 means that when the model predicts a turbine failure, it is correct most of the time. This helps avoid unnecessary inspections and wasted maintenance resources.

A failure recall of 0.85 means the model correctly identifies about 85% of actual turbine failures. This is a strong result because failure cases are rare in the dataset.

The failure F1-score of 0.90 shows that the model has a strong balance between detecting failures and avoiding false alarms.

Why This Is Important

The dataset is highly imbalanced, with only about 5.5% of observations representing actual failures. Because of this, accuracy alone can be misleading.

A model could achieve high accuracy by mostly predicting "No Failure," but that would not solve the business problem. The main objective is to detect real failures early.

Model 1 is important because it does not only achieve high accuracy; it also performs well on the minority failure class.

The validation recall of 85% for failures shows that the model can identify most actual failures before breakdown occurs. This is critical because missed failures,

or False Negatives, are the most expensive type of error in this project.

Business Impact

Model 1 can support ReneWind's predictive maintenance strategy by identifying turbines at risk of generator failure before complete breakdown.

Key Business Benefits

- Helps detect most failure events early
- Reduces unexpected turbine breakdowns
- Supports proactive maintenance planning
- Lowers emergency repair and replacement costs
- Reduces operational downtime
- Improves turbine reliability and availability
- Helps maintenance teams focus on high-risk turbines

The high precision of 96% for failure predictions also means the model is not creating excessive false alarms. This helps reduce unnecessary inspections and improves trust in the model's recommendations.

Remaining Business Risk

Although Model 1 performs well, failure recall is still 85%, not 100%.

This means approximately 15% of actual failures may still be missed.

These missed failures remain a business risk because they can lead to:

- generator breakdown,
- expensive replacement costs,
- production losses,
- and unplanned downtime.

Further model tuning should focus on improving failure recall while maintaining reasonable precision.

Overall Conclusion

Model 1 shows strong improvement potential and performs well on unseen validation data.

The model demonstrates:

- high overall accuracy,
- strong failure detection performance,
- high precision for failure predictions,
- stable training behavior,
- and limited signs of overfitting.

Overall, Model 1 is a strong candidate for predictive maintenance because it can help ReneWind detect generator failures earlier, reduce maintenance costs, and improve operational reliability.

Model 2

In []:

To introduce Regularization in our model, let's set the dropout to 50% after adding the first hidden layer. This step will randomly drop 50% of the neurons before proceeding to the next layer, reducing overfitting.

```
In [76]: # clears the current Keras session, resetting all layers and models pr
tf.keras.backend.clear_session()
```

In []:

```
In [77]: #Initializing the neural network
from tensorflow.keras.layers import Dropout
model_2 = Sequential()
model_2.add(Dense(64,activation="relu",input_dim=X_train.shape[1])) #
model_2.add(Dropout(0.5)) # Complete the code to define the dropout ra
model_2.add(Dense(32,activation = "relu")) # Complete the code to defi
model_2.add(Dense(16,activation = "relu")) # Complete the code to defi
model_2.add(Dense(1,activation="sigmoid")) # Complete the code to defi
```

```
In [78]: model_2.summary()
```

Model: "sequential"

Layer (type)	Output Shape	
dense (Dense)	(None, 64)	
dropout (Dropout)	(None, 64)	
dense_1 (Dense)	(None, 32)	
dense_2 (Dense)	(None, 16)	
dense_3 (Dense)	(None, 1)	

Total params: 5,249 (20.50 KB)

Trainable params: 5,249 (20.50 KB)

Non-trainable params: 0 (0.00 B)

```
In [79]: optimizer = tf.keras.optimizers.SGD() # defining SGD as the optimizer
# model_2.compile(loss='binary_crossentropy', optimizer=optimizer, met
# model_2.compile(loss='binary_crossentropy', optimizer=optimizer, met
# model_2.compile(loss='binary_crossentropy', optimizer=optimizer, met
# model_2.compile(loss='binary_crossentropy', optimizer=optimizer, met
```

```
In [81]: optimizer = tf.keras.optimizers.SGD()

model_2.compile(
    loss='binary_crossentropy',
    optimizer=optimizer,
    metrics=['Recall']
)

start = time.time()

history = model_2.fit(
    X_train,
    y_train,
    validation_data=(X_val, y_val),
    batch_size=batch_size,
    epochs=epochs
)

end=time.time()

print("Time taken in seconds ",end-start)
```

Epoch 1/56

63/63 ————— **0s** 2ms/step – Recall: 0.1475 – loss: 0.2711
– val_Recall: 0.1081 – val_loss: 0.1834

Epoch 2/56

63/63 ————— **0s** 663us/step – Recall: 0.1475 – loss: 0.209
8 – val_Recall: 0.1982 – val_loss: 0.1641

Epoch 3/56
63/63  **0s** 598us/step - Recall: 0.1959 - loss: 0.194
3 - val_Recall: 0.2523 - val_loss: 0.1527

Epoch 4/56
63/63  **0s** 701us/step - Recall: 0.2455 - loss: 0.182
5 - val_Recall: 0.3063 - val_loss: 0.1442

Epoch 5/56
63/63  **0s** 638us/step - Recall: 0.2455 - loss: 0.175
4 - val_Recall: 0.3468 - val_loss: 0.1380

Epoch 6/56
63/63  **0s** 698us/step - Recall: 0.2759 - loss: 0.165
7 - val_Recall: 0.3829 - val_loss: 0.1329

Epoch 7/56
63/63  **0s** 672us/step - Recall: 0.3198 - loss: 0.163
4 - val_Recall: 0.4099 - val_loss: 0.1287

Epoch 8/56
63/63  **0s** 660us/step - Recall: 0.3322 - loss: 0.155
7 - val_Recall: 0.4234 - val_loss: 0.1248

Epoch 9/56
63/63  **0s** 648us/step - Recall: 0.3378 - loss: 0.152
5 - val_Recall: 0.4550 - val_loss: 0.1216

Epoch 10/56
63/63  **0s** 658us/step - Recall: 0.3705 - loss: 0.149
7 - val_Recall: 0.4640 - val_loss: 0.1186

Epoch 11/56
63/63  **0s** 643us/step - Recall: 0.3784 - loss: 0.146
3 - val_Recall: 0.4730 - val_loss: 0.1161

Epoch 12/56
63/63  **0s** 715us/step - Recall: 0.4032 - loss: 0.141
9 - val_Recall: 0.4910 - val_loss: 0.1137

Epoch 13/56
63/63  **0s** 668us/step - Recall: 0.4043 - loss: 0.141
5 - val_Recall: 0.5000 - val_loss: 0.1115

Epoch 14/56
63/63  **0s** 708us/step - Recall: 0.4279 - loss: 0.136
6 - val_Recall: 0.5315 - val_loss: 0.1093

Epoch 15/56
63/63  **0s** 741us/step - Recall: 0.4133 - loss: 0.137
3 - val_Recall: 0.5586 - val_loss: 0.1072

Epoch 16/56
63/63  **0s** 817us/step - Recall: 0.4324 - loss: 0.137
4 - val_Recall: 0.5766 - val_loss: 0.1056

Epoch 17/56
63/63  **0s** 612us/step - Recall: 0.4572 - loss: 0.127
4 - val_Recall: 0.5901 - val_loss: 0.1040

Epoch 18/56
63/63  **0s** 624us/step - Recall: 0.4572 - loss: 0.128
1 - val_Recall: 0.6036 - val_loss: 0.1023

Epoch 19/56
63/63  **0s** 655us/step - Recall: 0.4561 - loss: 0.130
8 - val_Recall: 0.6081 - val_loss: 0.1008

Epoch 20/56

63/63 ————— **0s** 629us/step - Recall: 0.4764 - loss: 0.127
2 - val_Recall: 0.6126 - val_loss: 0.0993
Epoch 21/56

63/63 ————— **0s** 645us/step - Recall: 0.4899 - loss: 0.124
9 - val_Recall: 0.6171 - val_loss: 0.0979
Epoch 22/56

63/63 ————— **0s** 787us/step - Recall: 0.4932 - loss: 0.122
5 - val_Recall: 0.6396 - val_loss: 0.0965
Epoch 23/56

63/63 ————— **0s** 770us/step - Recall: 0.5056 - loss: 0.118
9 - val_Recall: 0.6261 - val_loss: 0.0952
Epoch 24/56

63/63 ————— **0s** 702us/step - Recall: 0.5034 - loss: 0.117
7 - val_Recall: 0.6441 - val_loss: 0.0941
Epoch 25/56

63/63 ————— **0s** 692us/step - Recall: 0.5169 - loss: 0.116
3 - val_Recall: 0.6486 - val_loss: 0.0929
Epoch 26/56

63/63 ————— **0s** 627us/step - Recall: 0.5259 - loss: 0.118
4 - val_Recall: 0.6577 - val_loss: 0.0919
Epoch 27/56

63/63 ————— **0s** 674us/step - Recall: 0.5248 - loss: 0.115
4 - val_Recall: 0.6622 - val_loss: 0.0907
Epoch 28/56

63/63 ————— **0s** 669us/step - Recall: 0.5428 - loss: 0.115
2 - val_Recall: 0.6622 - val_loss: 0.0897
Epoch 29/56

63/63 ————— **0s** 629us/step - Recall: 0.5495 - loss: 0.113
4 - val_Recall: 0.6622 - val_loss: 0.0888
Epoch 30/56

63/63 ————— **0s** 691us/step - Recall: 0.5394 - loss: 0.111
5 - val_Recall: 0.6757 - val_loss: 0.0878
Epoch 31/56

63/63 ————— **0s** 660us/step - Recall: 0.5552 - loss: 0.112
7 - val_Recall: 0.6757 - val_loss: 0.0868
Epoch 32/56

63/63 ————— **0s** 652us/step - Recall: 0.5484 - loss: 0.112
1 - val_Recall: 0.6892 - val_loss: 0.0860
Epoch 33/56

63/63 ————— **0s** 685us/step - Recall: 0.5811 - loss: 0.108
2 - val_Recall: 0.6847 - val_loss: 0.0854
Epoch 34/56

63/63 ————— **0s** 668us/step - Recall: 0.5664 - loss: 0.109
4 - val_Recall: 0.7027 - val_loss: 0.0845
Epoch 35/56

63/63 ————— **0s** 655us/step - Recall: 0.5608 - loss: 0.107
6 - val_Recall: 0.7027 - val_loss: 0.0838
Epoch 36/56

63/63 ————— **0s** 582us/step - Recall: 0.5800 - loss: 0.107
9 - val_Recall: 0.7027 - val_loss: 0.0832
Epoch 37/56

63/63 ————— **0s** 607us/step - Recall: 0.5890 - loss: 0.104

5 - val_Recall: 0.7117 - val_loss: 0.0824
Epoch 38/56
63/63  0s 647us/step - Recall: 0.6025 - loss: 0.106
6 - val_Recall: 0.7117 - val_loss: 0.0817
Epoch 39/56
63/63  0s 699us/step - Recall: 0.5867 - loss: 0.104
1 - val_Recall: 0.7117 - val_loss: 0.0811
Epoch 40/56
63/63  0s 583us/step - Recall: 0.6070 - loss: 0.102
8 - val_Recall: 0.7162 - val_loss: 0.0805
Epoch 41/56
63/63  0s 562us/step - Recall: 0.6070 - loss: 0.101
7 - val_Recall: 0.7162 - val_loss: 0.0800
Epoch 42/56
63/63  0s 579us/step - Recall: 0.5946 - loss: 0.099
3 - val_Recall: 0.7207 - val_loss: 0.0794
Epoch 43/56
63/63  0s 603us/step - Recall: 0.6014 - loss: 0.105
0 - val_Recall: 0.7252 - val_loss: 0.0789
Epoch 44/56
63/63  0s 620us/step - Recall: 0.6182 - loss: 0.100
6 - val_Recall: 0.7207 - val_loss: 0.0784
Epoch 45/56
63/63  0s 586us/step - Recall: 0.6025 - loss: 0.101
7 - val_Recall: 0.7297 - val_loss: 0.0778
Epoch 46/56
63/63  0s 630us/step - Recall: 0.6160 - loss: 0.100
3 - val_Recall: 0.7342 - val_loss: 0.0773
Epoch 47/56
63/63  0s 574us/step - Recall: 0.6115 - loss: 0.098
8 - val_Recall: 0.7342 - val_loss: 0.0768
Epoch 48/56
63/63  0s 599us/step - Recall: 0.6363 - loss: 0.098
3 - val_Recall: 0.7387 - val_loss: 0.0763
Epoch 49/56
63/63  0s 599us/step - Recall: 0.6374 - loss: 0.097
7 - val_Recall: 0.7387 - val_loss: 0.0759
Epoch 50/56
63/63  0s 593us/step - Recall: 0.6419 - loss: 0.093
0 - val_Recall: 0.7387 - val_loss: 0.0753
Epoch 51/56
63/63  0s 594us/step - Recall: 0.6227 - loss: 0.096
7 - val_Recall: 0.7432 - val_loss: 0.0749
Epoch 52/56
63/63  0s 577us/step - Recall: 0.6351 - loss: 0.095
1 - val_Recall: 0.7432 - val_loss: 0.0745
Epoch 53/56
63/63  0s 544us/step - Recall: 0.6475 - loss: 0.093
0 - val_Recall: 0.7432 - val_loss: 0.0740
Epoch 54/56
63/63  0s 595us/step - Recall: 0.6486 - loss: 0.093
6 - val_Recall: 0.7523 - val_loss: 0.0735

Epoch 55/56

63/63  **0s** 655us/step - Recall: 0.6464 - loss: 0.0940 - val_Recall: 0.7523 - val_loss: 0.0732

Epoch 56/56

63/63  **0s** 610us/step - Recall: 0.6464 - loss: 0.0970 - val_Recall: 0.7523 - val_loss: 0.0729

Time taken in seconds 2.8890140056610107

In []:

In []:

In []:

In []:

In []:

```
In [82]: start = time.time()
history = model_2.fit(X_train, y_train, validation_data=(X_val,y_val)
end=time.time()
```

Epoch 1/56

63/63  **0s** 1ms/step - Recall: 0.6475 - loss: 0.0934 - val_Recall: 0.7568 - val_loss: 0.0726

Epoch 2/56

63/63  **0s** 680us/step - Recall: 0.6577 - loss: 0.0913 - val_Recall: 0.7613 - val_loss: 0.0722

Epoch 3/56

63/63  **0s** 647us/step - Recall: 0.6453 - loss: 0.0891 - val_Recall: 0.7613 - val_loss: 0.0719

Epoch 4/56

63/63  **0s** 716us/step - Recall: 0.6678 - loss: 0.0882 - val_Recall: 0.7568 - val_loss: 0.0716

Epoch 5/56

63/63  **0s** 731us/step - Recall: 0.6532 - loss: 0.0908 - val_Recall: 0.7613 - val_loss: 0.0714

Epoch 6/56

63/63  **0s** 659us/step - Recall: 0.6712 - loss: 0.0880 - val_Recall: 0.7613 - val_loss: 0.0710

Epoch 7/56

63/63  **0s** 683us/step - Recall: 0.6498 - loss: 0.0923 - val_Recall: 0.7613 - val_loss: 0.0706

Epoch 8/56

63/63  **0s** 638us/step - Recall: 0.6622 - loss: 0.0873 - val_Recall: 0.7658 - val_loss: 0.0704

Epoch 9/56

63/63  **0s** 650us/step - Recall: 0.6723 - loss: 0.0898 - val_Recall: 0.7613 - val_loss: 0.0700

Epoch 10/56

63/63  **0s** 645us/step - Recall: 0.6644 - loss: 0.087

1 - val_Recall: 0.7658 - val_loss: 0.0698
Epoch 11/56
63/63  0s 701us/step - Recall: 0.6655 - loss: 0.089
5 - val_Recall: 0.7658 - val_loss: 0.0696
Epoch 12/56
63/63  0s 668us/step - Recall: 0.6565 - loss: 0.091
3 - val_Recall: 0.7658 - val_loss: 0.0693
Epoch 13/56
63/63  0s 711us/step - Recall: 0.6791 - loss: 0.086
7 - val_Recall: 0.7613 - val_loss: 0.0691
Epoch 14/56
63/63  0s 670us/step - Recall: 0.6779 - loss: 0.084
7 - val_Recall: 0.7613 - val_loss: 0.0688
Epoch 15/56
63/63  0s 622us/step - Recall: 0.6655 - loss: 0.088
9 - val_Recall: 0.7793 - val_loss: 0.0686
Epoch 16/56
63/63  0s 717us/step - Recall: 0.6802 - loss: 0.086
2 - val_Recall: 0.7748 - val_loss: 0.0684
Epoch 17/56
63/63  0s 660us/step - Recall: 0.6791 - loss: 0.086
9 - val_Recall: 0.7793 - val_loss: 0.0682
Epoch 18/56
63/63  0s 766us/step - Recall: 0.6824 - loss: 0.085
8 - val_Recall: 0.7748 - val_loss: 0.0680
Epoch 19/56
63/63  0s 679us/step - Recall: 0.6734 - loss: 0.085
2 - val_Recall: 0.7838 - val_loss: 0.0678
Epoch 20/56
63/63  0s 688us/step - Recall: 0.6734 - loss: 0.085
6 - val_Recall: 0.7838 - val_loss: 0.0676
Epoch 21/56
63/63  0s 705us/step - Recall: 0.6937 - loss: 0.083
6 - val_Recall: 0.7793 - val_loss: 0.0675
Epoch 22/56
63/63  0s 698us/step - Recall: 0.6779 - loss: 0.085
7 - val_Recall: 0.7793 - val_loss: 0.0672
Epoch 23/56
63/63  0s 648us/step - Recall: 0.6824 - loss: 0.083
6 - val_Recall: 0.7793 - val_loss: 0.0671
Epoch 24/56
63/63  0s 660us/step - Recall: 0.7072 - loss: 0.084
5 - val_Recall: 0.7748 - val_loss: 0.0670
Epoch 25/56
63/63  0s 671us/step - Recall: 0.6802 - loss: 0.080
8 - val_Recall: 0.7793 - val_loss: 0.0667
Epoch 26/56
63/63  0s 684us/step - Recall: 0.6892 - loss: 0.082
8 - val_Recall: 0.7838 - val_loss: 0.0666
Epoch 27/56
63/63  0s 621us/step - Recall: 0.6881 - loss: 0.084
3 - val_Recall: 0.7838 - val_loss: 0.0664

Epoch 28/56
63/63  **0s** 682us/step - Recall: 0.7005 - loss: 0.083
9 - val_Recall: 0.7973 - val_loss: 0.0663
Epoch 29/56
63/63  **0s** 638us/step - Recall: 0.7005 - loss: 0.082
5 - val_Recall: 0.7973 - val_loss: 0.0661
Epoch 30/56
63/63  **0s** 627us/step - Recall: 0.7050 - loss: 0.083
1 - val_Recall: 0.7883 - val_loss: 0.0659
Epoch 31/56
63/63  **0s** 710us/step - Recall: 0.6836 - loss: 0.082
3 - val_Recall: 0.7928 - val_loss: 0.0658
Epoch 32/56
63/63  **0s** 696us/step - Recall: 0.6847 - loss: 0.081
4 - val_Recall: 0.7973 - val_loss: 0.0656
Epoch 33/56
63/63  **0s** 709us/step - Recall: 0.6926 - loss: 0.080
8 - val_Recall: 0.7928 - val_loss: 0.0655
Epoch 34/56
63/63  **0s** 637us/step - Recall: 0.7016 - loss: 0.079
6 - val_Recall: 0.7928 - val_loss: 0.0655
Epoch 35/56
63/63  **0s** 635us/step - Recall: 0.6881 - loss: 0.081
3 - val_Recall: 0.8018 - val_loss: 0.0654
Epoch 36/56
63/63  **0s** 634us/step - Recall: 0.7027 - loss: 0.080
7 - val_Recall: 0.7928 - val_loss: 0.0652
Epoch 37/56
63/63  **0s** 655us/step - Recall: 0.6768 - loss: 0.081
9 - val_Recall: 0.8018 - val_loss: 0.0651
Epoch 38/56
63/63  **0s** 689us/step - Recall: 0.7038 - loss: 0.079
9 - val_Recall: 0.7928 - val_loss: 0.0648
Epoch 39/56
63/63  **0s** 648us/step - Recall: 0.7286 - loss: 0.076
7 - val_Recall: 0.7928 - val_loss: 0.0646
Epoch 40/56
63/63  **0s** 679us/step - Recall: 0.7072 - loss: 0.077
7 - val_Recall: 0.7973 - val_loss: 0.0644
Epoch 41/56
63/63  **0s** 706us/step - Recall: 0.7027 - loss: 0.079
0 - val_Recall: 0.7928 - val_loss: 0.0643
Epoch 42/56
63/63  **0s** 692us/step - Recall: 0.7038 - loss: 0.078
6 - val_Recall: 0.8063 - val_loss: 0.0642
Epoch 43/56
63/63  **0s** 716us/step - Recall: 0.7072 - loss: 0.079
3 - val_Recall: 0.8108 - val_loss: 0.0640
Epoch 44/56
63/63  **0s** 651us/step - Recall: 0.7162 - loss: 0.078
6 - val_Recall: 0.7973 - val_loss: 0.0639
Epoch 45/56

```


63/63 ————— 0s 701us/step - Recall: 0.7083 - loss: 0.079
1 - val_Recall: 0.7973 - val_loss: 0.0638
Epoch 46/56
63/63 ————— 0s 655us/step - Recall: 0.7151 - loss: 0.077
1 - val_Recall: 0.8108 - val_loss: 0.0637
Epoch 47/56
63/63 ————— 0s 735us/step - Recall: 0.7005 - loss: 0.082
2 - val_Recall: 0.8108 - val_loss: 0.0637
Epoch 48/56
63/63 ————— 0s 756us/step - Recall: 0.7173 - loss: 0.078
1 - val_Recall: 0.8018 - val_loss: 0.0635
Epoch 49/56
63/63 ————— 0s 752us/step - Recall: 0.7050 - loss: 0.078
1 - val_Recall: 0.8063 - val_loss: 0.0634
Epoch 50/56
63/63 ————— 0s 695us/step - Recall: 0.7106 - loss: 0.077
9 - val_Recall: 0.8153 - val_loss: 0.0634
Epoch 51/56
63/63 ————— 0s 624us/step - Recall: 0.7151 - loss: 0.075
4 - val_Recall: 0.8153 - val_loss: 0.0633
Epoch 52/56
63/63 ————— 0s 810us/step - Recall: 0.7331 - loss: 0.075
8 - val_Recall: 0.8063 - val_loss: 0.0632
Epoch 53/56
63/63 ————— 0s 729us/step - Recall: 0.7196 - loss: 0.076
2 - val_Recall: 0.8153 - val_loss: 0.0632
Epoch 54/56
63/63 ————— 0s 707us/step - Recall: 0.7354 - loss: 0.074
8 - val_Recall: 0.8063 - val_loss: 0.0630
Epoch 55/56
63/63 ————— 0s 751us/step - Recall: 0.7331 - loss: 0.074
5 - val_Recall: 0.8153 - val_loss: 0.0629
Epoch 56/56
63/63 ————— 0s 807us/step - Recall: 0.7185 - loss: 0.076
4 - val_Recall: 0.8153 - val_loss: 0.0629

```

```
In [83]: print("Time taken in seconds ",end=start)
```

```
Time taken in seconds 2.7847299575805664
```

```
In [84]: plot(history, 'loss')
```



Lets check the model performance of model_2 on training and validation data respectively.

```
In [85]: model_2_train_perf = model_performance_classification(model_2,X_train,
model_2_train_perf
```

500/500 ————— **0s** 149us/step

```
Out[85]:
```

	Accuracy	Recall	Precision	F1 Score
0	0.987938	0.902458	0.980119	0.937438

```
In [86]: model_2_val_perf = model_performance_classification(model_2,X_val,y_val)
model_2_val_perf
```

125/125 ————— **0s** 240us/step

```
Out[86]:
```

	Accuracy	Recall	Precision	F1 Score
0	0.9875	0.906467	0.970935	0.936026

```
In [87]: y_train_pred_2 = model_2.predict(X_train)
y_val_pred_2 = model_2.predict(X_val)
```

500/500 ————— **0s** 154us/step

125/125 ————— **0s** 156us/step

Lets check the classification report of model_2 on training and validation data respectively.

```
In [88]: print("Classification Report – Train data Model_2", end="\n\n")
cr_train_model_2 = classification_report(y_train,y_train_pred_2 > 0.5)
print(cr_train_model_2)
```

Classification Report – Train data Model_2

	precision	recall	f1-score	support
0.0	0.99	1.00	0.99	15112
1.0	0.97	0.81	0.88	888
accuracy			0.99	16000
macro avg	0.98	0.90	0.94	16000
weighted avg	0.99	0.99	0.99	16000

```
In [89]: print("Classification Report – Validation data Model_2", end="\n\n")
cr_val_model_2 = classification_report(y_val , y_val_pred_2 > 0.5)
print(cr_val_model_2)
```

Classification Report – Validation data Model_2

	precision	recall	f1-score	support
0.0	0.99	1.00	0.99	3778
1.0	0.95	0.82	0.88	222
accuracy			0.99	4000
macro avg	0.97	0.91	0.94	4000
weighted avg	0.99	0.99	0.99	4000

Model 2 – Neural Network with Dropout Regularization

Model Architecture

Model 2 was created to improve generalization and reduce overfitting by adding **Dropout regularization** after the first hidden layer.

Architecture

- Input features: 40 sensor variables
- Hidden Layer 1: 64 neurons with ReLU activation

- Dropout Layer: 50% dropout
- Hidden Layer 2: 32 neurons with ReLU activation
- Hidden Layer 3: 16 neurons with ReLU activation
- Output Layer: 1 neuron with Sigmoid activation
- Total trainable parameters: 5,249

The purpose of adding dropout is to randomly deactivate 50% of neurons during training. This forces the model to learn more general patterns instead of relying too heavily on specific neurons.

Model Loss Results

The loss plot shows that training loss is noisier than validation loss.

Key Observations

- Training loss decreases overall but fluctuates across epochs.
- Validation loss decreases steadily and smoothly.
- Validation loss is lower than training loss, which can happen when dropout is active during training but not during validation.
- There is no major increase in validation loss, so severe overfitting is not observed.

Interpretation

The model is learning useful patterns from the sensor data. The fluctuating training loss is expected because dropout randomly removes neurons during training, making the learning process less smooth. The steady validation loss suggests that dropout helped the model generalize to unseen data.

Model 2 Performance Results

Training Dataset Results

Metric	No Failure (0)	Failure (1)
Precision	0.99	0.97
Recall	1.00	0.81
F1-Score	0.99	0.88

Overall Training Performance

- Accuracy: 99%
 - Macro Average Precision: 0.98
 - Macro Average Recall: 0.90
 - Macro Average F1-Score: 0.94
 - Weighted Average F1-Score: 0.99
-

Validation Dataset Results

Metric	No Failure (0)	Failure (1)
Precision	0.99	0.95
Recall	1.00	0.82
F1-Score	0.99	0.88

Overall Validation Performance

- Accuracy: 99%
 - Macro Average Precision: 0.97
 - Macro Average Recall: 0.91
 - Macro Average F1-Score: 0.94
 - Weighted Average F1-Score: 0.99
-

What These Results Say

Model 2 performs very well overall, but its failure detection performance is slightly weaker than Model 1.

For the **No Failure** class:

- Recall is 1.00 on both training and validation data.
- This means the model is extremely good at identifying normal turbine operations.

For the **Failure** class:

- Training Recall is 0.81.
- Validation Recall is 0.82.

- This means the model correctly identifies about 82% of actual turbine failures on validation data.

For failure predictions:

- Validation Precision is 0.95.
- This means when the model predicts a failure, it is correct 95% of the time.

The Failure F1-Score is 0.88 on validation data, which shows a strong but slightly lower balance between Precision and Recall compared to Model 1.

Why This Is Important

1. Dropout Reduced Overfitting Risk

Dropout is designed to reduce overfitting by preventing the model from relying too heavily on specific neurons.

The loss curve suggests that:

- the model continues learning,
- validation loss decreases steadily,
- and there is no severe overfitting.

This is useful because predictive maintenance models need to perform well on future unseen turbine data, not just the training data.

2. Failure Detection Remains Strong but Lower Than Model 1

The validation Recall for the failure class is 0.82.

This means:

- the model detects most failures,
- but misses about 18% of actual failures.

Because missed failures are the most expensive error in this project, this is an important limitation.

3. High Precision Reduces False Alarms

The validation Precision for the failure class is 0.95.

This means:

- most predicted failures are true failures,
- and the model does not create too many unnecessary maintenance alerts.

This is valuable because excessive false alarms can waste maintenance resources and reduce trust in the model.

4. Accuracy Alone Is Not Enough

The model has 99% accuracy, but the dataset is highly imbalanced.

Since only about 5.5% of observations are failures, a high accuracy score does not fully explain how well the model identifies failures.

Recall and F1-score are more important for judging model usefulness in this business problem.

Business Impact

Model 2 can support predictive maintenance by identifying most generator failures before breakdown occurs.

Positive Business Impact

- Helps detect at-risk turbines early.
- Reduces unexpected generator breakdowns.
- Supports proactive maintenance scheduling.
- Helps reduce emergency repair and replacement costs.
- Maintains high precision, limiting unnecessary inspections.

Remaining Business Risk

Although Model 2 performs well, it misses about 18% of actual failures on validation data.

These missed failures are business-critical because they can lead to:

- expensive generator replacement,
- operational downtime,
- lost energy production,
- and higher maintenance costs.

Comparison to Model 1

Compared with Model 1:

Metric	Model 1 Validation	Model 2 Validation
Accuracy	99%	99%
Failure Precision	0.96	0.95
Failure Recall	0.85	0.82
Failure F1-Score	0.90	0.88

Model 2 has good performance, but Model 1 performs slightly better on the failure class.

This suggests that adding dropout improved regularization but may have slightly reduced the model’s ability to detect failures.

Overall Conclusion

Model 2 is a strong regularized neural network model with stable validation performance and limited overfitting.

However, compared with Model 1, it has slightly lower failure Recall and F1-score.

Since the main business objective is to minimize missed failures, Model 2 may not be the best model so far, but it is still useful because it demonstrates how dropout affects model performance and generalization.

Model 3

As we have are dealing with an imbalance in class distribution, we should also be

using class weights to allow the model to give proportionally more importance to the minority class.

```
In [90]: # Calculate class weights for imbalanced dataset
cw = (y_train.shape[0]) / np.bincount(y_train.astype(int)) # Convert y

# Create a dictionary mapping class indices to their respective class
cw_dict = {}
for i in range(cw.shape[0]):
    cw_dict[i] = cw[i]

cw_dict
```

```
Out[90]: {0: np.float64(1.0587612493382743), 1: np.float64(18.01801801801802)}
```

The history saving thread hit an unexpected error (OperationalError('at tempt to write a readonly database')).History will not be written to the database.

```
In [91]: # clears the current Keras session, resetting all layers and models pr
tf.keras.backend.clear_session()
```

```
In [92]: model_3 = Sequential()
model_3.add(Dense(64,activation="relu",input_dim=X_train.shape[1])) #
model_3.add(Dropout(0.5)) # Complete the code to define the dropout ra
model_3.add(Dense(32,activation="relu")) # Complete the code to define
model_3.add(Dense(16, activation = "relu")) # Complete the code to def
model_3.add(Dense(1,activation="sigmoid")) # Complete the code to defi
```

```
In [93]: model_3.summary()
```

Model: "sequential"

Layer (type)	Output Shape	
dense (Dense)	(None, 64)	
dropout (Dropout)	(None, 64)	
dense_1 (Dense)	(None, 32)	
dense_2 (Dense)	(None, 16)	
dense_3 (Dense)	(None, 1)	

Total params: 5,249 (20.50 KB)

Trainable params: 5,249 (20.50 KB)

Non-trainable params: 0 (0.00 B)

```
In [ ]:
```

```
In [94]: optimizer = tf.keras.optimizers.SGD() # defining SGD as the optimizer
# model_3.compile(loss='binary_crossentropy', optimizer=optimizer, met
# model_3.compile(loss='binary_crossentropy', optimizer=optimizer, met
# model_3.compile(loss='binary_crossentropy', optimizer=optimizer, met
# model_3.compile(loss='binary_crossentropy', optimizer=optimizer, met
```

```
In [96]: optimizer = tf.keras.optimizers.SGD()
```

```
model_3.compile(
    loss="binary_crossentropy",
    optimizer=optimizer,
    metrics=["Recall"]
)
```

```
In [97]: start = time.time()
history = model_3.fit(X_train, y_train, validation_data=(X_val,y_val)
end=time.time())
```

Epoch 1/56

63/63 ————— 0s 2ms/step – Recall: 0.6014 – loss: 1.2761
– val_Recall: 0.9144 – val_loss: 0.4974

Epoch 2/56

63/63 ————— 0s 663us/step – Recall: 0.8322 – loss: 0.957
0 – val_Recall: 0.8739 – val_loss: 0.3476

Epoch 3/56

63/63 ————— 0s 615us/step – Recall: 0.8536 – loss: 0.884
0 – val_Recall: 0.8874 – val_loss: 0.3413

Epoch 4/56

63/63 ————— 0s 645us/step – Recall: 0.8502 – loss: 0.803
3 – val_Recall: 0.8784 – val_loss: 0.2886

Epoch 5/56

63/63 ————— 0s 642us/step – Recall: 0.8525 – loss: 0.759
8 – val_Recall: 0.8694 – val_loss: 0.2756

Epoch 6/56

63/63 ————— 0s 631us/step – Recall: 0.8547 – loss: 0.716
2 – val_Recall: 0.8649 – val_loss: 0.2756

Epoch 7/56

63/63 ————— 0s 568us/step – Recall: 0.8536 – loss: 0.695
4 – val_Recall: 0.8784 – val_loss: 0.2989

Epoch 8/56

63/63 ————— 0s 636us/step – Recall: 0.8626 – loss: 0.654
3 – val_Recall: 0.8784 – val_loss: 0.3347

Epoch 9/56

63/63 ————— 0s 612us/step – Recall: 0.8637 – loss: 0.643
4 – val_Recall: 0.8649 – val_loss: 0.2055

Epoch 10/56

63/63 ————— 0s 619us/step – Recall: 0.8525 – loss: 0.647
1 – val_Recall: 0.8694 – val_loss: 0.2287

Epoch 11/56

63/63 ————— 0s 620us/step – Recall: 0.8649 – loss: 0.608
7 – val_Recall: 0.8694 – val_loss: 0.1969

Epoch 12/56
63/63  **0s** 647us/step - Recall: 0.8705 - loss: 0.584
0 - val_Recall: 0.8874 - val_loss: 0.2209

Epoch 13/56
63/63  **0s** 576us/step - Recall: 0.8795 - loss: 0.575
2 - val_Recall: 0.8874 - val_loss: 0.2307

Epoch 14/56
63/63  **0s** 616us/step - Recall: 0.8784 - loss: 0.555
1 - val_Recall: 0.8784 - val_loss: 0.2501

Epoch 15/56
63/63  **0s** 582us/step - Recall: 0.8716 - loss: 0.584
7 - val_Recall: 0.8829 - val_loss: 0.2261

Epoch 16/56
63/63  **0s** 607us/step - Recall: 0.8682 - loss: 0.544
5 - val_Recall: 0.8874 - val_loss: 0.2523

Epoch 17/56
63/63  **0s** 579us/step - Recall: 0.8795 - loss: 0.523
6 - val_Recall: 0.8784 - val_loss: 0.2058

Epoch 18/56
63/63  **0s** 632us/step - Recall: 0.8795 - loss: 0.547
8 - val_Recall: 0.8874 - val_loss: 0.2559

Epoch 19/56
63/63  **0s** 660us/step - Recall: 0.8727 - loss: 0.539
8 - val_Recall: 0.8829 - val_loss: 0.2290

Epoch 20/56
63/63  **0s** 689us/step - Recall: 0.8761 - loss: 0.530
7 - val_Recall: 0.8784 - val_loss: 0.2512

Epoch 21/56
63/63  **0s** 627us/step - Recall: 0.8840 - loss: 0.503
3 - val_Recall: 0.8829 - val_loss: 0.2531

Epoch 22/56
63/63  **0s** 634us/step - Recall: 0.8739 - loss: 0.508
7 - val_Recall: 0.8829 - val_loss: 0.2325

Epoch 23/56
63/63  **0s** 595us/step - Recall: 0.8795 - loss: 0.501
2 - val_Recall: 0.8874 - val_loss: 0.2048

Epoch 24/56
63/63  **0s** 650us/step - Recall: 0.8896 - loss: 0.478
0 - val_Recall: 0.8829 - val_loss: 0.2218

Epoch 25/56
63/63  **0s** 607us/step - Recall: 0.8795 - loss: 0.488
7 - val_Recall: 0.8739 - val_loss: 0.1868

Epoch 26/56
63/63  **0s** 569us/step - Recall: 0.8851 - loss: 0.486
1 - val_Recall: 0.8874 - val_loss: 0.2247

Epoch 27/56
63/63  **0s** 623us/step - Recall: 0.8795 - loss: 0.481
4 - val_Recall: 0.8784 - val_loss: 0.1935

Epoch 28/56
63/63  **0s** 678us/step - Recall: 0.8739 - loss: 0.487
0 - val_Recall: 0.8874 - val_loss: 0.1940

Epoch 29/56

63/63 ————— **0s** 691us/step - Recall: 0.8896 - loss: 0.470
1 - val_Recall: 0.8784 - val_loss: 0.1941
Epoch 30/56

63/63 ————— **0s** 753us/step - Recall: 0.8806 - loss: 0.466
6 - val_Recall: 0.8829 - val_loss: 0.1665
Epoch 31/56

63/63 ————— **0s** 709us/step - Recall: 0.8863 - loss: 0.460
4 - val_Recall: 0.8829 - val_loss: 0.1937
Epoch 32/56

63/63 ————— **0s** 700us/step - Recall: 0.8919 - loss: 0.449
7 - val_Recall: 0.8829 - val_loss: 0.1693
Epoch 33/56

63/63 ————— **0s** 611us/step - Recall: 0.8773 - loss: 0.472
8 - val_Recall: 0.8829 - val_loss: 0.1944
Epoch 34/56

63/63 ————— **0s** 615us/step - Recall: 0.8953 - loss: 0.445
8 - val_Recall: 0.8829 - val_loss: 0.1894
Epoch 35/56

63/63 ————— **0s** 707us/step - Recall: 0.8874 - loss: 0.456
0 - val_Recall: 0.8829 - val_loss: 0.1642
Epoch 36/56

63/63 ————— **0s** 689us/step - Recall: 0.8874 - loss: 0.448
9 - val_Recall: 0.8919 - val_loss: 0.1893
Epoch 37/56

63/63 ————— **0s** 1ms/step - Recall: 0.8818 - loss: 0.4606
- val_Recall: 0.8919 - val_loss: 0.1737
Epoch 38/56

63/63 ————— **0s** 725us/step - Recall: 0.8885 - loss: 0.450
4 - val_Recall: 0.8874 - val_loss: 0.1533
Epoch 39/56

63/63 ————— **0s** 689us/step - Recall: 0.8863 - loss: 0.459
8 - val_Recall: 0.8919 - val_loss: 0.2118
Epoch 40/56

63/63 ————— **0s** 718us/step - Recall: 0.8840 - loss: 0.452
6 - val_Recall: 0.8874 - val_loss: 0.1910
Epoch 41/56

63/63 ————— **0s** 704us/step - Recall: 0.8919 - loss: 0.444
9 - val_Recall: 0.8874 - val_loss: 0.1910
Epoch 42/56

63/63 ————— **0s** 668us/step - Recall: 0.8908 - loss: 0.440
9 - val_Recall: 0.8829 - val_loss: 0.1728
Epoch 43/56

63/63 ————— **0s** 708us/step - Recall: 0.8975 - loss: 0.435
0 - val_Recall: 0.8829 - val_loss: 0.1928
Epoch 44/56

63/63 ————— **0s** 745us/step - Recall: 0.8874 - loss: 0.440
2 - val_Recall: 0.8829 - val_loss: 0.1799
Epoch 45/56

63/63 ————— **0s** 716us/step - Recall: 0.8908 - loss: 0.425
2 - val_Recall: 0.8919 - val_loss: 0.1804
Epoch 46/56

63/63 ————— **0s** 745us/step - Recall: 0.8885 - loss: 0.438

```


6 - val_Recall: 0.8874 - val_loss: 0.1691
Epoch 47/56
63/63 ————— 0s 689us/step - Recall: 0.8885 - loss: 0.433
9 - val_Recall: 0.8874 - val_loss: 0.1990
Epoch 48/56
63/63 ————— 0s 618us/step - Recall: 0.8975 - loss: 0.425
7 - val_Recall: 0.8829 - val_loss: 0.2024
Epoch 49/56
63/63 ————— 0s 617us/step - Recall: 0.8919 - loss: 0.431
4 - val_Recall: 0.8874 - val_loss: 0.2231
Epoch 50/56
63/63 ————— 0s 654us/step - Recall: 0.8874 - loss: 0.415
0 - val_Recall: 0.8874 - val_loss: 0.1618
Epoch 51/56
63/63 ————— 0s 595us/step - Recall: 0.8930 - loss: 0.419
1 - val_Recall: 0.8874 - val_loss: 0.1571
Epoch 52/56
63/63 ————— 0s 598us/step - Recall: 0.8975 - loss: 0.412
7 - val_Recall: 0.8874 - val_loss: 0.1548
Epoch 53/56
63/63 ————— 0s 658us/step - Recall: 0.8919 - loss: 0.410
3 - val_Recall: 0.8919 - val_loss: 0.1508
Epoch 54/56
63/63 ————— 0s 669us/step - Recall: 0.8896 - loss: 0.424
2 - val_Recall: 0.8919 - val_loss: 0.1822
Epoch 55/56
63/63 ————— 0s 724us/step - Recall: 0.8941 - loss: 0.410
8 - val_Recall: 0.8964 - val_loss: 0.1498
Epoch 56/56
63/63 ————— 0s 668us/step - Recall: 0.8953 - loss: 0.418
8 - val_Recall: 0.8919 - val_loss: 0.1706

```

```
In [98]: print("Time taken in seconds ",end=start)
```

```
Time taken in seconds 2.9769158363342285
```

```
In [99]: plot(history, 'loss')
```

Lets check the model performance of model_3 on training and validation data respectively.

```
In [100... model_3_train_perf = model_performance_classification(model_3,X_train,
model_3_train_perf
```

500/500 ————— 0s 150us/step

```
Out[100... Accuracy Recall Precision F1 Score
0 0.983375 0.946151 0.905198 0.924538
```

```
In [101... model_3_val_perf = model_performance_classification(model_3,X_val,y_val
model_3_val_perf
```

125/125 ————— 0s 249us/step

```
Out[101... Accuracy Recall Precision F1 Score
0 0.984 0.940652 0.912777 0.926191
```

```
In [102... y_train_pred_3 = model_3.predict(X_train)
y_val_pred_3 = model_3.predict(X_val)
```

500/500 ————— 0s 161us/step
125/125 ————— 0s 157us/step

Lets check the classification report of model_3 on training and validation data respectively.

```
In [103... print("Classification Report - Train data Model_3", end="\n\n")
cr_train_model_3 = classification_report(y_train,y_train_pred_3 > 0.5)
print(cr_train_model_3)
```

Classification Report - Train data Model_3

	precision	recall	f1-score	support
0.0	0.99	0.99	0.99	15112
1.0	0.82	0.90	0.86	888
accuracy			0.98	16000
macro avg	0.91	0.95	0.92	16000
weighted avg	0.98	0.98	0.98	16000

```
In [104... print("Classification Report - Validation data Model_3", end="\n\n")
cr_val_model_3 = classification_report(y_val,y_val_pred_3 > 0.5)
print(cr_val_model_3)
```

Classification Report - Validation data Model_3

	precision	recall	f1-score	support
0.0	0.99	0.99	0.99	3778
1.0	0.83	0.89	0.86	222
accuracy			0.98	4000
macro avg	0.91	0.94	0.93	4000
weighted avg	0.98	0.98	0.98	4000

Model 3 – Neural Network with Dropout and Class Weights

Model Architecture

Model 3 was built using the same architecture as Model 2, but class weights were added to handle the imbalance in the target variable.

Architecture

- Input features: 40 sensor variables
- Hidden Layer 1: 64 neurons with ReLU activation

- Dropout Layer: 50% dropout
- Hidden Layer 2: 32 neurons with ReLU activation
- Hidden Layer 3: 16 neurons with ReLU activation
- Output Layer: 1 neuron with Sigmoid activation
- Total trainable parameters: 5,249

Class Weights

Class weights were applied because the dataset is highly imbalanced.

- Class 0 / No Failure weight: 1.06
- Class 1 / Failure weight: 18.02

This means the model gives much more importance to failure cases during training.

Model Loss Results

The loss plot shows that both training and validation loss decrease across epochs.

Key Observations

- Training loss starts high and decreases significantly over time.
- Validation loss also decreases and remains lower than training loss.
- Training loss is noisier because dropout and class weighting make learning more challenging.
- No major upward trend in validation loss is observed.

Interpretation

This suggests that:

- the model is learning effectively,
 - class weights are helping the model focus more on failure cases,
 - dropout is supporting regularization,
 - and the model does not show severe overfitting.
-

Model 3 Performance Results

Training Dataset Results

Metric	No Failure (0)	Failure (1)
Precision	0.99	0.82
Recall	0.99	0.90
F1-Score	0.99	0.86

Overall Training Performance

- Accuracy: 98%
 - Macro Average Precision: 0.91
 - Macro Average Recall: 0.95
 - Macro Average F1-Score: 0.92
 - Weighted Average F1-Score: 0.98
-

Validation Dataset Results

Metric	No Failure (0)	Failure (1)
Precision	0.99	0.83
Recall	0.99	0.89
F1-Score	0.99	0.86

Overall Validation Performance

- Accuracy: 98%
 - Macro Average Precision: 0.91
 - Macro Average Recall: 0.94
 - Macro Average F1-Score: 0.93
 - Weighted Average F1-Score: 0.98
-

What These Results Say

Model 3 shows strong improvement in detecting failure cases compared to the earlier dropout model.

For the **No Failure** class:

- Precision is 0.99.
- Recall is 0.99.
- F1-score is 0.99.

This means the model still performs very well in identifying normal turbine operations.

For the **Failure** class:

- Validation Precision is 0.83.
- Validation Recall is 0.89.
- Validation F1-score is 0.86.

A failure Recall of 0.89 means the model correctly identifies about 89% of actual turbine failures. This is a strong result because failures are rare and costly.

A failure Precision of 0.83 means that some predicted failures are false alarms, but most predicted failures are still correct.

The F1-score of 0.86 shows a reasonable balance between detecting failures and avoiding unnecessary alerts.

Why This Is Important

1. Class Weights Improved Failure Detection

Model 3 uses class weights to give more importance to the minority class, which is turbine failure.

This matters because only about 5.5% of the observations represent failures. Without class weighting, the model may focus too heavily on predicting the majority class, No Failure.

By applying a higher weight to failure cases, the model becomes more sensitive to actual generator failures.

2. Higher Recall Reduces Missed Failures

The validation Recall for failures increased to 0.89.

This means Model 3 misses fewer actual failures than models with lower failure

recall.

This is critical because missed failures, or False Negatives, are the most expensive type of error in this project.

3. Precision Decreased Compared to Earlier Models

The failure Precision is 0.83, which is lower than Model 1 and Model 2.

This means the model predicts more failures overall, which increases the number of false alarms.

However, this tradeoff may be acceptable because inspection costs are lower than replacement costs.

4. Accuracy Dropped Slightly but Business Value Improved

Model 3 has an accuracy of 98%, slightly lower than earlier models with 99% accuracy.

This drop is expected because the model is now more focused on identifying rare failure cases.

For this business problem, a small decrease in accuracy is acceptable if the model catches more actual failures.

Business Impact

Model 3 is valuable because it prioritizes failure detection.

Positive Business Impact

- Detects about 89% of actual failures on validation data.
- Reduces missed generator failures.
- Helps prevent expensive turbine breakdowns.
- Supports earlier maintenance intervention.
- Reduces replacement costs and operational downtime.

- Aligns well with the business goal of predictive maintenance.

Operational Tradeoff

- Precision decreased to 83%.
 - This means there may be more false alarms and inspections.
 - However, inspection costs are lower than replacement costs, so this may be a reasonable tradeoff.
-

Comparison to Previous Models

Metric	Model 1 Validation	Model 2 Validation	Model 3 Validation
Accuracy	99%	99%	98%
Failure Precision	0.96	0.95	0.83
Failure Recall	0.85	0.82	0.89
Failure F1-Score	0.90	0.88	0.86

Model 3 improves failure Recall compared to Model 1 and Model 2, but Precision decreases.

This means Model 3 is better at catching failures, but it may generate more false alarms.

Overall Conclusion

Model 3 is a strong model for the business objective because it improves failure detection.

Although accuracy and precision are slightly lower, the model has higher Recall for the failure class, which is the most important metric in this project.

Because missed failures are more expensive than extra inspections, Model 3 may be more valuable from a business perspective than models with higher precision but lower recall.

Model 4

Since we have used only SGD optimizer till now, let's use another kind of optimizer and observe its impact on the model performance.

```
In [106... # clears the current Keras session, resetting all layers and models pr
tf.keras.backend.clear_session()
```

```
In [107... #Initializing the neural network
model_4 = Sequential()
model_4.add(Dense(64,activation="relu",input_dim=X_train.shape[1])) #
model_4.add(Dense(32,activation="relu")) # Complete the code to define
model_4.add(Dense(1,activation="sigmoid")) # Complete the code to defi
```

```
In [108... model_4.summary()
```

Model: "sequential"

Layer (type)	Output Shape	
dense (Dense)	(None, 64)	
dense_1 (Dense)	(None, 32)	
dense_2 (Dense)	(None, 1)	

Total params: 4,737 (18.50 KB)

Trainable params: 4,737 (18.50 KB)

Non-trainable params: 0 (0.00 B)

```
In [109... optimizer = tf.keras.optimizers.Adam() # defining Adam as the optim
# model_4.compile(loss='binary_crossentropy', optimizer=optimizer, met
# model_4.compile(loss='binary_crossentropy', optimizer=optimizer, met
# model_4.compile(loss='binary_crossentropy', optimizer=optimizer, met
# model_4.compile(loss='binary_crossentropy', optimizer=optimizer, met
```

```
In [110... optimizer = tf.keras.optimizers.SGD()

model_4.compile(
    loss="binary_crossentropy",
    optimizer=optimizer,
    metrics=["Recall"]
)
```

```
In [ ]:
```

```
In [111... start = time.time()
history = model_4.fit(X_train, y_train, validation_data=(X_val,y_val)
end=time.time()
```

Epoch 1/56

63/63 ————— **0s** 1ms/step - Recall: 0.0225 - loss: 0.2044
- val_Recall: 0.0901 - val_loss: 0.1749
Epoch 2/56

63/63 ————— **0s** 529us/step - Recall: 0.1655 - loss: 0.1602
- val_Recall: 0.2297 - val_loss: 0.1464
Epoch 3/56

63/63 ————— **0s** 520us/step - Recall: 0.3041 - loss: 0.1373
- val_Recall: 0.3378 - val_loss: 0.1301
Epoch 4/56

63/63 ————— **0s** 559us/step - Recall: 0.3885 - loss: 0.1230
- val_Recall: 0.4685 - val_loss: 0.1194
Epoch 5/56

63/63 ————— **0s** 534us/step - Recall: 0.4527 - loss: 0.1129
- val_Recall: 0.4955 - val_loss: 0.1111
Epoch 6/56

63/63 ————— **0s** 514us/step - Recall: 0.5011 - loss: 0.1053
- val_Recall: 0.5676 - val_loss: 0.1050
Epoch 7/56

63/63 ————— **0s** 523us/step - Recall: 0.5484 - loss: 0.0993
- val_Recall: 0.6036 - val_loss: 0.1000
Epoch 8/56

63/63 ————— **0s** 523us/step - Recall: 0.5822 - loss: 0.0946
- val_Recall: 0.6261 - val_loss: 0.0960
Epoch 9/56

63/63 ————— **0s** 537us/step - Recall: 0.6194 - loss: 0.0906
- val_Recall: 0.6306 - val_loss: 0.0925
Epoch 10/56

63/63 ————— **0s** 562us/step - Recall: 0.6385 - loss: 0.0874
- val_Recall: 0.6577 - val_loss: 0.0897
Epoch 11/56

63/63 ————— **0s** 557us/step - Recall: 0.6577 - loss: 0.0846
- val_Recall: 0.6892 - val_loss: 0.0876
Epoch 12/56

63/63 ————— **0s** 591us/step - Recall: 0.6892 - loss: 0.0822
- val_Recall: 0.6937 - val_loss: 0.0851
Epoch 13/56

63/63 ————— **0s** 557us/step - Recall: 0.6948 - loss: 0.0801
- val_Recall: 0.7027 - val_loss: 0.0833
Epoch 14/56

63/63 ————— **0s** 534us/step - Recall: 0.7027 - loss: 0.0782
- val_Recall: 0.7297 - val_loss: 0.0817
Epoch 15/56

63/63 ————— **0s** 568us/step - Recall: 0.7061 - loss: 0.0765
- val_Recall: 0.7342 - val_loss: 0.0802
Epoch 16/56

63/63 ————— **0s** 572us/step - Recall: 0.7207 - loss: 0.0750
- val_Recall: 0.7432 - val_loss: 0.0789
Epoch 17/56

63/63 ————— **0s** 636us/step - Recall: 0.7331 - loss: 0.0736
- val_Recall: 0.7477 - val_loss: 0.0776
Epoch 18/56

63/63 ————— **0s** 652us/step - Recall: 0.7376 - loss: 0.072

3 - val_Recall: 0.7568 - val_loss: 0.0766
Epoch 19/56
63/63  0s 638us/step - Recall: 0.7466 - loss: 0.071
1 - val_Recall: 0.7658 - val_loss: 0.0759
Epoch 20/56
63/63  0s 595us/step - Recall: 0.7624 - loss: 0.070
1 - val_Recall: 0.7658 - val_loss: 0.0748
Epoch 21/56
63/63  0s 599us/step - Recall: 0.7624 - loss: 0.069
1 - val_Recall: 0.7703 - val_loss: 0.0738
Epoch 22/56
63/63  0s 620us/step - Recall: 0.7646 - loss: 0.068
1 - val_Recall: 0.7703 - val_loss: 0.0730
Epoch 23/56
63/63  0s 647us/step - Recall: 0.7714 - loss: 0.067
2 - val_Recall: 0.7703 - val_loss: 0.0723
Epoch 24/56
63/63  0s 588us/step - Recall: 0.7725 - loss: 0.066
4 - val_Recall: 0.7703 - val_loss: 0.0717
Epoch 25/56
63/63  0s 647us/step - Recall: 0.7748 - loss: 0.065
7 - val_Recall: 0.7748 - val_loss: 0.0710
Epoch 26/56
63/63  0s 589us/step - Recall: 0.7804 - loss: 0.064
9 - val_Recall: 0.7793 - val_loss: 0.0704
Epoch 27/56
63/63  0s 559us/step - Recall: 0.7838 - loss: 0.064
2 - val_Recall: 0.7838 - val_loss: 0.0698
Epoch 28/56
63/63  0s 552us/step - Recall: 0.7872 - loss: 0.063
5 - val_Recall: 0.7883 - val_loss: 0.0693
Epoch 29/56
63/63  0s 553us/step - Recall: 0.7939 - loss: 0.062
9 - val_Recall: 0.7838 - val_loss: 0.0687
Epoch 30/56
63/63  0s 602us/step - Recall: 0.7928 - loss: 0.062
3 - val_Recall: 0.7838 - val_loss: 0.0683
Epoch 31/56
63/63  0s 556us/step - Recall: 0.7984 - loss: 0.061
7 - val_Recall: 0.7838 - val_loss: 0.0678
Epoch 32/56
63/63  0s 536us/step - Recall: 0.7984 - loss: 0.061
1 - val_Recall: 0.7838 - val_loss: 0.0674
Epoch 33/56
63/63  0s 517us/step - Recall: 0.8041 - loss: 0.060
6 - val_Recall: 0.7928 - val_loss: 0.0670
Epoch 34/56
63/63  0s 535us/step - Recall: 0.8086 - loss: 0.060
1 - val_Recall: 0.7928 - val_loss: 0.0666
Epoch 35/56
63/63  0s 547us/step - Recall: 0.8131 - loss: 0.059
5 - val_Recall: 0.7883 - val_loss: 0.0662

Epoch 36/56
63/63  **0s** 554us/step - Recall: 0.8086 - loss: 0.059
1 - val_Recall: 0.7883 - val_loss: 0.0659
Epoch 37/56
63/63  **0s** 570us/step - Recall: 0.8097 - loss: 0.058
7 - val_Recall: 0.7928 - val_loss: 0.0656
Epoch 38/56
63/63  **0s** 550us/step - Recall: 0.8187 - loss: 0.058
3 - val_Recall: 0.7928 - val_loss: 0.0653
Epoch 39/56
63/63  **0s** 522us/step - Recall: 0.8164 - loss: 0.057
8 - val_Recall: 0.7928 - val_loss: 0.0650
Epoch 40/56
63/63  **0s** 545us/step - Recall: 0.8176 - loss: 0.057
4 - val_Recall: 0.7928 - val_loss: 0.0646
Epoch 41/56
63/63  **0s** 576us/step - Recall: 0.8198 - loss: 0.057
1 - val_Recall: 0.7928 - val_loss: 0.0644
Epoch 42/56
63/63  **0s** 527us/step - Recall: 0.8198 - loss: 0.056
7 - val_Recall: 0.7973 - val_loss: 0.0642
Epoch 43/56
63/63  **0s** 537us/step - Recall: 0.8232 - loss: 0.056
4 - val_Recall: 0.7973 - val_loss: 0.0639
Epoch 44/56
63/63  **0s** 512us/step - Recall: 0.8232 - loss: 0.056
0 - val_Recall: 0.8018 - val_loss: 0.0636
Epoch 45/56
63/63  **0s** 539us/step - Recall: 0.8209 - loss: 0.055
7 - val_Recall: 0.8018 - val_loss: 0.0634
Epoch 46/56
63/63  **0s** 541us/step - Recall: 0.8277 - loss: 0.055
3 - val_Recall: 0.8018 - val_loss: 0.0631
Epoch 47/56
63/63  **0s** 536us/step - Recall: 0.8209 - loss: 0.055
1 - val_Recall: 0.8063 - val_loss: 0.0630
Epoch 48/56
63/63  **0s** 541us/step - Recall: 0.8277 - loss: 0.054
8 - val_Recall: 0.8063 - val_loss: 0.0628
Epoch 49/56
63/63  **0s** 529us/step - Recall: 0.8266 - loss: 0.054
5 - val_Recall: 0.8063 - val_loss: 0.0626
Epoch 50/56
63/63  **0s** 533us/step - Recall: 0.8266 - loss: 0.054
2 - val_Recall: 0.8063 - val_loss: 0.0624
Epoch 51/56
63/63  **0s** 524us/step - Recall: 0.8300 - loss: 0.053
9 - val_Recall: 0.8018 - val_loss: 0.0622
Epoch 52/56
63/63  **0s** 550us/step - Recall: 0.8255 - loss: 0.053
6 - val_Recall: 0.8153 - val_loss: 0.0620
Epoch 53/56

```


63/63 ————— 0s 545us/step - Recall: 0.8288 - loss: 0.053
3 - val_Recall: 0.8153 - val_loss: 0.0618
Epoch 54/56
63/63 ————— 0s 566us/step - Recall: 0.8277 - loss: 0.053
1 - val_Recall: 0.8198 - val_loss: 0.0616
Epoch 55/56
63/63 ————— 0s 560us/step - Recall: 0.8288 - loss: 0.052
8 - val_Recall: 0.8198 - val_loss: 0.0615
Epoch 56/56
63/63 ————— 0s 546us/step - Recall: 0.8300 - loss: 0.052
6 - val_Recall: 0.8153 - val_loss: 0.0614

```

```
In [112... print("Time taken in seconds ",end=start)
```

Time taken in seconds 2.5611839294433594

```
In [113... plot(history, 'loss')
```



```
In [ ]:
```

```
In [ ]:
```

Lets check the model performance ofr model_4 on training and validation data respectively

```
In [114... model_4_train_perf = model_performance_classification(model_4,X_train,
```

```
model_4_train_perf
```

500/500 ————— 0s 146us/step

Out [114...

	Accuracy	Recall	Precision	F1 Score
0	0.989563	0.914978	0.983822	0.946428

In [115...

```
model_4_val_perf = model_performance_classification(model_4,X_val,y_val)
model_4_val_perf
```

125/125 ————— 0s 240us/step

Out [115...

	Accuracy	Recall	Precision	F1 Score
0	0.9885	0.906996	0.981184	0.940598

In [116...

```
y_train_pred_4 = model_4.predict(X_train)
y_val_pred_4 = model_4.predict(X_val)
```

500/500 ————— 0s 154us/step

125/125 ————— 0s 150us/step

Lets check the classification report of model_4 on raining and validation data respectively.

In [117...

```
print("Classification Report - Train data Model_4", end="\n\n")
cr_train_model_4 = classification_report(y_train,y_train_pred_4 > 0.5)
print(cr_train_model_4)
```

Classification Report - Train data Model_4

	precision	recall	f1-score	support
0.0	0.99	1.00	0.99	15112
1.0	0.98	0.83	0.90	888
accuracy			0.99	16000
macro avg	0.98	0.91	0.95	16000
weighted avg	0.99	0.99	0.99	16000

In [118...

```
print("Classification Report - Validation data Model_4", end="\n\n")
cr_val_model_4 = classification_report(y_val,y_val_pred_4 > 0.5)
print(cr_val_model_4)
```

Classification Report – Validation data Model_4

	precision	recall	f1-score	support
0.0	0.99	1.00	0.99	3778
1.0	0.97	0.82	0.89	222
accuracy			0.99	4000
macro avg	0.98	0.91	0.94	4000
weighted avg	0.99	0.99	0.99	4000

Model 4 – Neural Network Using Adam Optimizer

Model Architecture

Model 4 was created to test the impact of using a different optimizer. Earlier models mainly used the SGD optimizer, while Model 4 uses the Adam optimizer.

Architecture

- Input features: 40 sensor variables
- Hidden Layer 1: 64 neurons with ReLU activation
- Hidden Layer 2: 32 neurons with ReLU activation
- Output Layer: 1 neuron with Sigmoid activation
- Total trainable parameters: 4,737

Adam was used because it often converges faster than SGD and can handle gradient updates more efficiently.

Model Loss Results

The loss curve shows that both training loss and validation loss decrease smoothly across epochs.

Key Observations

- Training loss decreases sharply at the beginning and then gradually stabilizes.
- Validation loss also decreases steadily.
- The training and validation loss curves remain close to each other.

- There is no major divergence between the curves.

Interpretation

This suggests that:

- the model is learning effectively,
 - training is stable,
 - the Adam optimizer helps the model converge smoothly,
 - and there is limited evidence of severe overfitting.
-

Model 4 Performance Results

Training Dataset Results

Metric	No Failure (0)	Failure (1)
Precision	0.99	0.98
Recall	1.00	0.83
F1-Score	0.99	0.90

Overall Training Performance

- Accuracy: 99%
 - Macro Average Precision: 0.98
 - Macro Average Recall: 0.91
 - Macro Average F1-Score: 0.95
 - Weighted Average F1-Score: 0.99
-

Validation Dataset Results

Metric	No Failure (0)	Failure (1)
Precision	0.99	0.97
Recall	1.00	0.82
F1-Score	0.99	0.89

Overall Validation Performance

- Accuracy: 99%
 - Macro Average Precision: 0.98
 - Macro Average Recall: 0.91
 - Macro Average F1-Score: 0.94
 - Weighted Average F1-Score: 0.99
-

What These Results Say

Model 4 performs very well overall on both training and validation data.

For the **No Failure** class:

- Validation precision is 0.99.
- Validation recall is 1.00.
- Validation F1-score is 0.99.

This means the model is extremely strong at identifying normal turbine operations.

For the **Failure** class:

- Validation precision is 0.97.
- Validation recall is 0.82.
- Validation F1-score is 0.89.

A failure precision of 0.97 means that when Model 4 predicts a turbine failure, it is usually correct.

A failure recall of 0.82 means the model correctly identifies about 82% of actual turbine failures, but still misses about 18% of failure cases.

The failure F1-score of 0.89 shows that the model has a strong balance between precision and recall, although recall could still be improved.

Why This Is Important

1. Adam Optimizer Produced Stable Learning

The smooth loss curve indicates that Adam helped the model learn efficiently and consistently.

This is important because stable training improves confidence that the model is learning meaningful turbine sensor patterns rather than reacting to noise.

2. High Precision Reduces False Alarms

The validation precision for failures is 0.97.

This means most predicted failures are true failures.

From an operational standpoint, this reduces:

- unnecessary inspections,
 - unnecessary maintenance activity,
 - and wasted technician time.
-

3. Recall Still Needs Improvement

The failure recall is 0.82.

This means the model misses around 18% of actual failures.

Since missed failures are the most expensive error in this project, this remains a key limitation.

4. Accuracy Alone Is Not Enough

Although Model 4 has 99% accuracy, the dataset is highly imbalanced.

Since failures represent only a small portion of the data, accuracy can look very strong even when some failures are missed.

For this project, recall and F1-score are more meaningful than accuracy alone.

Business Impact

Model 4 can support predictive maintenance by identifying most turbine failures before breakdown.

Positive Business Impact

- Helps identify high-risk turbines early.
- Reduces unexpected breakdowns.
- Supports maintenance planning.
- Minimizes unnecessary inspections due to high precision.
- Improves operational reliability.

Remaining Business Risk

Model 4 still misses about 18% of actual failures.

These missed failures may result in:

- generator breakdown,
- expensive replacement costs,
- unplanned downtime,
- emergency maintenance,
- and production losses.

Comparison to Previous Models

Metric	Model 1 Validation	Model 2 Validation	Model 3 Validation	Model 4 Validation
Accuracy	99%	99%	98%	99%
Failure Precision	0.96	0.95	0.83	0.97
Failure Recall	0.85	0.82	0.89	0.82
Failure F1- Score	0.90	0.88	0.86	0.89

Model 4 has the highest failure precision among the models so far, but its failure recall is lower than Model 3 and Model 1.

This means Model 4 is very reliable when it predicts a failure, but it does not catch as many actual failures as Model 3.

Overall Conclusion

Model 4 is a strong and stable neural network model using the Adam optimizer.

It demonstrates:

- smooth learning,
- high overall accuracy,
- excellent precision for failure prediction,
- and limited signs of overfitting.

However, because the main business objective is to minimize missed failures, Model 4 may not be the best model if Recall is prioritized.

Model 4 is useful when the business wants very reliable failure alerts with fewer false positives, but Model 3 may be more useful when the goal is to catch more actual failures.

In []:

Model 5

This time we will add more layers and dropout while using a different optimizer.

```
In [119... # clears the current Keras session, resetting all layers and models pr
tf.keras.backend.clear_session()
```

```
In [120... #Initializing the neural network
from tensorflow.keras.layers import Dropout
model_5 = Sequential()
model_5.add(Dense(128,activation="relu",input_dim=X_train.shape[1])) #
model_5.add(Dropout(0.5)) #Complete the code to define the dropout rat
model_5.add(Dense(64,activation="relu")) # Complete the code to define
model_5.add(Dense(32, activation = "relu")) # Complete the code to defi
model_5.add(Dense(1,activation="sigmoid")) # Complete the code to defi
```

```
In [121... model_5.summary()
```

Model: "sequential"

Layer (type)	Output Shape	
dense (Dense)	(None, 128)	
dropout (Dropout)	(None, 128)	
dense_1 (Dense)	(None, 64)	
dense_2 (Dense)	(None, 32)	
dense_3 (Dense)	(None, 1)	

Total params: 15,617 (61.00 KB)

Trainable params: 15,617 (61.00 KB)

Non-trainable params: 0 (0.00 B)

```
In [122... optimizer = tf.keras.optimizers.Adam()    # defining Adam as the optim
# model_5.compile(loss='binary_crossentropy', optimizer=optimizer, met
# model_5.compile(loss='binary_crossentropy', optimizer=optimizer, met
# model_5.compile(loss='binary_crossentropy', optimizer=optimizer, met
# model_5.compile(loss='binary_crossentropy', optimizer=optimizer, met
```

In []:

In []:

```
In [123... optimizer = tf.keras.optimizers.SGD()

model_5.compile(
    loss="binary_crossentropy",
    optimizer=optimizer,
    metrics=["Recall"]
)
```

```
In [124... start = time.time()
history = model_5.fit(X_train, y_train, validation_data=(X_val,y_val)
end=time.time()
```

Epoch 1/56

63/63 ————— 0s 2ms/step - Recall: 0.0495 - loss: 0.2604
- val_Recall: 0.0270 - val_loss: 0.1796

Epoch 2/56

63/63 ————— 0s 857us/step - Recall: 0.0349 - loss: 0.198
3 - val_Recall: 0.1396 - val_loss: 0.1586

Epoch 3/56

63/63 ————— 0s 922us/step - Recall: 0.0788 - loss: 0.183
4 - val_Recall: 0.2297 - val_loss: 0.1442

Epoch 4/56

63/63 ————— 0s 813us/step - Recall: 0.1441 - loss: 0.165

2 - val_Recall: 0.3649 - val_loss: 0.1344
Epoch 5/56
63/63  0s 858us/step - Recall: 0.1892 - loss: 0.160
9 - val_Recall: 0.4189 - val_loss: 0.1271
Epoch 6/56
63/63  0s 847us/step - Recall: 0.2432 - loss: 0.150
4 - val_Recall: 0.4550 - val_loss: 0.1194
Epoch 7/56
63/63  0s 965us/step - Recall: 0.2782 - loss: 0.145
5 - val_Recall: 0.4820 - val_loss: 0.1137
Epoch 8/56
63/63  0s 896us/step - Recall: 0.3164 - loss: 0.139
0 - val_Recall: 0.5135 - val_loss: 0.1087
Epoch 9/56
63/63  0s 932us/step - Recall: 0.3502 - loss: 0.137
4 - val_Recall: 0.5450 - val_loss: 0.1051
Epoch 10/56
63/63  0s 921us/step - Recall: 0.3874 - loss: 0.132
3 - val_Recall: 0.5586 - val_loss: 0.1019
Epoch 11/56
63/63  0s 1ms/step - Recall: 0.4122 - loss: 0.1315
- val_Recall: 0.5856 - val_loss: 0.0996
Epoch 12/56
63/63  0s 852us/step - Recall: 0.4459 - loss: 0.125
3 - val_Recall: 0.5991 - val_loss: 0.0965
Epoch 13/56
63/63  0s 972us/step - Recall: 0.4550 - loss: 0.124
1 - val_Recall: 0.6171 - val_loss: 0.0942
Epoch 14/56
63/63  0s 847us/step - Recall: 0.4707 - loss: 0.123
9 - val_Recall: 0.6171 - val_loss: 0.0922
Epoch 15/56
63/63  0s 877us/step - Recall: 0.5079 - loss: 0.115
8 - val_Recall: 0.6261 - val_loss: 0.0904
Epoch 16/56
63/63  0s 960us/step - Recall: 0.4944 - loss: 0.116
5 - val_Recall: 0.6486 - val_loss: 0.0892
Epoch 17/56
63/63  0s 917us/step - Recall: 0.5146 - loss: 0.115
9 - val_Recall: 0.6441 - val_loss: 0.0876
Epoch 18/56
63/63  0s 981us/step - Recall: 0.5327 - loss: 0.112
2 - val_Recall: 0.6712 - val_loss: 0.0863
Epoch 19/56
63/63  0s 895us/step - Recall: 0.5236 - loss: 0.114
4 - val_Recall: 0.6802 - val_loss: 0.0851
Epoch 20/56
63/63  0s 911us/step - Recall: 0.5462 - loss: 0.108
3 - val_Recall: 0.6847 - val_loss: 0.0840
Epoch 21/56
63/63  0s 865us/step - Recall: 0.5518 - loss: 0.108
4 - val_Recall: 0.6937 - val_loss: 0.0830

Epoch 22/56
63/63  **0s** 1ms/step - Recall: 0.5721 - loss: 0.1067
- val_Recall: 0.6937 - val_loss: 0.0821

Epoch 23/56
63/63  **0s** 1000us/step - Recall: 0.5619 - loss: 0.1090
- val_Recall: 0.6937 - val_loss: 0.0810

Epoch 24/56
63/63  **0s** 928us/step - Recall: 0.5800 - loss: 0.1042
- val_Recall: 0.7162 - val_loss: 0.0800

Epoch 25/56
63/63  **0s** 856us/step - Recall: 0.5777 - loss: 0.1054
- val_Recall: 0.7162 - val_loss: 0.0794

Epoch 26/56
63/63  **0s** 920us/step - Recall: 0.6002 - loss: 0.1038
- val_Recall: 0.7162 - val_loss: 0.0786

Epoch 27/56
63/63  **0s** 985us/step - Recall: 0.6137 - loss: 0.0992
- val_Recall: 0.7342 - val_loss: 0.0779

Epoch 28/56
63/63  **0s** 1ms/step - Recall: 0.6137 - loss: 0.0998
- val_Recall: 0.7342 - val_loss: 0.0773

Epoch 29/56
63/63  **0s** 1ms/step - Recall: 0.6047 - loss: 0.1021
- val_Recall: 0.7387 - val_loss: 0.0764

Epoch 30/56
63/63  **0s** 878us/step - Recall: 0.6205 - loss: 0.0972
- val_Recall: 0.7568 - val_loss: 0.0760

Epoch 31/56
63/63  **0s** 963us/step - Recall: 0.6351 - loss: 0.0997
- val_Recall: 0.7523 - val_loss: 0.0753

Epoch 32/56
63/63  **0s** 930us/step - Recall: 0.6239 - loss: 0.0971
- val_Recall: 0.7523 - val_loss: 0.0747

Epoch 33/56
63/63  **0s** 833us/step - Recall: 0.6284 - loss: 0.0959
- val_Recall: 0.7613 - val_loss: 0.0739

Epoch 34/56
63/63  **0s** 996us/step - Recall: 0.6498 - loss: 0.0920
- val_Recall: 0.7658 - val_loss: 0.0734

Epoch 35/56
63/63  **0s** 993us/step - Recall: 0.6498 - loss: 0.0944
- val_Recall: 0.7658 - val_loss: 0.0729

Epoch 36/56
63/63  **0s** 962us/step - Recall: 0.6329 - loss: 0.0946
- val_Recall: 0.7658 - val_loss: 0.0722

Epoch 37/56
63/63  **0s** 953us/step - Recall: 0.6689 - loss: 0.0934
- val_Recall: 0.7658 - val_loss: 0.0718

Epoch 38/56
63/63  **0s** 925us/step - Recall: 0.6441 - loss: 0.0929
- val_Recall: 0.7793 - val_loss: 0.0714

Epoch 39/56

63/63 ————— **0s** 920us/step - Recall: 0.6655 - loss: 0.088
9 - val_Recall: 0.7793 - val_loss: 0.0708
Epoch 40/56

63/63 ————— **0s** 857us/step - Recall: 0.6678 - loss: 0.091
6 - val_Recall: 0.7748 - val_loss: 0.0704
Epoch 41/56

63/63 ————— **0s** 883us/step - Recall: 0.6734 - loss: 0.089
4 - val_Recall: 0.7838 - val_loss: 0.0699
Epoch 42/56

63/63 ————— **0s** 876us/step - Recall: 0.6824 - loss: 0.089
0 - val_Recall: 0.7838 - val_loss: 0.0694
Epoch 43/56

63/63 ————— **0s** 908us/step - Recall: 0.6723 - loss: 0.089
8 - val_Recall: 0.7838 - val_loss: 0.0690
Epoch 44/56

63/63 ————— **0s** 893us/step - Recall: 0.6813 - loss: 0.089
3 - val_Recall: 0.7838 - val_loss: 0.0685
Epoch 45/56

63/63 ————— **0s** 840us/step - Recall: 0.6959 - loss: 0.087
0 - val_Recall: 0.7838 - val_loss: 0.0683
Epoch 46/56

63/63 ————— **0s** 938us/step - Recall: 0.6836 - loss: 0.087
0 - val_Recall: 0.7838 - val_loss: 0.0679
Epoch 47/56

63/63 ————— **0s** 823us/step - Recall: 0.6892 - loss: 0.085
0 - val_Recall: 0.7838 - val_loss: 0.0677
Epoch 48/56

63/63 ————— **0s** 806us/step - Recall: 0.6757 - loss: 0.085
5 - val_Recall: 0.7838 - val_loss: 0.0673
Epoch 49/56

63/63 ————— **0s** 875us/step - Recall: 0.6926 - loss: 0.085
8 - val_Recall: 0.7928 - val_loss: 0.0671
Epoch 50/56

63/63 ————— **0s** 912us/step - Recall: 0.6937 - loss: 0.086
0 - val_Recall: 0.7883 - val_loss: 0.0667
Epoch 51/56

63/63 ————— **0s** 936us/step - Recall: 0.6914 - loss: 0.084
4 - val_Recall: 0.7928 - val_loss: 0.0664
Epoch 52/56

63/63 ————— **0s** 807us/step - Recall: 0.7050 - loss: 0.084
4 - val_Recall: 0.7883 - val_loss: 0.0662
Epoch 53/56

63/63 ————— **0s** 833us/step - Recall: 0.7061 - loss: 0.080
9 - val_Recall: 0.7973 - val_loss: 0.0659
Epoch 54/56

63/63 ————— **0s** 954us/step - Recall: 0.6914 - loss: 0.083
9 - val_Recall: 0.8063 - val_loss: 0.0654
Epoch 55/56

63/63 ————— **0s** 896us/step - Recall: 0.7061 - loss: 0.082
7 - val_Recall: 0.8063 - val_loss: 0.0652
Epoch 56/56


63/63 ————— **0s** 931us/step - Recall: 0.7140 - loss: 0.081

8 - val_Recall: 0.8063 - val_loss: 0.0649

```
In [125... print("Time taken in seconds ",end=start)
```

Time taken in seconds 3.8337950706481934

```
In [126... plot(history,'loss')
```



Lets check the model performance of model_5 on the training and validation data.

```
In [127... model_5_train_perf = model_performance_classification(model_5,X_train,
model_5_train_perf)
```

500/500 ————— 0s 155us/step

```
Out[127... 
```

	Accuracy	Recall	Precision	F1 Score
0	0.986625	0.891694	0.977778	0.929937

```
In [128... model_5_val_perf = model_performance_classification(model_5,X_val,y_val,
model_5_val_perf)
```

125/125 ————— 0s 248us/step

Out [128... **Accuracy** **Recall** **Precision** **F1 Score**

0 0.98725 0.902094 0.972971 0.934294

In [129... `y_train_pred_5 = model_5.predict(X_train)`
`y_val_pred_5 = model_5.predict(X_val)`

500/500  **0s** 168us/step
125/125  **0s** 163us/step

Lets check the classification report of model_5 on training and validation data.

In [130... `print("Classification Report - Train data Model_2", end="\n\n")`
`cr_train_model_5 = classification_report(y_train,y_train_pred_5 > 0.5)`
`print(cr_train_model_5)`

Classification Report - Train data Model_2

	precision	recall	f1-score	support
0.0	0.99	1.00	0.99	15112
1.0	0.97	0.78	0.87	888
accuracy			0.99	16000
macro avg	0.98	0.89	0.93	16000
weighted avg	0.99	0.99	0.99	16000

In [131... `print("Classification Report - Validation data Model_2", end="\n\n")`
`cr_val_model_5 = classification_report(y_val,y_val_pred_5 > 0.5)`
`print(cr_val_model_5)`

Classification Report - Validation data Model_2

	precision	recall	f1-score	support
0.0	0.99	1.00	0.99	3778
1.0	0.96	0.81	0.88	222
accuracy			0.99	4000
macro avg	0.97	0.90	0.93	4000
weighted avg	0.99	0.99	0.99	4000

Model 5 – Deeper Neural Network with Dropout and Adam Optimizer

Model Architecture

Model 5 was built to test whether adding more model complexity, dropout, and the Adam optimizer would improve performance.

Architecture

- Input features: 40 sensor variables
- Hidden Layer 1: 128 neurons with ReLU activation
- Dropout Layer: 50% dropout
- Hidden Layer 2: 64 neurons with ReLU activation
- Hidden Layer 3: 32 neurons with ReLU activation
- Output Layer: 1 neuron with Sigmoid activation
- Total trainable parameters: 15,617

This model has more trainable parameters than earlier models, meaning it has greater learning capacity.

Model Loss Results

The loss curve shows that both training and validation loss decrease across epochs.

Key Observations

- Training loss starts higher and steadily decreases.
- Validation loss also decreases smoothly.
- Validation loss remains lower than training loss, which can happen when dropout is used during training.
- No major upward trend in validation loss is observed.

Interpretation

This suggests that:

- the model is learning effectively,
 - training is stable,
 - dropout helps control overfitting,
 - and the model generalizes reasonably well to validation data.
-

Model 5 Performance Results

Training Dataset Results

Metric	No Failure (0)	Failure (1)
Precision	0.99	0.97
Recall	1.00	0.78
F1-Score	0.99	0.87

Overall Training Performance

- Accuracy: 99%
 - Macro Average Precision: 0.98
 - Macro Average Recall: 0.89
 - Macro Average F1-Score: 0.93
 - Weighted Average F1-Score: 0.99
-

Validation Dataset Results

Metric	No Failure (0)	Failure (1)
Precision	0.99	0.96
Recall	1.00	0.81
F1-Score	0.99	0.88

Overall Validation Performance

- Accuracy: 99%
 - Macro Average Precision: 0.97
 - Macro Average Recall: 0.90
 - Macro Average F1-Score: 0.93
 - Weighted Average F1-Score: 0.99
-

What These Results Say

Model 5 performs very well overall, but it is weaker at detecting failures compared with some earlier models.

For the **No Failure** class:

- Validation precision is 0.99.
- Validation recall is 1.00.
- Validation F1-score is 0.99.

This means the model is extremely strong at identifying normal turbine operations.

For the **Failure** class:

- Validation precision is 0.96.
- Validation recall is 0.81.
- Validation F1-score is 0.88.

A failure precision of 0.96 means that when the model predicts a turbine failure, it is usually correct.

A failure recall of 0.81 means the model correctly identifies about 81% of actual turbine failures, but misses around 19% of failure cases.

The failure F1-score of 0.88 indicates a good balance between precision and recall, but recall is lower than desired for this business problem.

Why This Is Important

1. More Complexity Did Not Necessarily Improve Failure Detection

Model 5 has more layers and more trainable parameters than earlier models.

However, the validation recall for the failure class is only 0.81, which is lower than Model 1 and Model 3.

This shows that increasing model complexity does not automatically improve performance.

2. Dropout Helped Reduce Overfitting

The loss curve does not show a major divergence between training and validation loss.

This suggests that dropout helped prevent severe overfitting, even though the

model is more complex.

3. High Precision Reduces False Alarms

The validation precision for failures is 0.96.

This means most predicted failures are true failures.

From a maintenance standpoint, this reduces:

- unnecessary inspections,
 - wasted technician time,
 - and unnecessary maintenance costs.
-

4. Recall Remains the Key Limitation

The validation recall for failures is 0.81.

This means the model misses approximately 19% of actual failures.

Since False Negatives are the most expensive error in this project, this is an important limitation.

Business Impact

Model 5 can support predictive maintenance, but it may not be the best model for minimizing missed failures.

Positive Business Impact

- Provides reliable failure alerts due to high precision.
- Reduces unnecessary inspections.
- Supports proactive maintenance planning.
- Helps identify many at-risk turbines before breakdown.

Remaining Business Risk

- Misses about 19% of actual failures.
- Missed failures may lead to:

- unexpected generator breakdown,
- expensive replacement costs,
- emergency maintenance,
- operational downtime,
- and production losses.

Comparison to Previous Models

Metric	Model 1 Validation	Model 2 Validation	Model 3 Validation	Model 4 Validation	Model 5 Validation
Accuracy	99%	99%	98%	99%	99%
Failure Precision	0.96	0.95	0.83	0.97	0.96
Failure Recall	0.85	0.82	0.89	0.82	0.81
Failure F1-Score	0.90	0.88	0.86	0.89	0.88

Model 5 has strong precision but lower recall than Model 1 and Model 3.

This means Model 5 is reliable when it predicts a failure, but it misses more actual failures than some earlier models.

Overall Conclusion

Model 5 is a strong model overall, but it is not the strongest model for the main business objective.

It demonstrates:

- high accuracy,
- high precision,
- stable learning,
- and reasonable generalization.

However, its lower failure recall makes it less ideal for a predictive maintenance problem where missing failures is very costly.

Model 5 may be useful when minimizing false alarms is a priority, but Model 3

remains more aligned with the business goal of detecting as many failures as possible.

Model 6

Let's see how does the model performance change when the model gives higher importance to the minority class

```
In [152... # clears the current Keras session, resetting all layers and models pr
tf.keras.backend.clear_session()
tf.keras.backend.clear_session()
```

```
In [153... model_6 = Sequential()
model_6.add(Dense(128,activation="relu",input_dim=X_train.shape[1])) #
model_6.add(Dropout(0.5)) # Complete the code to define the dropout ra
model_6.add(Dense(64,activation="relu")) # Complete the code to define
model_6.add(Dense(32, activation = "relu")) # Complete the code to def
model_6.add(Dense(1,activation="sigmoid")) # Complete the code to defi
```

```
In [154... model_6.summary()
```

Model: "sequential"

Layer (type)	Output Shape	
dense (Dense)	(None, 128)	
dropout (Dropout)	(None, 128)	
dense_1 (Dense)	(None, 64)	
dense_2 (Dense)	(None, 32)	
dense_3 (Dense)	(None, 1)	

Total params: 15,617 (61.00 KB)

Trainable params: 15,617 (61.00 KB)

Non-trainable params: 0 (0.00 B)

```
In [155... optimizer = tf.keras.optimizers.SGD()
# model_6.compile(loss='binary_crossentropy', optimizer=optimizer, met
# model_6.compile(loss='binary_crossentropy', optimizer=optimizer, met
# model_6.compile(loss='binary_crossentropy', optimizer=optimizer, met
# model_6.compile(loss='binary_crossentropy', optimizer=optimizer, met
```

In []:

```
In [156... ptimizer = tf.keras.optimizers.SGD()

model_6.compile(
    loss="binary_crossentropy",
    optimizer=optimizer,
    metrics=["Recall"]
)
```

```
In [157... start = time.time()
history = model_3.fit(X_train, y_train, validation_data=(X_val,y_val)
end=time.time())
```

Epoch 1/56

63/63  **0s** 1ms/step - Recall: 0.9009 - loss: 0.3760
- val_Recall: 0.8874 - val_loss: 0.1310

Epoch 2/56

63/63  **0s** 702us/step - Recall: 0.8975 - loss: 0.3636
- val_Recall: 0.8964 - val_loss: 0.1674

Epoch 3/56

63/63  **0s** 651us/step - Recall: 0.8986 - loss: 0.3704
- val_Recall: 0.8829 - val_loss: 0.1404

Epoch 4/56

63/63  **0s** 653us/step - Recall: 0.8975 - loss: 0.3796
- val_Recall: 0.8919 - val_loss: 0.1477

Epoch 5/56

63/63  **0s** 620us/step - Recall: 0.9043 - loss: 0.3738
- val_Recall: 0.8874 - val_loss: 0.1393

Epoch 6/56

63/63  **0s** 633us/step - Recall: 0.9077 - loss: 0.3616
- val_Recall: 0.8919 - val_loss: 0.1488

Epoch 7/56

63/63  **0s** 581us/step - Recall: 0.9043 - loss: 0.3665
- val_Recall: 0.8874 - val_loss: 0.1412

Epoch 8/56

63/63  **0s** 665us/step - Recall: 0.8964 - loss: 0.3723
- val_Recall: 0.8964 - val_loss: 0.1573

Epoch 9/56

63/63  **0s** 590us/step - Recall: 0.8998 - loss: 0.3595
- val_Recall: 0.8919 - val_loss: 0.1499

Epoch 10/56

63/63  **0s** 591us/step - Recall: 0.9043 - loss: 0.3647
- val_Recall: 0.9009 - val_loss: 0.1417

Epoch 11/56

63/63  **0s** 611us/step - Recall: 0.9054 - loss: 0.3620
- val_Recall: 0.9009 - val_loss: 0.1762

Epoch 12/56

63/63  **0s** 667us/step - Recall: 0.9043 - loss: 0.3641
- val_Recall: 0.8919 - val_loss: 0.1626

Epoch 13/56

63/63  **0s** 639us/step - Recall: 0.9020 - loss: 0.3567
- val_Recall: 0.8919 - val_loss: 0.1789

Epoch 14/56

63/63 ————— **0s** 636us/step - Recall: 0.9009 - loss: 0.367
4 - val_Recall: 0.8919 - val_loss: 0.1935
Epoch 15/56

63/63 ————— **0s** 582us/step - Recall: 0.9077 - loss: 0.365
2 - val_Recall: 0.8964 - val_loss: 0.1597
Epoch 16/56

63/63 ————— **0s** 598us/step - Recall: 0.9043 - loss: 0.356
7 - val_Recall: 0.8919 - val_loss: 0.1281
Epoch 17/56

63/63 ————— **0s** 649us/step - Recall: 0.9009 - loss: 0.363
9 - val_Recall: 0.8919 - val_loss: 0.1374
Epoch 18/56

63/63 ————— **0s** 612us/step - Recall: 0.9009 - loss: 0.358
7 - val_Recall: 0.8919 - val_loss: 0.1474
Epoch 19/56

63/63 ————— **0s** 638us/step - Recall: 0.9009 - loss: 0.356
2 - val_Recall: 0.8874 - val_loss: 0.1264
Epoch 20/56

63/63 ————— **0s** 668us/step - Recall: 0.9032 - loss: 0.352
4 - val_Recall: 0.8919 - val_loss: 0.1525
Epoch 21/56

63/63 ————— **0s** 608us/step - Recall: 0.9054 - loss: 0.362
8 - val_Recall: 0.8919 - val_loss: 0.1517
Epoch 22/56

63/63 ————— **0s** 642us/step - Recall: 0.9054 - loss: 0.346
7 - val_Recall: 0.8874 - val_loss: 0.1267
Epoch 23/56

63/63 ————— **0s** 598us/step - Recall: 0.9009 - loss: 0.361
2 - val_Recall: 0.8874 - val_loss: 0.1417
Epoch 24/56

63/63 ————— **0s** 602us/step - Recall: 0.8964 - loss: 0.361
5 - val_Recall: 0.8919 - val_loss: 0.1486
Epoch 25/56

63/63 ————— **0s** 609us/step - Recall: 0.9032 - loss: 0.363
9 - val_Recall: 0.8874 - val_loss: 0.1457
Epoch 26/56

63/63 ————— **0s** 633us/step - Recall: 0.9009 - loss: 0.365
8 - val_Recall: 0.8829 - val_loss: 0.1167
Epoch 27/56

63/63 ————— **0s** 610us/step - Recall: 0.9009 - loss: 0.357
4 - val_Recall: 0.8874 - val_loss: 0.1417
Epoch 28/56

63/63 ————— **0s** 602us/step - Recall: 0.9009 - loss: 0.365
2 - val_Recall: 0.8874 - val_loss: 0.1341
Epoch 29/56

63/63 ————— **0s** 647us/step - Recall: 0.9043 - loss: 0.362
2 - val_Recall: 0.8874 - val_loss: 0.1373
Epoch 30/56

63/63 ————— **0s** 614us/step - Recall: 0.9088 - loss: 0.348
1 - val_Recall: 0.8874 - val_loss: 0.1328
Epoch 31/56

63/63 ————— **0s** 614us/step - Recall: 0.9020 - loss: 0.355

4 - val_Recall: 0.8874 - val_loss: 0.1277
Epoch 32/56
63/63  0s 698us/step - Recall: 0.9009 - loss: 0.356
2 - val_Recall: 0.8874 - val_loss: 0.1359
Epoch 33/56
63/63  0s 610us/step - Recall: 0.8953 - loss: 0.361
2 - val_Recall: 0.8874 - val_loss: 0.1612
Epoch 34/56
63/63  0s 680us/step - Recall: 0.9009 - loss: 0.337
9 - val_Recall: 0.8874 - val_loss: 0.1295
Epoch 35/56
63/63  0s 590us/step - Recall: 0.9032 - loss: 0.349
3 - val_Recall: 0.8874 - val_loss: 0.1273
Epoch 36/56
63/63  0s 626us/step - Recall: 0.9065 - loss: 0.346
5 - val_Recall: 0.8874 - val_loss: 0.1245
Epoch 37/56
63/63  0s 678us/step - Recall: 0.9065 - loss: 0.349
0 - val_Recall: 0.8874 - val_loss: 0.1449
Epoch 38/56
63/63  0s 633us/step - Recall: 0.9032 - loss: 0.337
8 - val_Recall: 0.8919 - val_loss: 0.1598
Epoch 39/56
63/63  0s 672us/step - Recall: 0.9009 - loss: 0.360
6 - val_Recall: 0.8874 - val_loss: 0.1425
Epoch 40/56
63/63  0s 627us/step - Recall: 0.9065 - loss: 0.343
8 - val_Recall: 0.8919 - val_loss: 0.1213
Epoch 41/56
63/63  0s 652us/step - Recall: 0.9043 - loss: 0.365
2 - val_Recall: 0.8919 - val_loss: 0.1226
Epoch 42/56
63/63  0s 586us/step - Recall: 0.9065 - loss: 0.347
9 - val_Recall: 0.8919 - val_loss: 0.1357
Epoch 43/56
63/63  0s 644us/step - Recall: 0.9133 - loss: 0.351
5 - val_Recall: 0.8919 - val_loss: 0.1459
Epoch 44/56
63/63  0s 644us/step - Recall: 0.9054 - loss: 0.337
6 - val_Recall: 0.8919 - val_loss: 0.1256
Epoch 45/56
63/63  0s 630us/step - Recall: 0.9088 - loss: 0.338
8 - val_Recall: 0.8919 - val_loss: 0.1318
Epoch 46/56
63/63  0s 628us/step - Recall: 0.9088 - loss: 0.339
6 - val_Recall: 0.8919 - val_loss: 0.1304
Epoch 47/56
63/63  0s 590us/step - Recall: 0.9020 - loss: 0.347
0 - val_Recall: 0.8919 - val_loss: 0.1340
Epoch 48/56
63/63  0s 581us/step - Recall: 0.9088 - loss: 0.346
8 - val_Recall: 0.8919 - val_loss: 0.1361

```


Epoch 49/56
63/63 ————— 0s 618us/step - Recall: 0.9065 - loss: 0.339
4 - val_Recall: 0.8874 - val_loss: 0.1149
Epoch 50/56
63/63 ————— 0s 651us/step - Recall: 0.9043 - loss: 0.340
2 - val_Recall: 0.8919 - val_loss: 0.1247
Epoch 51/56
63/63 ————— 0s 608us/step - Recall: 0.9077 - loss: 0.344
2 - val_Recall: 0.8874 - val_loss: 0.1282
Epoch 52/56
63/63 ————— 0s 654us/step - Recall: 0.9054 - loss: 0.346
7 - val_Recall: 0.8964 - val_loss: 0.1544
Epoch 53/56
63/63 ————— 0s 624us/step - Recall: 0.8986 - loss: 0.357
6 - val_Recall: 0.8874 - val_loss: 0.1226
Epoch 54/56
63/63 ————— 0s 612us/step - Recall: 0.9032 - loss: 0.344
8 - val_Recall: 0.8874 - val_loss: 0.1380
Epoch 55/56
63/63 ————— 0s 629us/step - Recall: 0.9009 - loss: 0.358
8 - val_Recall: 0.8874 - val_loss: 0.1127
Epoch 56/56
63/63 ————— 0s 600us/step - Recall: 0.8998 - loss: 0.347
0 - val_Recall: 0.8874 - val_loss: 0.1239

```

```
In [159... print("Time taken in seconds ",end=start)
```

```
Time taken in seconds 2.6208670139312744
```

```
In [160... plot(history,'loss')
```



Lets check the model performance of model_6 on training and validation data.

```
In [161...] model_6_train_perf = model_performance_classification(model_6,X_train,
model_6_train_perf
```

500/500 ————— **0s** 153us/step

```
Out[161...]      Accuracy    Recall    Precision    F1 Score
0  0.263125  0.45198  0.486387  0.235716
```

```
In [162...] model_6_val_perf = model_performance_classification(model_6,X_val,y_val,
model_6_val_perf
```

125/125 ————— **0s** 213us/step

```
Out[162...]      Accuracy    Recall    Precision    F1 Score
0  0.2535  0.464903  0.489643  0.229777
```

```
In [163...] y_train_pred_6 = model_6.predict(X_train)
y_val_pred_6 = model_6.predict(X_val)
```

500/500 ————— **0s** 161us/step

125/125 ————— **0s** 168us/step

Lets check the classification report of model_6 on both training and validation

data.

```
In [164... print("Classification Report – Train data Model_3", end="\n\n")
cr_train_model_6 = classification_report(y_train,y_train_pred_6 > 0.5)
print(cr_train_model_6)
```

Classification Report – Train data Model_3

	precision	recall	f1-score	support
0.0	0.92	0.24	0.38	15112
1.0	0.05	0.66	0.09	888
accuracy			0.26	16000
macro avg	0.49	0.45	0.24	16000
weighted avg	0.88	0.26	0.36	16000

```
In [165... print("Classification Report – Validation data Model_3", end="\n\n")
cr_val_model_6 = classification_report(y_val,y_val_pred_6 > 0.5)
print(cr_val_model_6)
```

Classification Report – Validation data Model_3

	precision	recall	f1-score	support
0.0	0.93	0.23	0.36	3778
1.0	0.05	0.70	0.09	222
accuracy			0.25	4000
macro avg	0.49	0.46	0.23	4000
weighted avg	0.88	0.25	0.35	4000

In []:

Model 6 – Deeper Neural Network with Dropout and Class Weights

Model Architecture

Model 6 was created to test how model performance changes when the neural network gives higher importance to the minority class, which is turbine failure.

Architecture

- Input features: 40 sensor variables

- Hidden Layer 1: 128 neurons with ReLU activation
- Dropout Layer: 50% dropout
- Hidden Layer 2: 64 neurons with ReLU activation
- Hidden Layer 3: 32 neurons with ReLU activation
- Output Layer: 1 neuron with Sigmoid activation
- Total trainable parameters: 15,617

This model has the same deeper structure as Model 5 but uses class weights to make the model pay more attention to failure cases.

Model Loss Results

The loss curve for Model 6 behaves differently from earlier models.

Key Observations

- Training loss remains relatively high and fluctuates across epochs.
- Validation loss is lower than training loss but also fluctuates.
- The loss curves do not show the same smooth improvement observed in some earlier models.
- The unstable loss pattern suggests the model is having difficulty learning a balanced decision boundary.

Interpretation

The combination of a deeper model, dropout, SGD optimizer, and class weights appears to make training unstable. Class weights strongly push the model toward the minority class, but in this case the model overcompensates and predicts too many failures.

Model 6 Performance Results

Training Dataset Results

Metric	No Failure (0)	Failure (1)
Precision	0.92	0.05
Recall	0.24	0.66

F1-Score 0.38 0.09

Overall Training Performance

- Accuracy: 26%
 - Macro Average Precision: 0.49
 - Macro Average Recall: 0.45
 - Macro Average F1-Score: 0.24
 - Weighted Average F1-Score: 0.36
-

Validation Dataset Results

Metric	No Failure (0)	Failure (1)
Precision	0.93	0.05
Recall	0.23	0.70
F1-Score	0.36	0.09

Overall Validation Performance

- Accuracy: 25%
 - Macro Average Precision: 0.49
 - Macro Average Recall: 0.46
 - Macro Average F1-Score: 0.23
 - Weighted Average F1-Score: 0.35
-

What These Results Say

Model 6 performs poorly overall, despite trying to give more importance to failure cases.

For the **No Failure** class:

- Validation Precision is 0.93.
- Validation Recall is only 0.23.
- This means the model fails to correctly identify most normal turbine observations.

For the **Failure** class:

- Validation Precision is 0.05.
- Validation Recall is 0.70.
- F1-score is only 0.09.

The failure recall of 0.70 means the model identifies 70% of actual failures, but the precision of 0.05 means that almost all predicted failures are false alarms.

This indicates that the model is predicting too many observations as failures.

Why This Is Important

1. Class Weights Can Improve Recall, But May Hurt Precision

Class weights are intended to help the model focus more on rare failure cases.

However, in Model 6, the model appears to overcorrect toward predicting failures.

This improves sensitivity to failures somewhat but creates too many false positives.

2. High Recall Alone Is Not Enough

Although the model detects 70% of actual failures, it does so with extremely low precision.

This means maintenance teams would receive too many false failure alerts.

A model like this would be difficult to trust operationally.

3. Accuracy Drops Sharply

Validation accuracy is only 25%, which is much worse than earlier models.

This confirms that Model 6 is not learning a useful separation between failure and no-failure cases.

4. Poor F1-Score Shows Weak Balance

The validation failure F1-score is only 0.09.

This indicates the model does not balance failure detection and false alarms well.

For this project, a useful model needs both:

- strong Recall,
- and reasonable Precision.

Model 6 does not achieve that balance.

Business Impact

Model 6 would not be a good candidate for deployment.

Negative Business Impact

- Produces too many false failure alerts.
- Could lead to excessive unnecessary inspections.
- May overload maintenance teams.
- Reduces trust in the predictive maintenance system.
- Poorly identifies normal turbine operations.
- Does not provide reliable decision support.

Remaining Risk

Although Model 6 catches some failures, it creates too many false alarms to be useful.

- This could increase:
- inspection costs,
 - maintenance workload,
 - operational inefficiency,
 - and alert fatigue.

Comparison to Previous Models

Metric	Model 3 Validation	Model 5 Validation	Model 6 Validation
--------	--------------------	--------------------	--------------------

Accuracy	98%	99%	25%
Failure Precision	0.83	0.96	0.05
Failure Recall	0.89	0.81	0.70
Failure F1-Score	0.86	0.88	0.09

Model 6 performs substantially worse than previous models.

Compared with Model 3:

- Recall is lower.
- Precision is dramatically worse.
- Overall accuracy is much lower.
- F1-score is much weaker.

This shows that heavier class weighting with this deeper architecture is not effective.

Overall Conclusion

Model 6 does not provide a good balance between failure detection and false alarms.

Although the model was designed to focus more on the minority failure class, it overcorrected and produced too many false positives.

Because of its low accuracy, extremely low precision, and poor F1-score, Model 6 is not recommended as the final model.

Model 3 remains stronger for the business objective because it provides higher failure recall while maintaining much better precision.

Model Performance Comparison and Final Model Selection

Now, in order to select the final model, we will compare the performances of all the models for the training and test sets.

Training Performance Comparison

In [166... `# training performance comparison`

```
models_train_comp_df = pd.concat(
    [
        model_0_train_perf.T,
        model_1_train_perf.T,
        model_2_train_perf.T,
        model_3_train_perf.T,
        model_4_train_perf.T,
        model_5_train_perf.T,
        model_6_train_perf.T
    ],
    axis=1,
)
models_train_comp_df.columns = [
    "Model 0",
    "Model 1",
    "Model 2",
    "Model 3",
    "Model 4",
    "Model 5",
    "Model 6"
]
print("Training set performance comparison:")
models_train_comp_df
```

Training set performance comparison:

Out[166...	Model 0	Model 1	Model 2	Model 3	Model 4	Model 5	Model 6
Accuracy	0.989313	0.991500	0.987938	0.983375	0.989563	0.986625	0.263
Recall	0.913786	0.931903	0.902458	0.946151	0.914978	0.891694	0.451
Precision	0.982432	0.985853	0.980119	0.905198	0.983822	0.977778	0.486
F1 Score	0.945145	0.957088	0.937438	0.924538	0.946428	0.929937	0.235

Validation Performance Comparison

In [167... *# Validation performance comparison*

```
models_val_comp_df = pd.concat(
    [
        model_0_val_perf.T,
        model_1_val_perf.T,
        model_2_val_perf.T,
        model_3_val_perf.T,
        model_4_val_perf.T,
        model_5_val_perf.T,
        model_6_val_perf.T
    ],
    axis=1,
)
```

```

    ],
    axis=1,
)
models_val_comp_df.columns = [
    "Model 0",
    "Model 1",
    "Model 2",
    "Model 3",
    "Model 4",
    "Model 5",
    "Model 6"
]
print("Validation set performance comparison:")
models_val_comp_df

```

Validation set performance comparison:

	Model 0	Model 1	Model 2	Model 3	Model 4	Model 5	Model 6
Accuracy	0.988750	0.989500	0.987500	0.984000	0.988500	0.987250	0.987500
Recall	0.911368	0.922365	0.906467	0.940652	0.906996	0.902094	0.906467
Precision	0.979010	0.975123	0.970935	0.912777	0.981184	0.972971	0.970935
F1 Score	0.942291	0.946991	0.936026	0.926191	0.940598	0.934294	0.922910

Checking the performance of the best model on the test set

```

In [170... # best_model = model_0 ## Uncomment this line in case the best model is
# best_model = model_1 ## Uncomment this line in case the best model is
# best_model = model_2 ## Uncomment this line in case the best model is
best_model = model_3 ## Uncomment this line in case the best model is
# best_model = model_4 ## Uncomment this line in case the best model is
# best_model = model_5 ## Uncomment this line in case the best model is
# best_model = model_6 ## Uncomment this line in case the best model is

```

```

In [171... # Test set performance for the best model
best_model_test_perf = model_performance_classification(best_model,X_test)
best_model_test_perf

```

157/157  0s 304us/step

	Accuracy	Recall	Precision	F1 Score
0	0.9872	0.933203	0.945096	0.939058

```

In [172... y_test_pred_best = best_model.predict(X_test)

cr_test_best_model = classification_report(y_test, y_test_pred_best>0.
print(cr_test_best_model)

```

157/157					0s 182us/step
	precision	recall	f1-score	support	
0.0	0.99	0.99	0.99	4718	
1.0	0.90	0.87	0.88	282	
accuracy			0.99	5000	
macro avg	0.95	0.93	0.94	5000	
weighted avg	0.99	0.99	0.99	5000	

In []:

Final Model Performance Summary and Model Selection

Objective of the Project

The primary goal of this project was to build a predictive maintenance model capable of identifying wind turbine generator failures using sensor-based operational data.

Because turbine failures are rare but operationally expensive, the focus was placed on:

- maximizing Recall for the failure class,
- minimizing False Negatives,
- while maintaining strong overall model stability and precision.

Summary of Models Evaluated

Several neural network architectures were developed and compared throughout the project.

The experiments included:

- baseline neural networks,
- deeper architectures,
- dropout regularization,
- class weighting,
- and optimizer tuning.

The following models were tested:

Model	Key Characteristics
Model 0	Baseline neural network
Model 1	Additional hidden layer
Model 2	Dropout regularization
Model 3	Dropout + Class Weights
Model 4	Alternative optimizer setup
Model 5	Deeper architecture with more neurons
Model 6	Aggressive class weighting

Validation Performance Comparison

Model	Accuracy	Recall	Precision	F1 Score
Model 0	0.9888	0.9114	0.9790	0.9423
Model 1	0.9895	0.9224	0.9751	0.9470
Model 2	0.9875	0.9065	0.9709	0.9360
Model 3	0.9840	0.9407	0.9128	0.9262
Model 4	0.9885	0.9070	0.9812	0.9406
Model 5	0.9873	0.9021	0.9730	0.9343
Model 6	0.2535	0.4649	0.4896	0.2298

Final Model Selection – Model 3

Model 3 was selected as the final model because it achieved the strongest balance between:

- failure detection capability,
- generalization performance,
- and operational usefulness.

Although some models achieved slightly higher accuracy or precision, Model 3 produced the highest Recall among the strong-performing models.

This is critically important because:

- missing a true turbine failure is far more costly than generating a false alarm.

Final Test Set Performance – Model 3

Overall Test Performance

Metric	Value
Accuracy	98.72%
Recall	93.32%
Precision	94.51%
F1 Score	93.91%

Detailed Classification Report – Test Data

Class	Precision	Recall	F1-Score	Support
No Failure (0)	0.99	0.99	0.99	4718
Failure (1)	0.90	0.87	0.88	282

What These Results Mean

Strong Overall Performance

The final model achieved:

- 98.72% overall accuracy,
- strong precision,
- strong recall,
- and high F1-score.

This indicates that the model can effectively distinguish between healthy turbines

and failing turbines.

Failure Detection Capability

For the failure class:

- Precision = 0.90
- Recall = 0.87
- F1-score = 0.88

Interpretation

A Recall of 0.87 means:

- the model correctly identifies approximately 87% of true turbine failures.

A Precision of 0.90 means:

- when the model predicts a failure, it is correct about 90% of the time.

This is a strong operational balance.

Generalization Performance

The model performed consistently across:

- training data,
- validation data,
- and test data.

This indicates:

- limited overfitting,
 - stable learning,
 - and strong generalization to unseen turbine data.
-

Why This Is Important

1. Early Failure Detection

The model successfully identifies most turbine failures before breakdown occurs.

This enables:

- proactive maintenance,
 - earlier intervention,
 - and reduced operational risk.
-

2. Reduction of False Negatives

False Negatives are the most dangerous error in this project because missed failures may lead to:

- generator damage,
- emergency shutdowns,
- expensive repairs,
- and production losses.

Model 3 significantly reduces this risk through strong Recall performance.

3. Balanced Operational Performance

The model maintains:

- high Recall,
- strong Precision,
- and high F1-score.

This balance is important because:

- excessive false alarms waste resources,
- while missed failures increase operational risk.

Model 3 provides a practical operational tradeoff.

Business Impact

Predictive Maintenance Optimization

The final model can support:

- predictive maintenance scheduling,
 - earlier maintenance planning,
 - and condition-based monitoring.
-

Reduced Operational Costs

Improved failure prediction can help reduce:

- unplanned downtime,
 - emergency maintenance costs,
 - replacement expenses,
 - and productivity losses.
-

Improved Turbine Reliability

Early detection of abnormal turbine behavior supports:

- improved equipment reliability,
 - safer operations,
 - and more consistent energy production.
-

Improved Resource Allocation

Maintenance teams can:

- prioritize high-risk turbines,
 - optimize inspection schedules,
 - and reduce unnecessary maintenance activity.
-

Key Takeaways

- Sensor data contains strong predictive signals for turbine failures.
- Neural networks performed effectively for predictive maintenance classification.
- Model 3 provided the best business-aligned balance between Recall and

Precision.

- Class weighting improved minority class detection when applied appropriately.
 - Excessive class weighting (Model 6) caused instability and excessive false alarms.
 - Strong Recall is more valuable than maximizing Accuracy alone in this problem.
-

Final Conclusion

Model 3 is the recommended final model for turbine failure prediction.

The model demonstrates:

- high predictive performance,
- strong generalization,
- meaningful failure detection capability,
- and strong operational usefulness.

Its ability to detect most failures while maintaining reasonable precision makes it well suited for real-world predictive maintenance deployment in wind turbine operations.

Actionable Insights and Recommendations

In []:

- Write down actionable insights here
- Write down business recommendations here